


# Data models and query languages

Toni Espinosa 17/02/25



# Why are data models important?

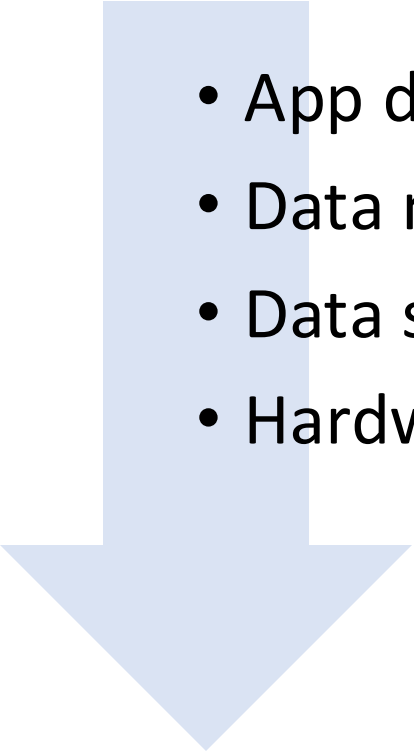
How software is  
designed and  
developed

website users?  
dynamic web content?

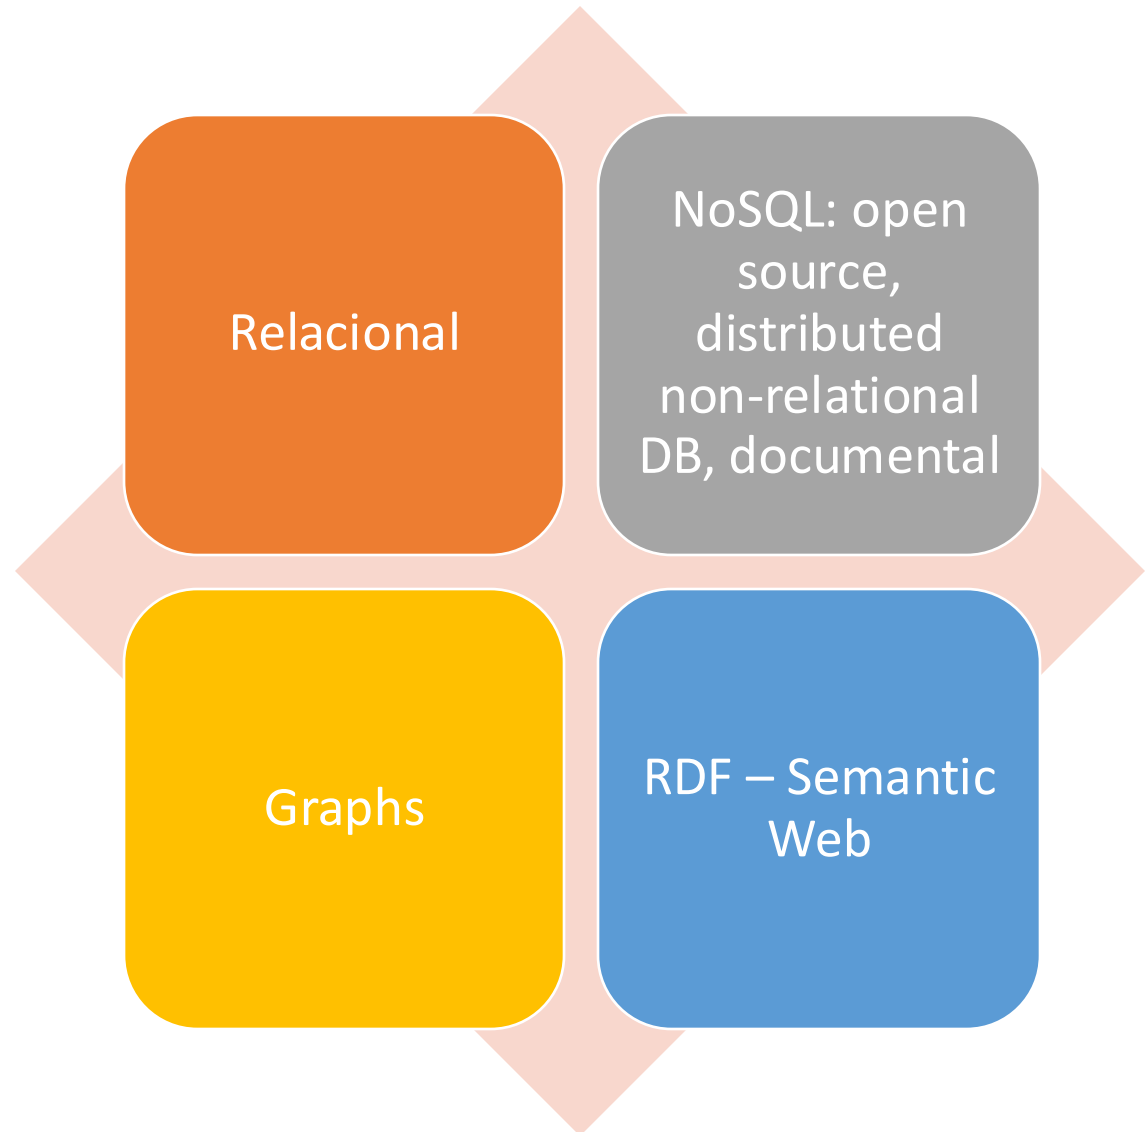
How do we think  
about the  
problem to solve?

Applications are operations  
applied to data -> new data

# Data and points of view

- 
- App developer: Objects, data structures, APIs
  - Data model storage: JSON, XML, relational tables
  - Data system internals: Bytes to memory/disk/cloud
  - Hardware: electric current, magnetic fields, light pulses

# Most common data models



# How to store a Linkedin profile?

Do you see data entities?



**Bill Gates**

Greater Seattle Area | Philanthropy

## **Summary**

Co-chair of the Bill & Melinda Gates Foundation. Chairman, Microsoft Corporation. Voracious reader. Avid traveler. Active blogger.

## **Experience**

Co-chair • Bill & Melinda Gates Foundation  
*2000 – Present*

Co-founder, Chairman • Microsoft  
*1975 – Present*

## **Education**

Harvard University  
*1973 – 1975*

Lakeside School, Seattle

## **Contact Info**

Blog: [thegatesnotes.com](http://thegatesnotes.com)  
Twitter: @BillGates

# How to store a Linkedin profile?

Do you see data entities?

- users
- regions
- industries
- positions
- education
- contact info



**Bill Gates**

Greater Seattle Area | Philanthropy

## **Summary**

Co-chair of the Bill & Melinda Gates Foundation. Chairman, Microsoft Corporation. Voracious reader. Avid traveler. Active blogger.

## **Experience**

Co-chair • Bill & Melinda Gates Foundation  
*2000 – Present*

Co-founder, Chairman • Microsoft  
*1975 – Present*

## **Education**

Harvard University  
*1973 – 1975*

Lakeside School, Seattle

## **Contact Info**

Blog: [thegatesnotes.com](http://thegatesnotes.com)  
Twitter: @BillGates

# Relational model



Bill Gates

Greater Seattle Area | Philanthropy

**Summary**

Co-chair of the Bill & Melinda Gates Foundation. Chairman, Microsoft Corporation. Voracious reader. Avid traveler. Active blogger.

**Experience**

Co-chair • Bill & Melinda Gates Foundation  
2000 – Present

Co-founder, Chairman • Microsoft  
1975 – Present

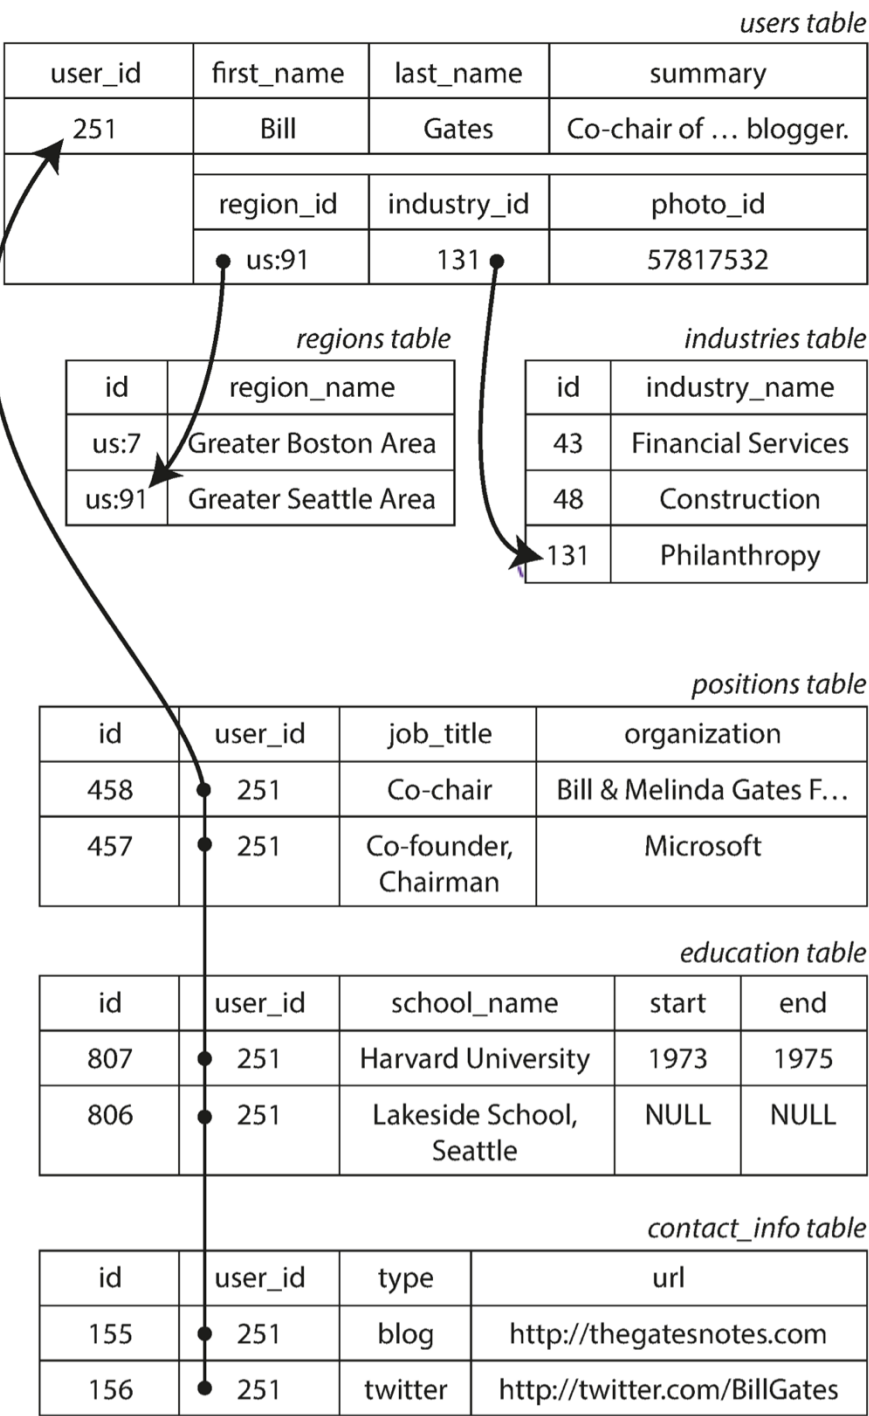
**Education**

Harvard University  
1973 – 1975

Lakeside School, Seattle

**Contact Info**

Blog: thegatesnotes.com  
Twitter: @BillGates



# Relational model limitations

Volume	Write operations	License
Handle large volumes of data: beyond disk capacity	Manage high volume of data write operations per second: write throughput	Avoid expensive license cost of complex products: Oracle



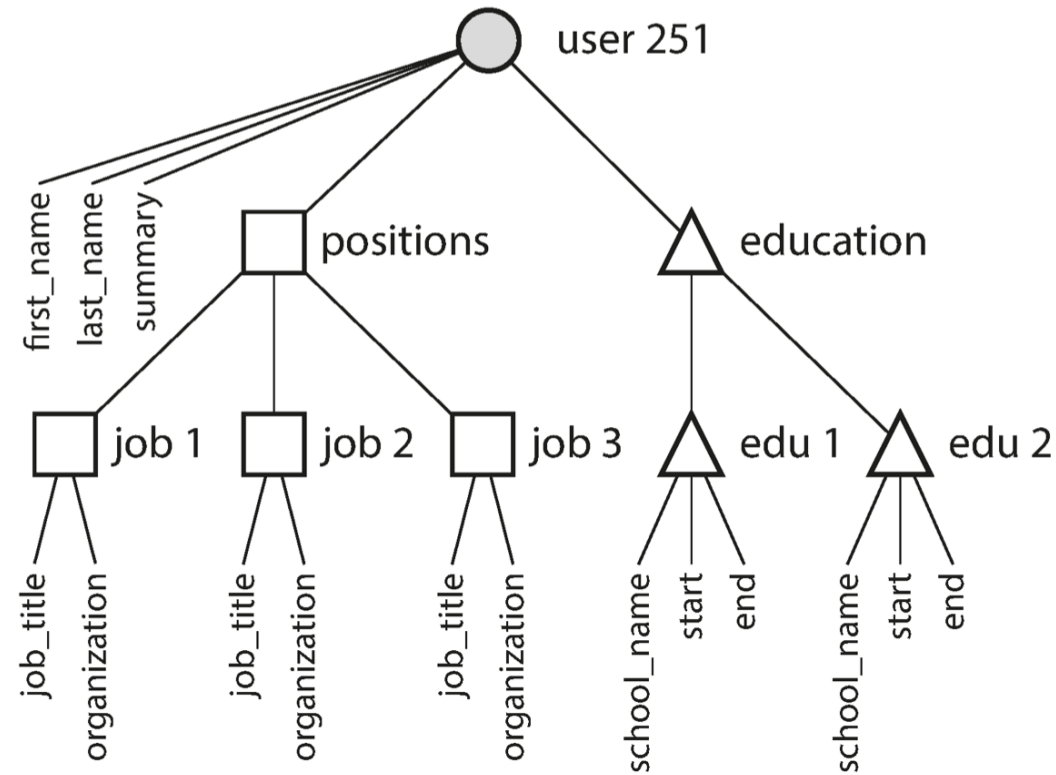
# Relational model limitations

- Handle large volumes of data: beyond disk capacity
- Manage high volume of data write operations per second: write throughput
- Avoid expensive license cost of complex products: Oracle

## New needs:

- Specific search operations
- Dynamic and expressive data models: easy to change schema

# one-to-many relationships from the user profile



# JSON Tree representation is explicit

```
{
  "user_id": 251,
  "first_name": "Bill",
  "last_name": "Gates",
  "summary": "Co-chair of the Bill & Melinda Gates... Active blogger.",
  "region_id": "us:91",
  "industry_id": 131,
  "photo_url": "/p/7/000/253/05b/308dd6e.jpg",
  "positions": [
    {"job_title": "Co-chair", "organization": "Bill & Melinda Gates Foundation"},
    {"job_title": "Co-founder, Chairman", "organization": "Microsoft"}
  ],
  "education": [
    {"school_name": "Harvard University", "start": 1973, "end": 1975},
    {"school_name": "Lakeside School, Seattle", "start": null, "end": null}
  ],
  "contact_info": {
    "blog": "http://thegatesnotes.com",
    "twitter": "http://twitter.com/BillGates"
  }
}
```

Need to explore a complete linkedin profile

Mysql DB:

**multiple queries** or multi-join  
between users and other tables

JSON:

everything in a single query ->  
good **data locality**

# Data locality for queries

- A document is usually stored as a **single continuous string**: encoded as JSON, XML, or plain text
- Need to access different parts of entire document
- If data is in a Single JSON document:
  - faster access because it is **already in memory**
- If data is split across multiple tables:
  - **multiple index lookups** are required to retrieve all data
  - requires more **disk search operations**

# Need to explore a complete linkedin profile

But, if we use JSON as a database  
how do we manage **entity relationships**?



**Bill Gates**

Greater Seattle Area | Philanthropy

## **Summary**

Co-chair of the Bill & Melinda Gates Foundation. Chairman, Microsoft Corporation. Voracious reader. Avid traveler. Active blogger.

## **Experience**

Co-chair • Bill & Melinda Gates Foundation  
*2000 – Present*

Co-founder, Chairman • Microsoft  
*1975 – Present*

## **Education**


Harvard University  
*1973 – 1975*

Lakeside School, Seattle

## **Contact Info**

Blog: [thegatesnotes.com](http://thegatesnotes.com)  
Twitter: @BillGates

# Why do we need numerical identifiers?



```
{
  "user_id": 251,
  "first_name": "Bill",
  "last_name": "Gates",
  "summary": "Co-chair of the Bill & Melinda Gates... Active blogger.",
  "region_id": "us:91",
  "industry_id": 131,
  "photo_url": "/p/7/000/253/05b/308dd6e.jpg",
  "positions": [
    {"job_title": "Co-chair", "organization": "Bill & Melinda Gates Foundation"},
    {"job_title": "Co-founder, Chairman", "organization": "Microsoft"}
  ],
  "education": [
    {"school_name": "Harvard University", "start": 1973, "end": 1975},
    {"school_name": "Lakeside School, Seattle", "start": null, "end": null}
  ],
  "contact_info": {
    "blog": "http://thegatesnotes.com",
    "twitter": "http://twitter.com/BillGates"
  }
}
```

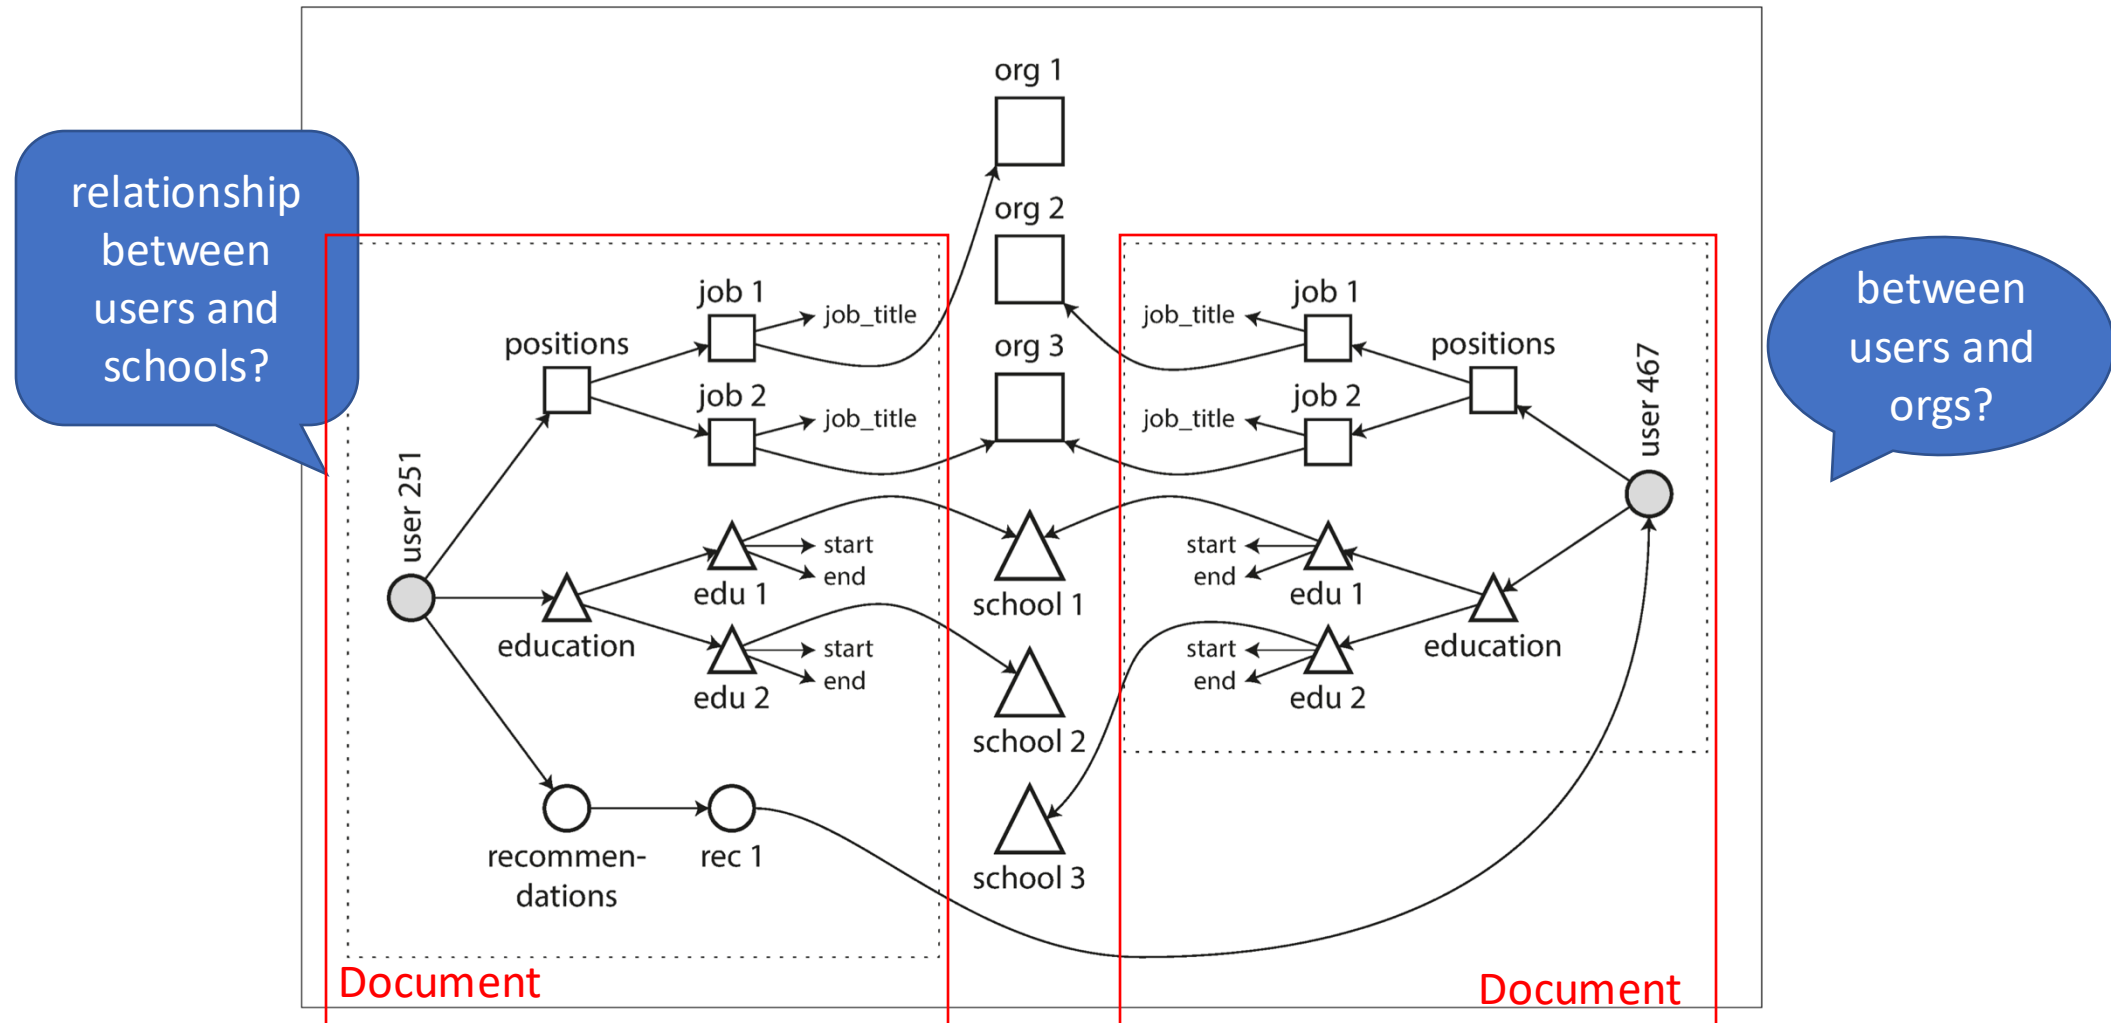
# Entity relationships

region\_id and industry\_id are identifiers. Why?

- Manage **ambiguity**: John Smith
- Easier **updates**: company data can change
- Easier to **translate**: information is isolated
- Easier to **search** for categories than in plain text



# organizations and schools as entities



# Best data model?

Which data model leads to a simpler application code?

one-to-many / many-to-many queries

## Management problems


Documental DBs are not very useful for entity join queries

Relational DB join operation is a good tool for these queries

Which is the best model for my application?

# Ask yourself: which are the most common operations?

- Data is mostly documents
  - one-to-many relationship tree
  - Load the tree just once
  - Not many reference outside the tree
- Some entities and need to analyse relationships
  - most queries are many-to-many
  - Information structure will not change too much in the future
  - Not a large amount of information changing in time
  - Not thousands of entities



# Data change management

Is it possible to modify the data model?

Which is the cost of evolving the models?

# Dynamic data schema

- We need to change data format
- Example: separate user name and surname

how do we deal with this change in each data model type?

document (dynamic) / relational (static)

# documental DB – dynamic model

- Write new documents with new fields: name and surname
- App will manage difference between model documents

```
if(user && user.name && !user.first_name) {  
    //old docs do not have surname  
    user.first_name = user.name.split(" ")[0]  
}
```

# documental DB – dynamic model

- Each change forces a modification of the whole database
- **Migration process**

```
ALTER TABLE users ADD COLUMN first_name text;  
UPDATE users SET first_name =  
    substring_index(name, ' ', 1);
```



# Pros and Cons


- UPDATE has to **modify each element** of a table or more
- Schema modifications can be **very slow**
- Many times data is heterogeneous
  - Many kinds of objects make impractical to have a table per object
  - Data structure is **determined by an external entity** that introduces changes in every new version: REST API

# schema-on-read / schema-on-write

- **schema on read:** Data structure is interpreted when reading
  - Data structure is implicit
  - Similar to Python's dynamic types
  - Maximum flexibility
- **schema on write:** all stored data must follow a pre-defined model
  - Explicit data schema
  - Similar to C++ static types
  - Solid structure

# Link data types and schema

1. JSON google maps
2. BBVA bank account
3. UAB registered students
4. Twitter message
5. Education ministry fellowships



A. Schema on read

B. Schema on write

# New DB and hybrid functionality



PostgreSQL, MySQL i DB2 support  
XML and JSON documents

search, index, modification of files as  
they were tables



RethinkDB

allows document join operations



MongoDB

client join operations



# Query languages

declarative / imperative

SQL, MapReduce, graphs, RDF

# Query language

```
function getDolphins() {  
  var dolphins=[];  
  for(var i=0;i<animals.length;i++){  
    if(animals[i].family==="dolphin"){  
      dolphins.push(animals[i]);  
    }  
  }  
  return dolphins;  
}
```

```
SELECT *  
FROM animals  
WHERE family='dolphin';
```



which is best?

# imperative

```
function getDolphins(){  
  var dolphins=[];  
  for(var i=0;i<animals.length;i++){  
    if(animals[i].family==="dolphin"){  
      dolphins.push(animals[i]);  
    }  
  }  
  return dolphins;  
}
```

Data structure OPERATIONS and  
THEIR ORDER

# declarative

```
SELECT *  
FROM animals  
WHERE family='dolphin';
```

DATA PATTERNS  
DATA TRANSFORMATIONS



# CHOOSE CHARACTERISTICS

1. Hide implementation details
2. Define an animal order
3. Does not need any order in the operations to do
4. Define a strict order of operations, not easy to make it parallel



declarative (SQL)

imperative (Python)

# CSS Declarative Query

```
<ul>
  <li class="selected"> ❶
    <p>Sharks</p> ❷
    <ul>
      <li>Great White Shark</li>
      <li>Tiger Shark</li>
      <li>Hammerhead Shark</li>
    </ul>
  </li>
  <li>
    <p>Whales</p>
    <ul>
      <li>Blue Whale</li>
      <li>Humpback Whale</li>
      <li>Fin Whale</li>
    </ul>
  </li>
</ul>
```

1. selected item marked with CSS class = "selected"

2. <p>Sharks</p> is the title of the page

# Page selected is blue

```
li.selected > p {  
  background-color: blue;  
}
```

**Declarative language:**

which part of the document do we apply the property (blue)

# Selected page in blue?

```
<ul>
  <li class="selected"> ❶
    <p>Sharks</p> ❷
    <ul>
      <li>Great White Shark</li>
      <li>Tiger Shark</li>
      <li>Hammerhead Shark</li>
    </ul>
  </li>
  <li>
    <p>Whales</p>
    <ul>
      <li>Blue Whale</li>
      <li>Humpback Whale</li>
      <li>Fin Whale</li>
    </ul>
  </li>
</ul>
```

```
li.selected > p {
    background-color: blue;
}
```

**li.selected > p**

element patterns to apply blue style:

elements <p> with parent <li>  
with CSS class **selected**

# same example with XSL format

li[@class='selected']/p

same to CSS selector:

li.selected > p

```
<xsl:template match="li[@class='selected']/p">
  <fo:block background-color="blue">
    <xsl:apply-templates/>
  </fo:block>
</xsl:template>
```

# Managing styles with javaScript (DOM)

```
var liElements = document.getElementsByTagName("li");
for (var i = 0; i < liElements.length; i++) {
    if (liElements[i].className === "selected") {
        var children = liElements[i].childNodes;
        for(var j = 0; j < children.length; j++) {
            var child = children[j];

            if(child.nodeType === Node.ELEMENT_NODE&&child.tagName === "P")
            {
                child.setAttribute("style", "background-color: blue");
            }
        }
    }
}
```


Which is the best option?

- CSS (declarative)
- javascript with DOM (imperative)





# Graph data models



# Graph data models

one-to-many relationships:  
document model could be used

many-to-many relationships in  
our data

- Simple cases can be solved with relational models
- Complex connection patterns: graph models

# Graph: nodes and links

## Social graphs:

- Nodes are people, links are relationships (who knows who)

## Web graph

- Nodes are pages connected with html hyperlinks

## Map graph

- Nodes are street crossings, links are streets between them

# Graph algorithms

Google PageRank

Vehicle navigation systems

Social networks: value of users in terms of their connectivity

- Nodes: people, places, actions, events, comments
- Links: friendships, who is commenting, actions related to places, who attended events, ...

# Data models and programming languages

## Models

- Property graph: Neo4j, Titan, InfiniteGraf
- Triple-Store: Datomic, AllegroGraph

## Declarative languages for graphs

- Cypher, SPARQL, Datalog

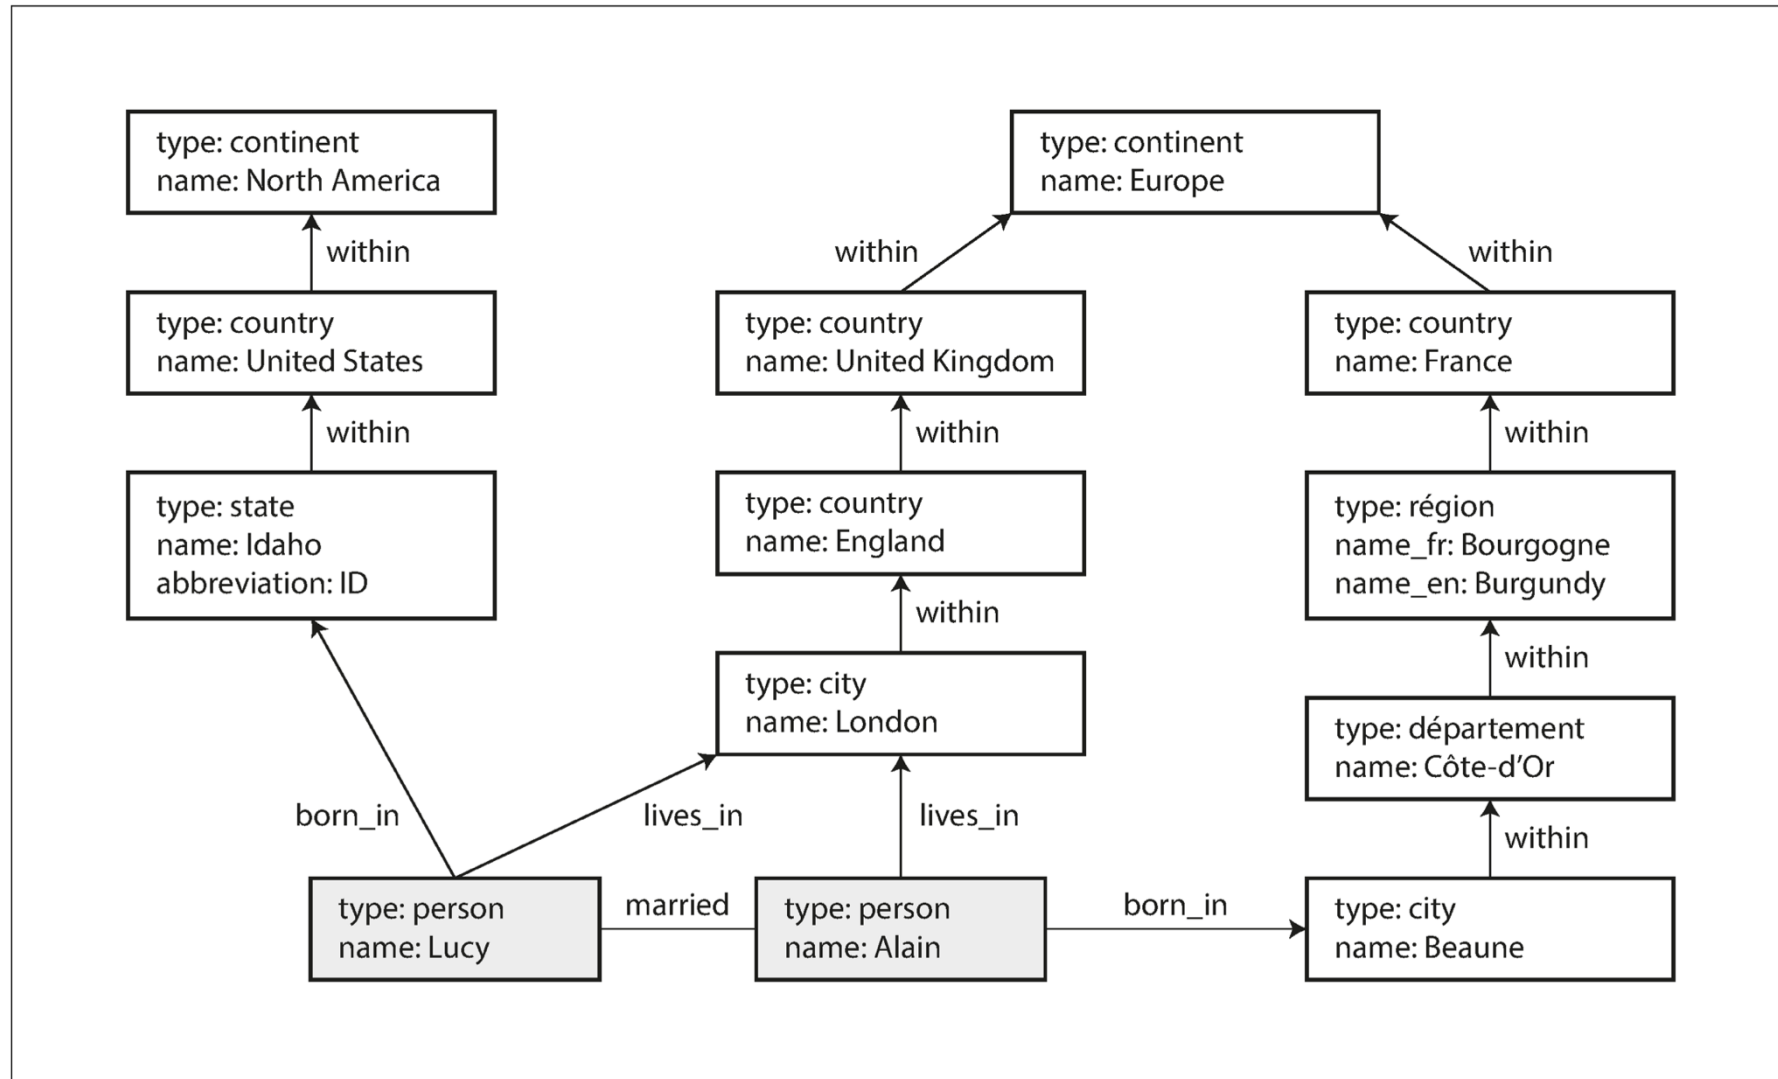
## Imperative languages for graphs

- Gremlin

## Environments for graph processing

- Pregel

# Social network example



# Property graph: tables

## Node/vertex

- Node identifier
- Output links
- Input links
- Properties (key-value)

## Link/edge

- Edge identifier
- initial node → final node
- relationship description
- Properties (key-value)

```
CREATE TABLE vertices (  
    vertex_id integer PRIMARY KEY,  
    properties json  
);
```

```
CREATE TABLE edges (  
    edge_id integer PRIMARY KEY,  
    tail_vertex integer REFERENCES vertices (vertex_id),  
    head_vertex integer REFERENCES vertices (vertex_id),  
    label text,  
    properties json  
);
```

```
CREATE INDEX edges_tails ON edges (tail_vertex);  
CREATE INDEX edges_heads ON edges (head_vertex);
```

# It's easy to express complex relationships

- Each node can be connected with any other with a new link
- Graph **can be traversed** using links
- Use links to label **different kinds of relationships** in the same graph
- Example: we can adapt to different regional structures of countries
  - USA: counties and states
  - France: departements and regions
- Good for evolvability: as you add features to your application, a graph can easily be extended to accommodate changes in data structures.
  - Easy to add new functionality with links, nodes and labels



# How to describe entities

- Describe a place on earth

```
(USA:Location {name: 'United States', type: 'country' } )
```

```
(California:Location {name 'California', type: 'state' } )
```

- Describe a person

```
(Person:Lucy {name: 'Lucy' } )
```

# Describe entity relationships

- California is a US state

`(California) -[:WITHIN]-> (USA)`

- Lucy was born in California

`(Lucy) -: [BORN_IN]-> (California)`

# Cypher language for graphs

## CREATE

```
(NAmerica:Location {name:'North America', type:'continent'}),  
(USA:Location      {name:'United States', type:'country'  } ),  
(Idaho:Location     {name:'Idaho',         type:'state'     } ),  
(Lucy:Person        {name:'Lucy'  } ),
```

# Cypher language for graphs

## CREATE

```
(NAmerica:Location {name:'North America', type:'continent'}),  
(USA:Location      {name:'United States', type:'country'  } ),  
(Idaho:Location     {name:'Idaho',         type:'state'     } ),  
(Lucy:Person        {name:'Lucy' } ),  
(Idaho) -[:WITHIN]-> (USA)  -[:WITHIN]-> (NAmerica),  
(Lucy)  -[:BORN_IN]-> (Idaho)
```

# Cypher graf model queries

Names of people who went from US to Europe?

**search for nodes** with these links:

- **BORN\_IN** link pointing to USA
- **LIVING\_IN** link pointing to Europe
- Our objective is to extract a **specific property**:
  - **name** property

# person BORN\_IN WITHIN Location USA

Find a node (person) that meets condition:

- 1-person has a BORN\_IN link that:
  - has WITHIN links that can be followed to
  - a node of Location type
  - with a name property equal to "USA"

# person LIVES\_IN WITHIN Location Europe

Finds a node (person) that follows two conditions:

1-person has a **BORN\_IN** link with a number of **WITHIN** links that connect to a node of **Location** type with a name property equal to **"United States"**

2-the same person node has a **LIVES\_IN** link with any number of **WITHIN** links that connect to a node of Location type with name = **'Europe'**

# Cypher MATCH

1-person has a **BORN\_IN** link with a number of **WITHIN** links that connect to a node of **Location** type with a name property equal to “**United States**”

2-the same person node has a **LIVES\_IN** link with any number of **WITHIN** links that connect to a node of Location type with name = ‘**Europe**’

## **MATCH**

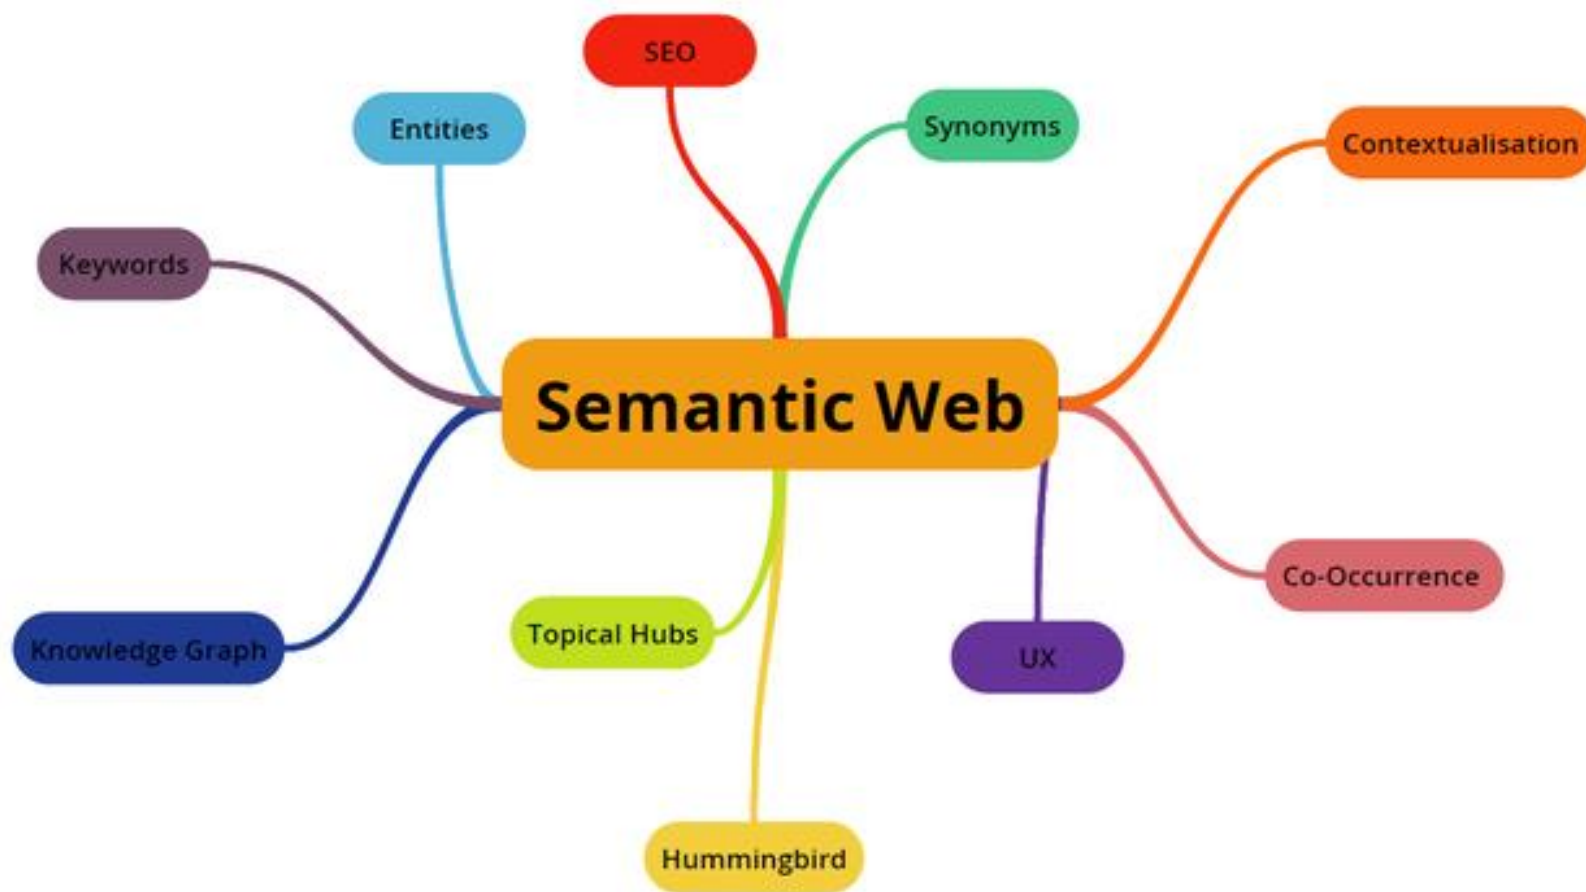
```
(person) -[:BORN_IN]-> () -[:WITHIN*0..]-> (us:Location {name:'United States'}),  
(person) -[:LIVES_IN]-> () -[:WITHIN*0..]-> (eu:Location {name:'Europe'})
```

**RETURN** person.name





# Web as a database



# Semantic web

- Websites publish information for humans (HTML)
- Data web: why not publish information for machines?
- Web seen as a largely distributed DataBase
- Resource Description Framework (RDF): mechanism for publishing data in a consistent format

# RDF example / XML syntax

```
<rdf:RDF xmlns="urn:example:"  
  xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#">
```

```
  <Location rdf:nodeID="idaho">  
    <name>Idaho</name>  
    <type>state</type>  
    <within>  
      <Location rdf:nodeID="usa">  
        <name>United States</name>  
        <type>country</type>  
        <within>  
          <Location rdf:nodeID="america">  
            <name>North America</name>  
            <type>continent</type>  
          </Location>  
        </within>  
      </Location>  
    </within>  
  </Location>  
  
  <Person rdf:nodeID="lucy">  
    <name>Lucy</name>  
    <bornIn rdf:nodeID="idaho"/>  
  </Person>  
</rdf:RDF>
```

# Metadata with JSON-LD

- Machine-readable
- Context: [how to interpret data](#)
  - @id, @type, @language, @value, ...
- `<script type="application/ld+json">` element inside the `<head>` of an HTML page
- *Permissionless*: anyone can take the metadata and run with it.

```
{
  "@context": "https://schema.org",
  "@type": "BlogPosting",
  "headline": "From Shell to Excel – with a little bit of HTTPS",
  "url": "https://csvbase.com/blog/10",
  "description": "Write once, read everywhere",
  "author": {
    "@type": "Person",
    "name": "Cal Paterson",
    "email": "cal@calpaterson.com",
    "url": "https://calpaterson.com/about.html"
  },
  "image": "https://csvbase.com/blog-static/excel.png",
  "datePublished": "2024-08-12",
  "dateCreated": "2024-08-12",
  "dateModified": "2024-08-12"
}
```

What is JSON-LD? <https://www.youtube.com/watch?v=vioCbTo3C-4>

reference:

Chapter 2: Data Models and Query Languages

**“Designing Data Intensive Applications”**

MARTIN KLEPPMAN – O'REILLY – 2017