

In-memory databases

Redis for large data projects

TE 03/04/25

Schedule of sessions

- 3 April: In-memory databases introduction
- 7 April: *Redis Tutorial*
- 10 April : *Redis lab*
- ***Redis Assignment final date: April 27th***

Simple initial starting point

- Data storage requirements in a **single address-space**
- **Complexity** of the solution is **reduced**
 - No need to page information out of memory
 - No need to design consistency methods
- Fast access to large memory space
- Data structures no need to be optimized for disk anymore

Main problems to solve

- **Durability**: what to do when **system shutdowns**?
- **Capacity**: what to do when data needs are **larger than memory space**?

Distributed in-memory solutions

- Shared-Nothing architecture
- Capacity problem:
 - Add more servers to the system
- Durability problem:
 - Create and distribute data copies
 - Increase redundancy
- Oracle Coherence, VoltDB, SAP HANA

Redundancy: now we have new problems!

- Complex queries will need to read all distributed data copies
- Need of a **distributed global join**
 - data update: different values of the same query
 - old values must be overwritten by new in each copy
- Avoid the need of cross-partition join queries
 - NoSQL databases use local data

Technology advances help database systems

- many-core processors
- TB of DRAM
- TB of local SSD memory storage
 - 100s slower than memory, but 10x faster than disk

What do these changes mean?

- If your dataset fits in memory: better use an in-memory DB
 - Much faster than disk
- RAM Memory+SSD become a new standard
 - Already default in Cloud Database services (AWS EBS volumes)
- majority of DB are not so big

High-memory AWS EC2 instances

Name	Logical Processors*	Memory (GiB)	Instance Storage (GB)	Network Bandwidth (Gbps)	EBS Bandwidth (Gbps)
u-3tb1.56xlarge	224	3,072	EBS-Only	50	19
u-6tb1.56xlarge	224	6,144	EBS-Only	100	38
u-6tb1.112xlarge	448	6,144	EBS-Only	100	38
u-6tb1.metal**	448	6,144	EBS-Only	100	38
u-9tb1.112xlarge	448	9,216	EBS-Only	100	38
u-9tb1.metal**	448	9,216	EBS-Only	100	38
u-12tb1.112xlarge	448	12,288	EBS-Only	100	38
u-12tb1.metal**	448	12,288	EBS-Only	100	38
u-18tb1.112xlarge	448	18,432	EBS-Only	100	38
u-18tb1.metal	448	18,432	EBS-Only	100	38
u-24tb1.112xlarge	448	24,576	EBS-Only	100	38
u-24tb1.metal	448	24,576	EBS-Only	100	38



- Memory as main storage support
 - If dataset is bigger than memory: cannot use Redis
 - Not a substitute for general DBMS, good for memory intensive apps
- Disk only for persistence
- Data structure oriented: more complex than simple {key,value} stores
- No need to switch to Redis: use it for partial solutions
- Open source + Redis Labs support
- Redis annual conference: redislabs.com/redisconf/sessions

Why Redis?

- Available as a service in cloud providers (aws,google,azure)
 - <https://aws.amazon.com/redis>
- Who's using Redis? Twitter, Snapchat, Github, Coinbase, Shopify, ...
 - <https://techstacks.io/tech/redis>
- Project presentations from Redisconf
 - Data warehousing – IBM
 - Machine learning – RedisAI – Atlassian
 - Redis on the 5G Edge – Verizon
 - Microservices and Redis – AWS
 - Redis and Apache Kafka – Microsoft

Features and functionality of databases

Name	Type	Data storage options	Query types	Additional features
Redis	In-memory non-relational database	Strings, lists, sets, hashes, sorted sets	Commands for each data type for common access patterns, with bulk operations, and partial transaction support	Publish/Subscribe, master/slave replication, disk persistence, scripting (stored procedures)
memcached	In-memory key-value cache	Mapping of keys to values	Commands for create, read, update, delete, and a few others	Multithreaded server for additional performance
MySQL	Relational database	Databases of tables of rows, views over tables, spatial and third-party extensions	SELECT, INSERT, UPDATE, DELETE, functions, stored procedures	ACID compliant (with InnoDB), master/slave and master/master replication
PostgreSQL	Relational database	Databases of tables of rows, views over tables, spatial and third-party extensions, customizable types	SELECT, INSERT, UPDATE, DELETE, built-in functions, custom stored procedures	ACID compliant, master/slave replication, multi-master replication (third party)
MongoDB	On-disk non-relational document store	Databases of tables of schema-less BSON documents	Commands for create, read, update, delete, conditional queries, and more	Supports map-reduce operations, master/slave replication, sharding, spatial indexes

No tables! individual objects in memory

RELATIONAL TABLE

short	website
uab	www.uab.cat/enginyeria
google	www.google.com
yahoo	www.yahoo.com

REDIS IN-MEMORY OBJECTS

uab	www.uab.cat/enginyeria
google	www.google.com
yahoo	www.yahoo.com

Ready for changes

- React to user events with new structures
- Views are dynamic: no need to redeploy new database schema
- No migration process needed



Redis

IN ACTION

Josiah L. Carlson

FOREWORD BY Salvatore Sanfilippo

Redis by example

Redis in Action, Josiah L. Carlson,
Manning, 2013

Build a simple reddit.com

HASH, ZSET, SET



reddit

Redis Ubuntu installation

```
sudo apt install -y redis
```

```
redis-server -v
```

```
>Redis server v=5.0.7 sha=000000:0 malloc=jemalloc-5.2.1  
bits=64 build=636cde3b5c7a3923
```

Interact with Redis CLI

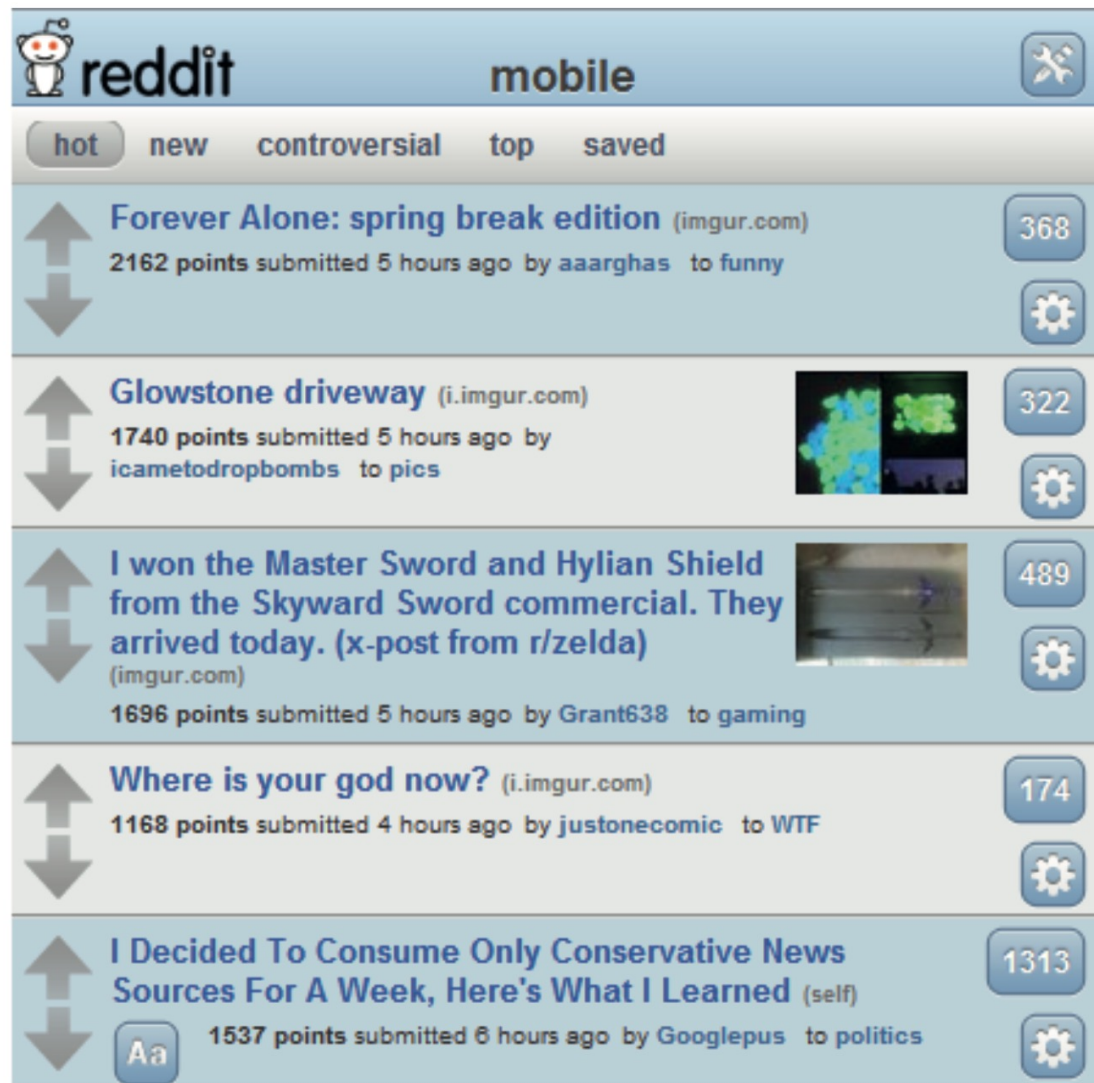
```
$redis-cli
```

```
127.0.0.1>ping
```

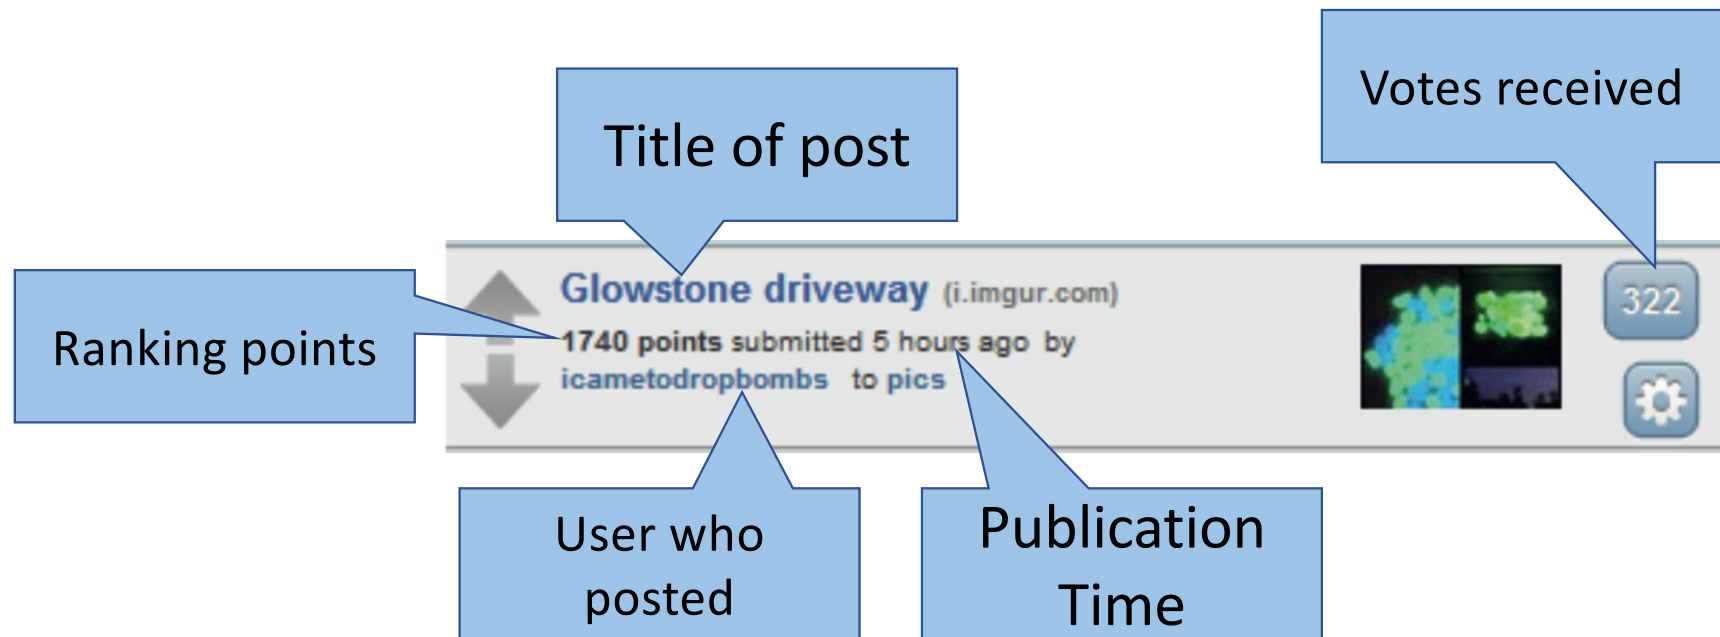
```
PONG
```

```
127.0.0.1>quit
```

Articles can be
voted on by
clicking on up
and down arrows.



Anatomy of a Reddit article

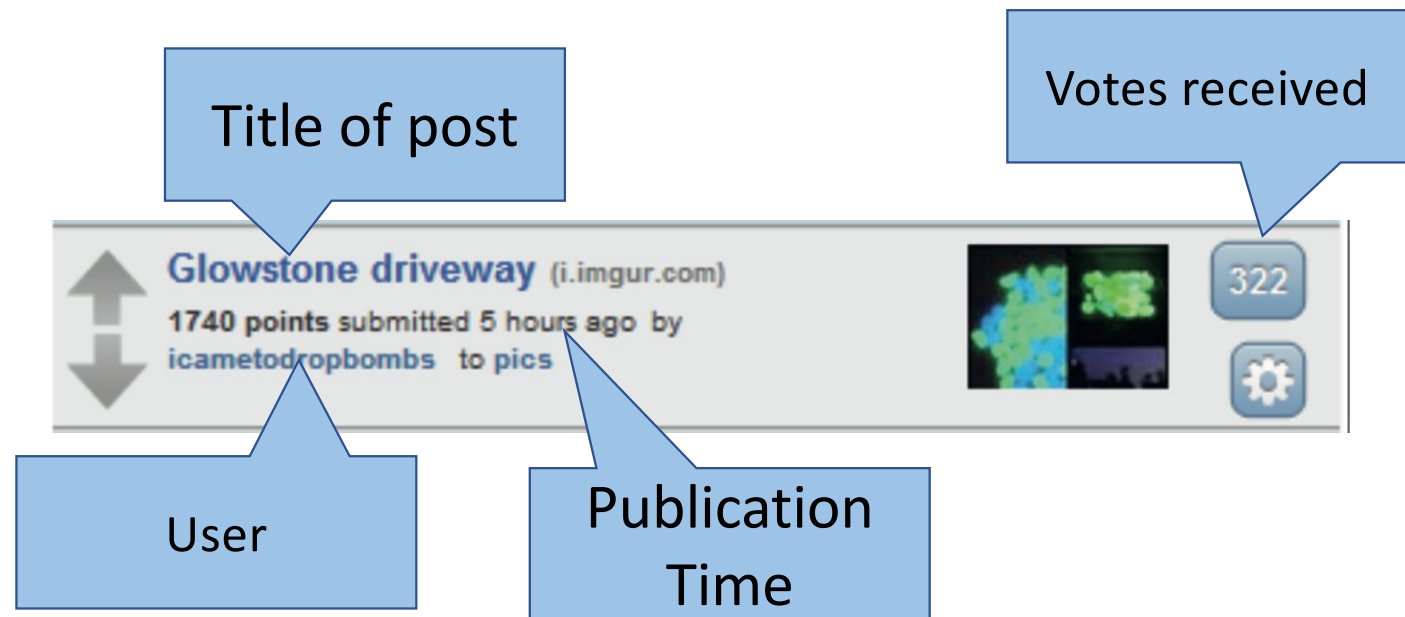


Article list ranking by users vote

- Web page to vote on published articles
- Posts have a ranking score (more votes -> more ranking)
- Posts are published following ranking (highest ranking -> position 1)
- Ranking score will decrease over time automatically (aging)
 - timestamp of post
 - score value = timestamp + number of votes * 500
 - what happens if no one votes for an article?

Data model entities: articles

Articles: title, user, votes received, publication time



Data model entities: ranking

- Article **ranking**: article id, article score



Data model entities: voters

- **Voters:** users who voted for an article
 - users can only vote once per article

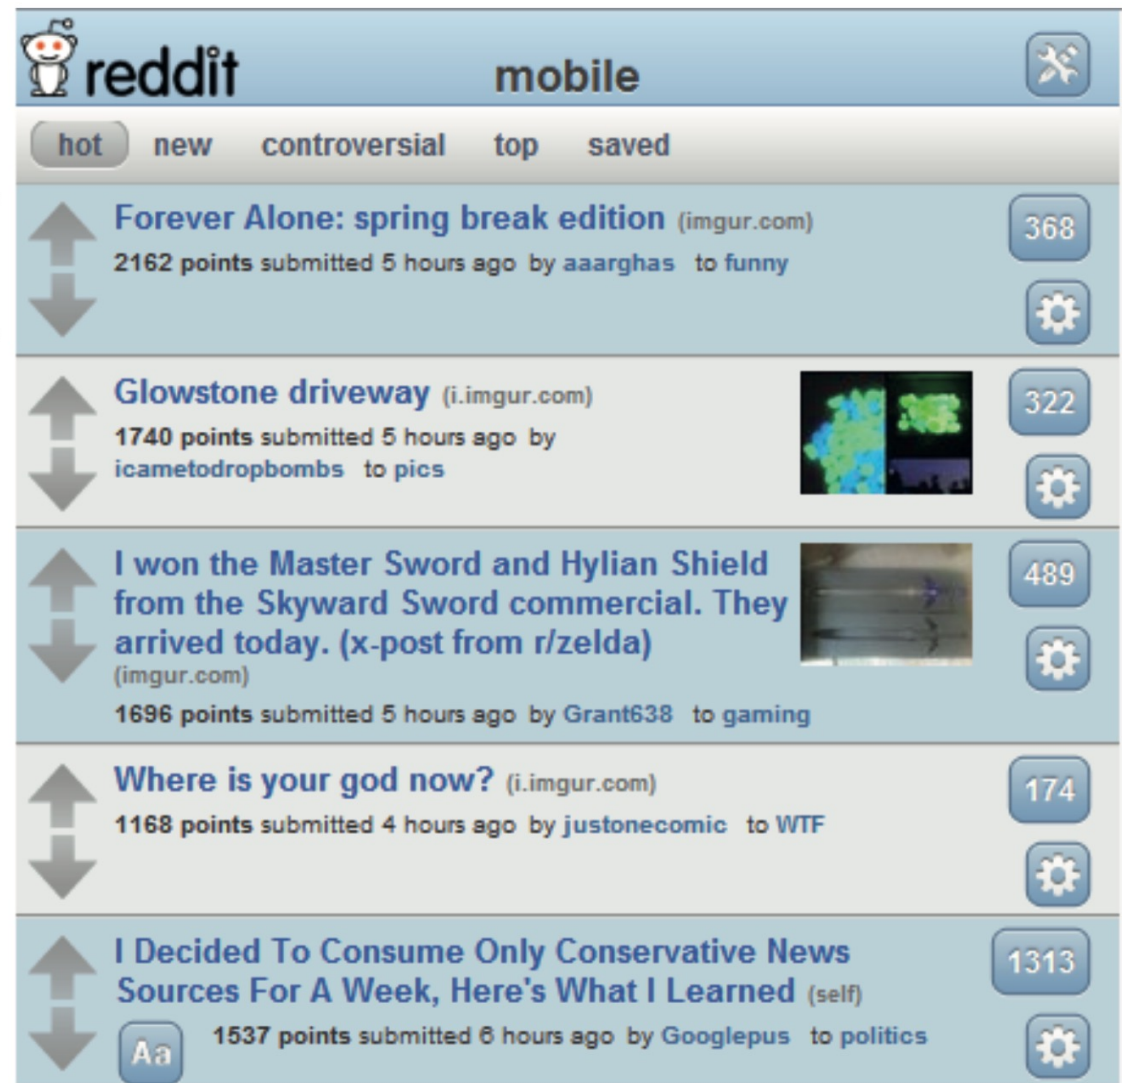


Data model entities

- Articles
 - Article rankings
 - Voters
-
- Redis data structures
 - How to store article, ranking and voters information?
 - How to read article information to implement applications?
 - Implementation of Python functions with Redis CLI/API

Reddit ranking example view

Articles can be
voted on by
clicking on up
and down arrows.



How to store articles?

How to sort articles by score?

Learning how to use Redis Hash and Zset

1 – Redis Hash

To keep elements of our dataset: one element, one object

Redis hash example: Article:00001

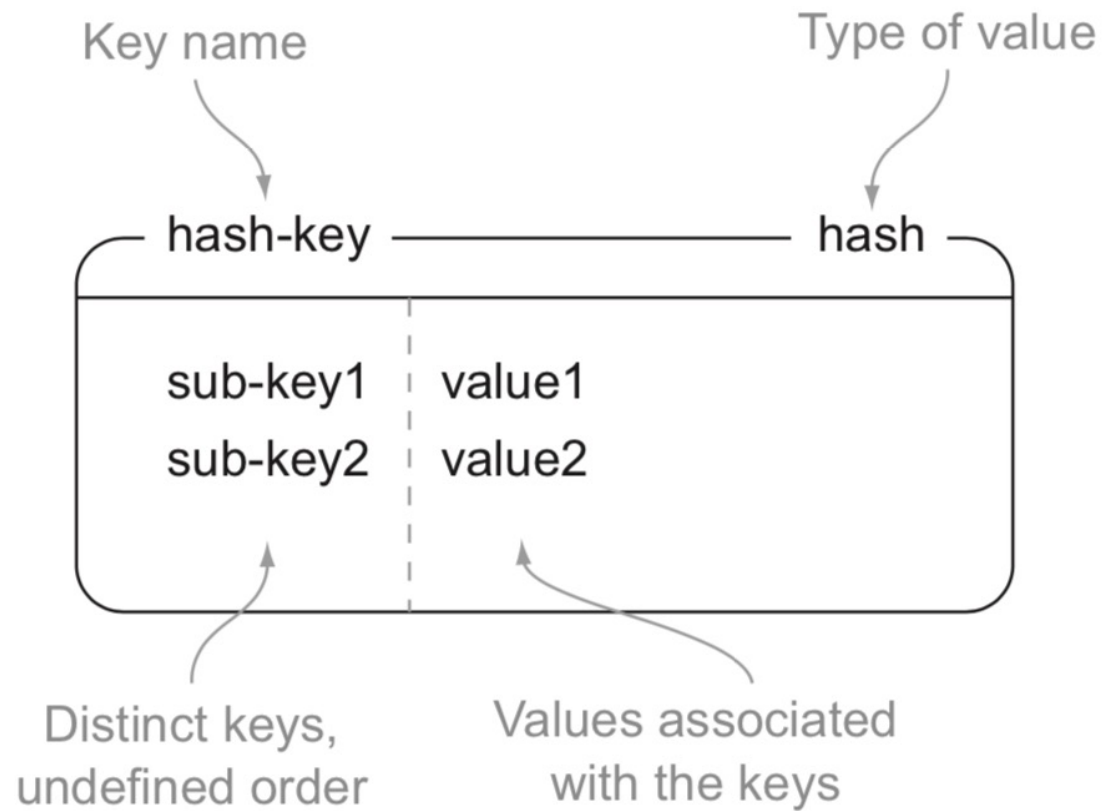
Object Key Name <article:number>

Subkeys for my article objects

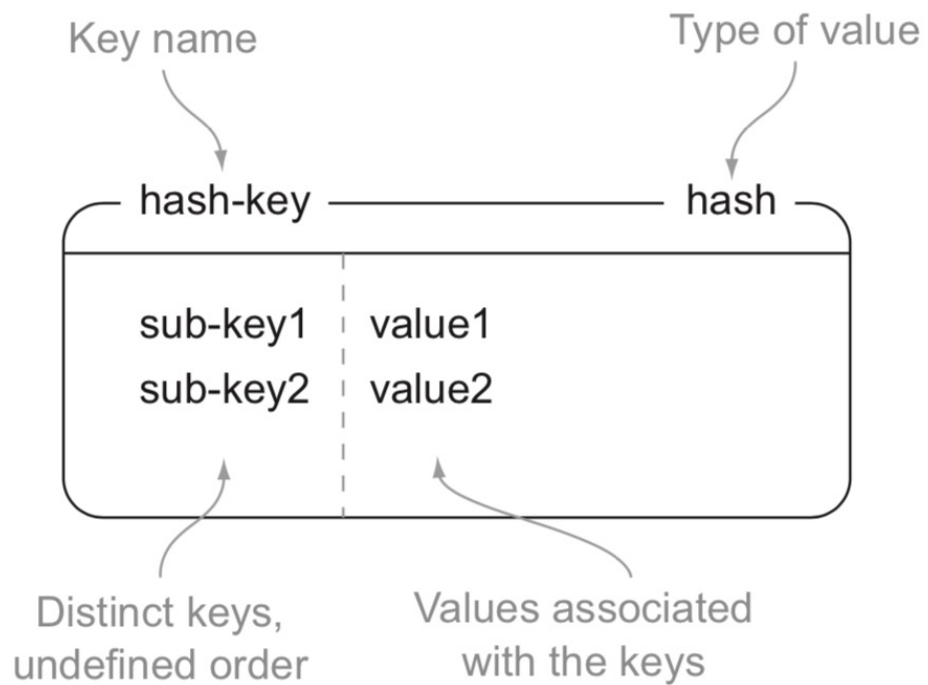
article:00001

- **Title**
 - *title: effects of palm oil in cancer*
- **Link** to article
 - *link: <http://pubmed.org/12345>*
- **user** who posted it
 - *user: toni*
- **Timestamp**
 - *timestamp: 21/10/2023 10:00*
- **Votes** received
 - *Votes: 0*

Redis HASH



Redis HASH



- New hash key: <article:1>
 - title: *effects of palm oil in cancer*
 - votes: 0
- Redis operations
 - **HSET**: create new object
 - **HGET**: get hash-key from object
 - **HMGET**: get all values for an object
 - **HDEL**: eliminate subkey

Article object creation

Redis **HASH object**: object id

Object attributes:

- Title
- user who posted it
- Timestamp
- Votes received

- Object id: ?
- Object information
 - Title
 - User
 - Time
 - Votes

key design <article:00001>

- Very long keys are not a good idea: key lookup will take a long time
 - 12345678912345678
- Very short keys are not a good idea: too many lookup matches
 - 1234 vs user:1234
- Use an informal key schema
- "object-type:id" is a good idea, as in "user:1000" or "article:92617"

Article object creation

Redis **HASH object**: object id

- Title
- user who posted it
- Timestamp
- Votes received

- Object id:
<article:92617>
- Object information
 - Title
 - <title, *genomic analysis*>
 - User
 - <user, user-83721>
 - Time
 - <time, 1331382699>
 - Votes
 - <votes, 1>

Article information structure

Redis **HASH** – <article:number>

- Title
- Link to article
- user who posted it
- Timestamp
- Votes received

article:92617 — hash	
title	Go to statement considered harmful
link	http://goo.gl/kZUSu
poster	user:83271
time	1331382699.33
votes	528

Add a new article to database – Redis CLI

```
HSET <key> <attrib name> <attrib value>
```

```
HSET article:1 title "data analysis" user user:1 votes 1
```

```
HGET <key> <attrib name>
```

```
HGET article:1 title
```

```
"data analysis"
```

```
HMGET <key>
```

```
HMGET article:1
```

Hash operations using Python API

article:92617 ————— hash	
title	Go to statement considered harmful
link	http://goo.gl/kZUSu
poster	user:83271
time	1331382699.33
votes	528

```
myarticle = 'article:92617'
```

- **create new object:**

```
hmset(myarticle,  
      {'title': "data analysis",  
       'time': "01/11/2021 10:10",  
       'votes':1,})
```

- **read all values for an object:**

```
hgetall(myarticle)
```

Add a new article to database - Python

```
now=round(time.time()*1000)
article_id="1234"
article="article:"+article_id
client.hset(article, mapping={
    "title": "data analysis",
    "time": now,
    "votes": 1,
})
```

Hash objects summary

- Good for keeping collections of individual elements of our data model like rows of our articles table
- Must create a key for each object: <article:123456>
- Store attributes with label keys: “title”: “data analysis”
Multi attribute with json format
- Get values at anytime searching for object key

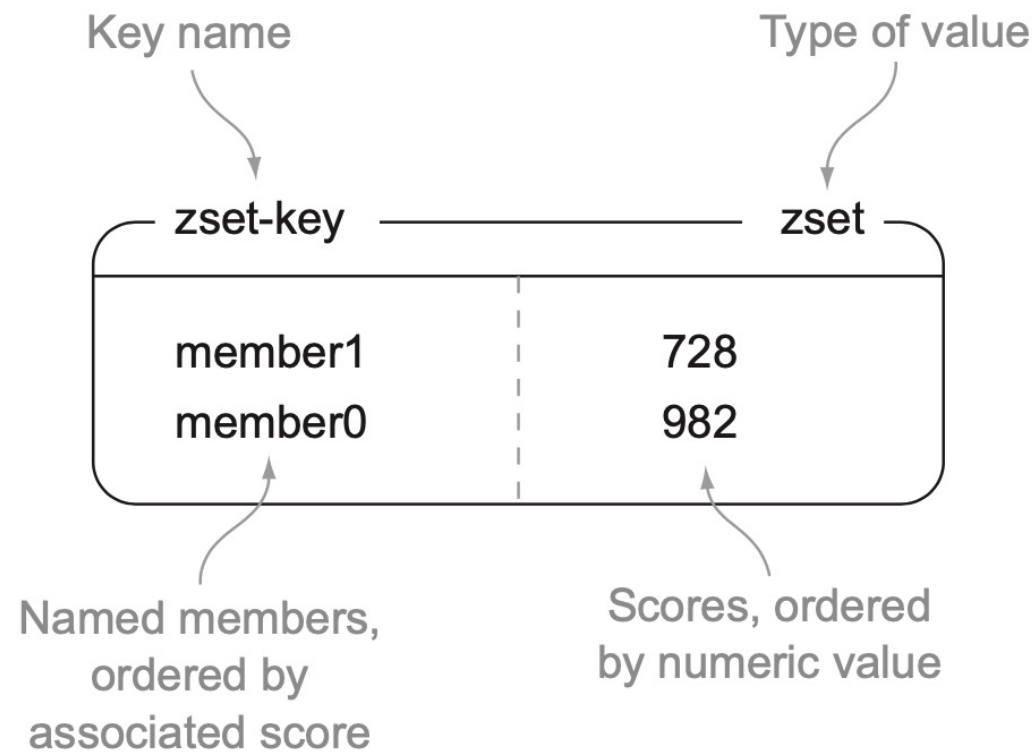
ZSETS – sorted lists of elements

global list of relevant objects

Article ranking

- Must keep a **top list view** of articles with its score
 - Ranking is a sorted collection of our articles
 - Article, score of article
- Get most ranked articles
 - by position: give me the top 10 highest score items
 - by ranking: all articles with score between (900 and 1000)
- use **ZSET**: sorted **set** of elements (by score)

Redis sorted sets **ZSET**



Sorted set of articles with ZSET

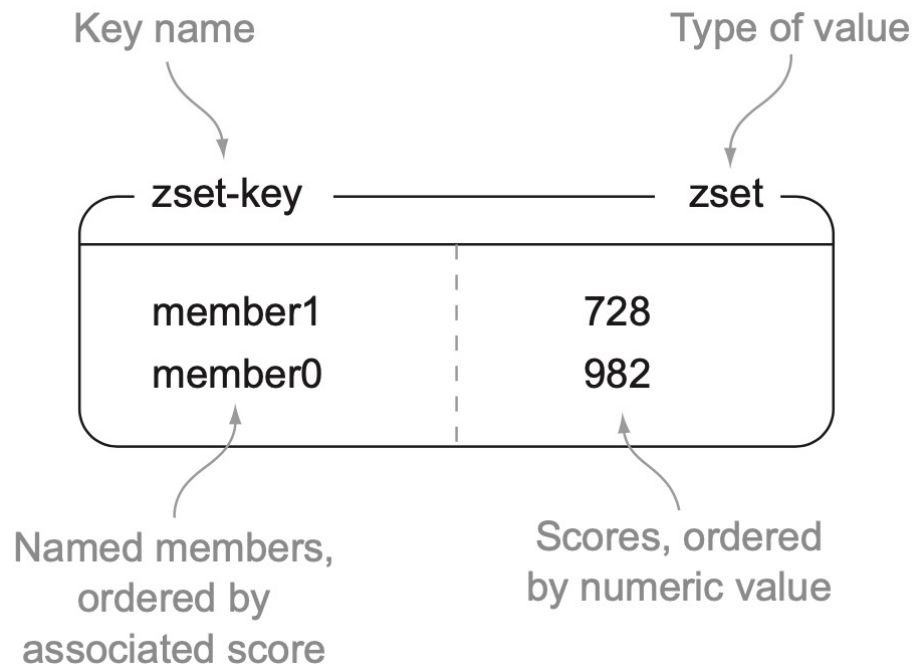
Keep all articles sorted by score

- ZSET name **score:**
- article id (key), article score (value)
 - article:100408, 1332065417.47

score:	zset
article:100635	1332164063.49
article:100408	1332174713.47
article:100716	1332225027.26

A score-ordered ZSET of articles

Redis sorted sets ZSET



Operations with ZSET

- **ZADD**: add a new item
- **ZSCORE**: get score from item
- **ZRANGE**: get TOP items
- **ZRANGEBYSCORE**: get items with some score

Article time of posting with ZSET

Keep the time of article posting

- ZSET name **time:**
- Article id, Linux timestamp
 - article:100635, 1332164063

Timestamp:

- number of seconds since 1/1/1970

time:	zset
article:100408	1332065417.47
article:100635	1332075503.49
article:100716	1332082035.26

A time-ordered ZSET of articles

ZSET operations

new article score

ZADD <key of zset> <member-score member-id>

ZADD score: 500 article:1

get article score

ZSCORE <key of zset> <member-id>

ZSCORE score: article:1

get 10 first elements

ZRANGE <key> <first position> <last position>

ZRANGE score: 0 9

Add a new article to global score - Python

```
myarticle="article:1234"  
votes=500
```

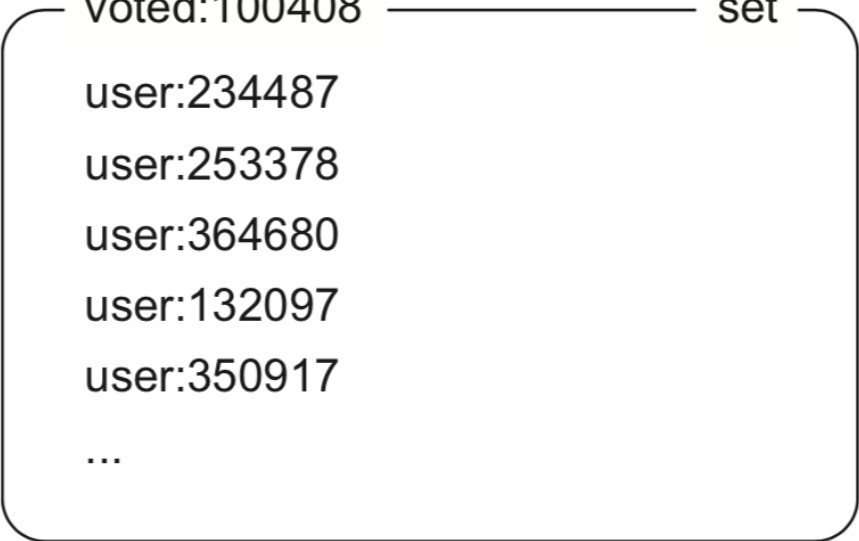
```
score_zset="score:"  
client.zadd(score_zset, {myarticle : votes})
```


SETS: keep groups of unique elements

One element per SET, unique values only

Article voters structure with SET

- List of users who voted for an article
- Prevents users voting more than once for the same article
- One SET for each article
 - *voted:<article id>*

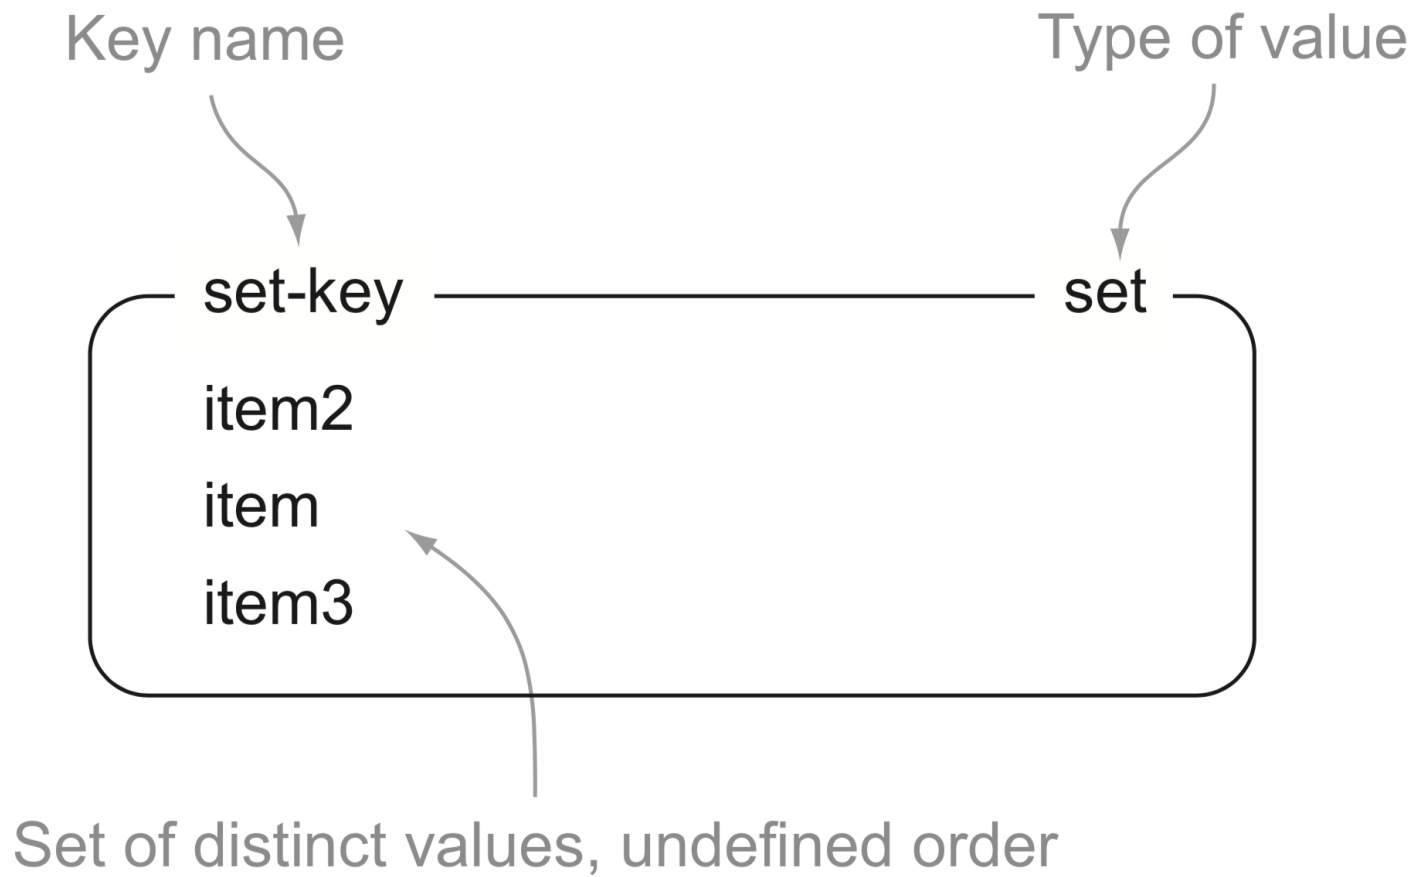


The diagram shows a rounded rectangular box representing a Redis SET. The text inside the box is as follows:

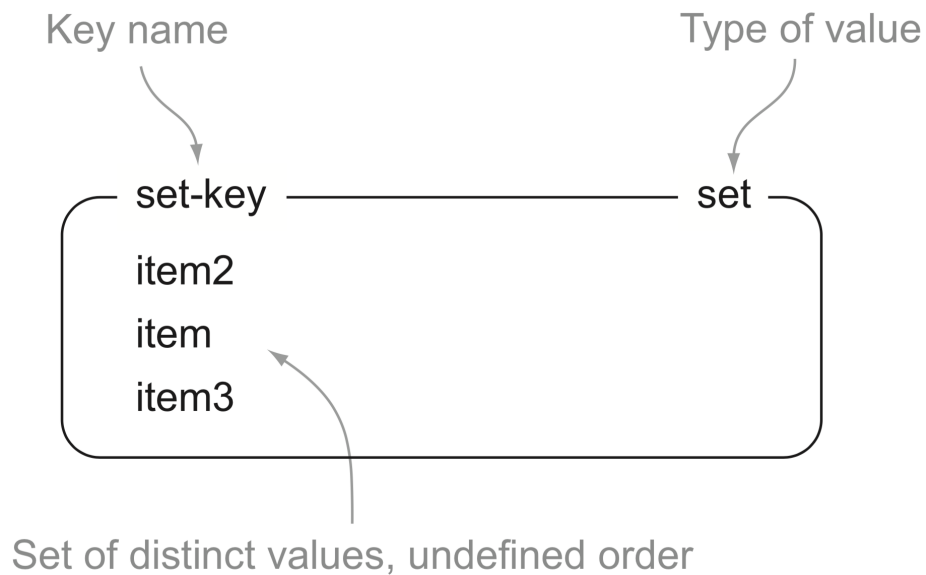
```
voted:100408  
user:234487  
user:253378  
user:364680  
user:132097  
user:350917  
...
```

A horizontal line connects the text 'voted:100408' on the left to the text 'set' on the right, indicating the key and the data type.

Redis SET structure



Redis SET structure



SET common operations:

- **SADD**: Add a new item to SET
- **SREM**: Remove item
- **SISMEMBER**: check if item is in SET
- **SMEMBERS**: get all items from SET
- **SINTER**: intersection
- **SCARD**: cardinality

Add a new user to list of voters for an article

Add user:1234 to the list of article:1 voters

SADD <set key> <member-id>

SADD voted:1 user:1234

get all users who voted for article 1

SMEMBERS <set key>

SMEMBERS voted:1

check if user:1234 voted for article 1

SISMEMBER <key> <member-id>

SISMEMBER voted:1 user:1234

Add a new user to list of voters for an article

```
myarticle= 'article:'+article_id  
voted = 'voted:'+article_id  
user = 'user:1234'  
  
sadd(voted, user)
```


Data structures for our Reddit example

- **Articles**: individual HASH per article (**article:1234**)
 - keep article details
- **Article ranking**: global ZSET for ranking and to store publish time
 - **score**:
 - keep global article ranking
- **Voters**: individual SET per article
 - **voted:1**
 - keep list of voters to limit one vote per user



Posting and voting articles

Reacting to users operations to update data model

Application events

- External
 - user creates a new article
 - user votes for an article
- Internal
 - produce current top 10 ranking articles
- How must our Redis design deal with these events?
 - Application reacts to external events then executes some code
 - Keep always a coherent view of the data model

User creates a new article

1. Build a new article-id
2. Store article attributes in new <article:id> object
3. Create voted SET to allow other users vote for this article
4. Add initial score and posting time for new article

New article – build an article ID

- User posts a new article
- build a new article id

```
counter = counter + 1  
myarticle = 'article:' + counter
```

article:1234

New article – build a new article hash object

```
HMSET  article:1234  
      title:'this is the title'  
      time:1111  
      votes:1
```

New article – build a new article hash object

```
counter = counter + 1
myarticle = 'article:' + counter
now = round(time.time())

client.hset(myarticle, mapping={
    'title': 'this is the article title',
    'time': now,
    'votes': 1,
})
```

New article: new voted SET

Add post user-id to a new **voted:** SET for the new article

```
SADD voted:1234 user:1
```


add article with initial score

- Add initial score to **score:** global ZSET

```
ZADD score: 500 article:1234
```

add article with initial score

- Add initial score to **score:** global ZSET

```
myarticle='article:'+counter
```

```
VOTE_SCORE= 500
```

```
zadd('score:', {myarticle : now + VOTE_SCORE})
```

add article initial posting time

- Add initial posting time to **time:** global ZSET

```
ZADD time: 12345679 article:1234
```

add article initial posting time

- Add initial posting time to **time:** global ZSET

```
now=round(time.time())
```

```
zadd('time:', {myarticle: now})
```

new article operations summary

1. Build an article-id
 - Increment global article counter
2. Add article details into <article:id> HASH
 - HMSET article-id title: 'this is my title' votes:1})
3. Create voted SET
 - Add post user ID to the SET with SADD voted:article-id user id
 - give expiration date of one week with EXPIRE
4. Add initial score and posting time
 - ZADD score: article-id SCORE
 - ZADD time: article-id publication time

```
counter=counter+1
myarticle = 'article:'+counter
now = round(time.time())
client.hset(myarticle, mapping={
    'title': title,
    'time': now,
    'votes':1,
})
```

```
voted = 'voted:'+counter
client.sadd(voted, user)
client.expire(voted, ONE_WEEK_IN_SECONDS)
```

```
client.zadd('score:', {myarticle : now+VOTE_SCORE})
client.zadd('time:', {myarticle : now})
```


User wants to vote for an article


Steps to manage voting action:

1. Verify that article was posted within last week in **article** HASH
2. Add the user to article **voted:** SET
Only one vote per user is allowed
3. Increment the score of the article in **score:** SET
4. Update vote count in **article** HASH

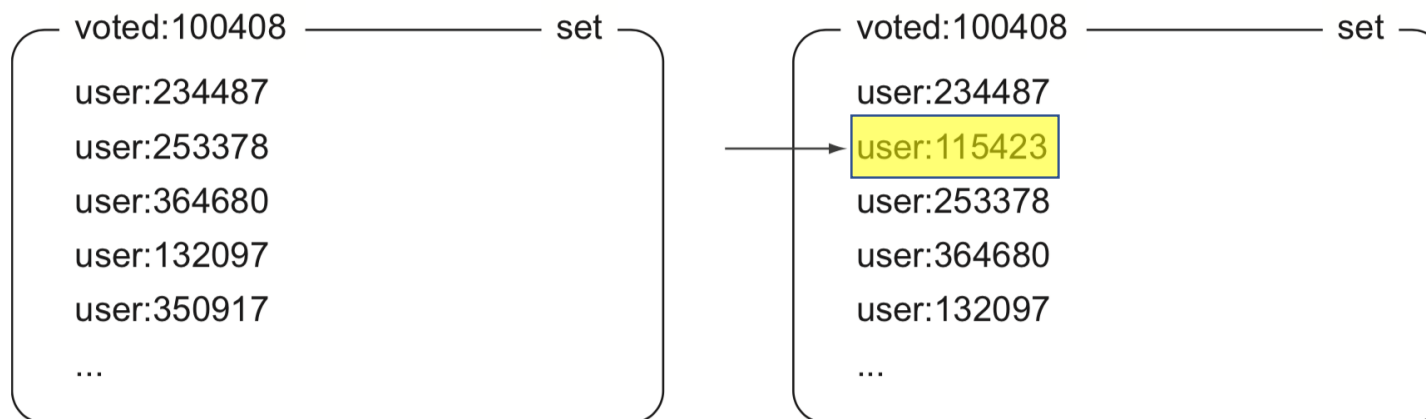
is article published within last week?
YES=add +1 to votes

article:92617 ————— hash	
title	Go to statement considered harmful
link	http://goo.gl/kZUSu
poster	user:83271
time	1331382699.33
votes	528

user:115423 votes for *article:100408*



Article 100408 got a new vote, so its score was increased.



Since user 115423 voted on the article, they are added to the voted SET.

Why new score is not +1?

*SCORE = number of seconds in a day (86.400) /
number of votes required (200) to last a full day*

SCORE = 432

User wants to vote for an article

- Verify that article was posted within last week in **article** HASH
 - **time:<article id>** must be less than one week old
- Add the user to article **voted** SET
 - **voted:<article id>**, add <user id>
- Increment the score of the article in **score** SET
 - **score:<article id>**, = score + SCORE INCREMENT (432)
- Update vote count in **article** HASH
 - **article:<article id>**, votes = votes + 1

check if article is less than a week old

- get article creation time from `<time:myarticle>` ZSET
- get current time with `time.time()`
- creation time = `ZSCORE(<time:myarticle>)`
- IS (current time – creation time) < one week in seconds?
 - YES: article can receive votes
- Now we must check if user has not voted for this article before

check if user has voted for myarticle

- value=`SISMEMBER voted:<myarticle> user`
 - if value = 1 user has already voted
 - If value = 0
 - add user to voters list
 - `SADD(voted:<myarticle>, user)`
 - update article score
 - `ZINCRBY(score:<myarticle>, 432)`
 - update vote count
 - `HINCRBY(article:<myarticle>, votes, 1)`

article vote implementation

```
ONE_WEEK_IN_SECONDS = 7 * 86400
```

```
VOTE_SCORE = 432
```

```
ctime=round(time.time())
```

```
cutoff = ctime - ONE_WEEK_IN_SECONDS
```

```
if conn.zscore('time:', article) > cutoff:
```

```
    article_id = article.partition(':')[ -1]
```

```
    conn.sadd('voted:' + article_id, user):
```

```
        conn.zincrby('score:', myarticle, VOTE_SCORE})
```

```
    conn.hincrby(myarticle, 'votes', 1)
```

Build TOP 10 articles by popularity

- Fetch current top scoring articles
- ZRANGE to fetch articles
- with highest decreasing score
- from position 0 to 9

```
ids = zrange("score:", 0, 9, REV)
```


Build TOP 10 articles by popularity

- Fetch most recent articles
- HGETALL to get info about each article

```
ids = xrange("score:", 0, 9, REV)
for id in ids:
    article_data = hgetall(id)
```

```
//GET TOP ARTICLES BY SCORE
```

```
ids = conn.zrange('score:', 0, top-1, 'REV')
```

```
articles = []
```

```
for id in ids:
```

```
    article_data = conn.hgetall(id)
```

```
    article_data['id'] = id
```

```
    articles.append(article_data)
```

```
return articles
```

Next tutorial

build a simple web service with Redis: structures + functions