


Neural Networks and Deep Learning

MLP & Backpropagation


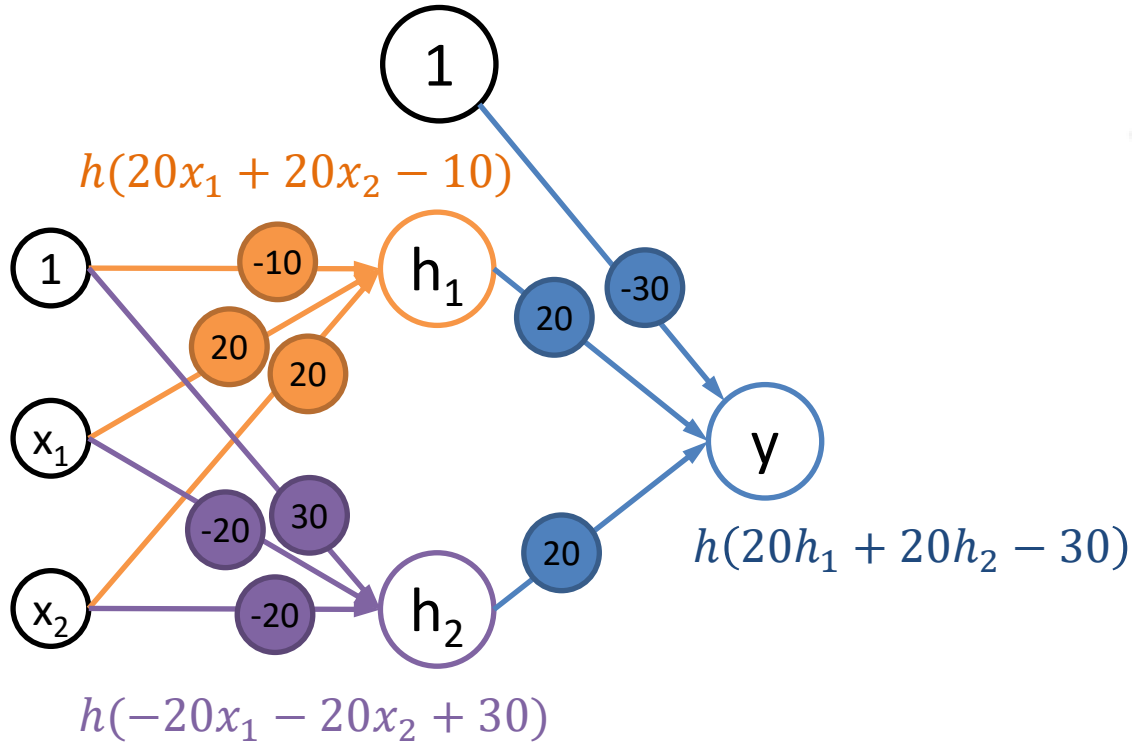
Credit to Dimosthenis Karatzas

BUILDING MORE COMPLEX NETWORKS

Deep Neural Network



Solving the XOR problem - revisited



OR

$$\sigma(20 * \mathbf{0} + 20 * \mathbf{0} - 10) = \mathbf{0}$$

$$\sigma(20 * \mathbf{1} + 20 * \mathbf{1} - 10) = \mathbf{1}$$

$$\sigma(20 * \mathbf{0} + 20 * \mathbf{1} - 10) = \mathbf{1}$$

$$\sigma(20 * \mathbf{1} + 20 * \mathbf{0} - 10) = \mathbf{1}$$

NAND

$$\sigma(-20 * \mathbf{0} - 20 * \mathbf{0} + 30) = \mathbf{1}$$

$$\sigma(-20 * \mathbf{1} - 20 * \mathbf{1} + 30) = \mathbf{0}$$

$$\sigma(-20 * \mathbf{0} - 20 * \mathbf{1} + 30) = \mathbf{1}$$

$$\sigma(-20 * \mathbf{1} - 20 * \mathbf{0} + 30) = \mathbf{1}$$

AND


$$\sigma(20 * \mathbf{0} + 20 * \mathbf{1} - 30) = \mathbf{0}$$

$$\sigma(20 * \mathbf{1} + 20 * \mathbf{0} - 30) = \mathbf{0}$$


$$\sigma(20 * \mathbf{1} + 20 * \mathbf{1} - 30) = \mathbf{1}$$

$$\sigma(20 * \mathbf{0} + 20 * \mathbf{0} - 30) = \mathbf{0}$$

What do hidden layers do?




An architecture with hidden units





An architecture with hidden units

Initial hidden layers would give you low-level information, adding subsequent hidden layers the system can encode higher-level features

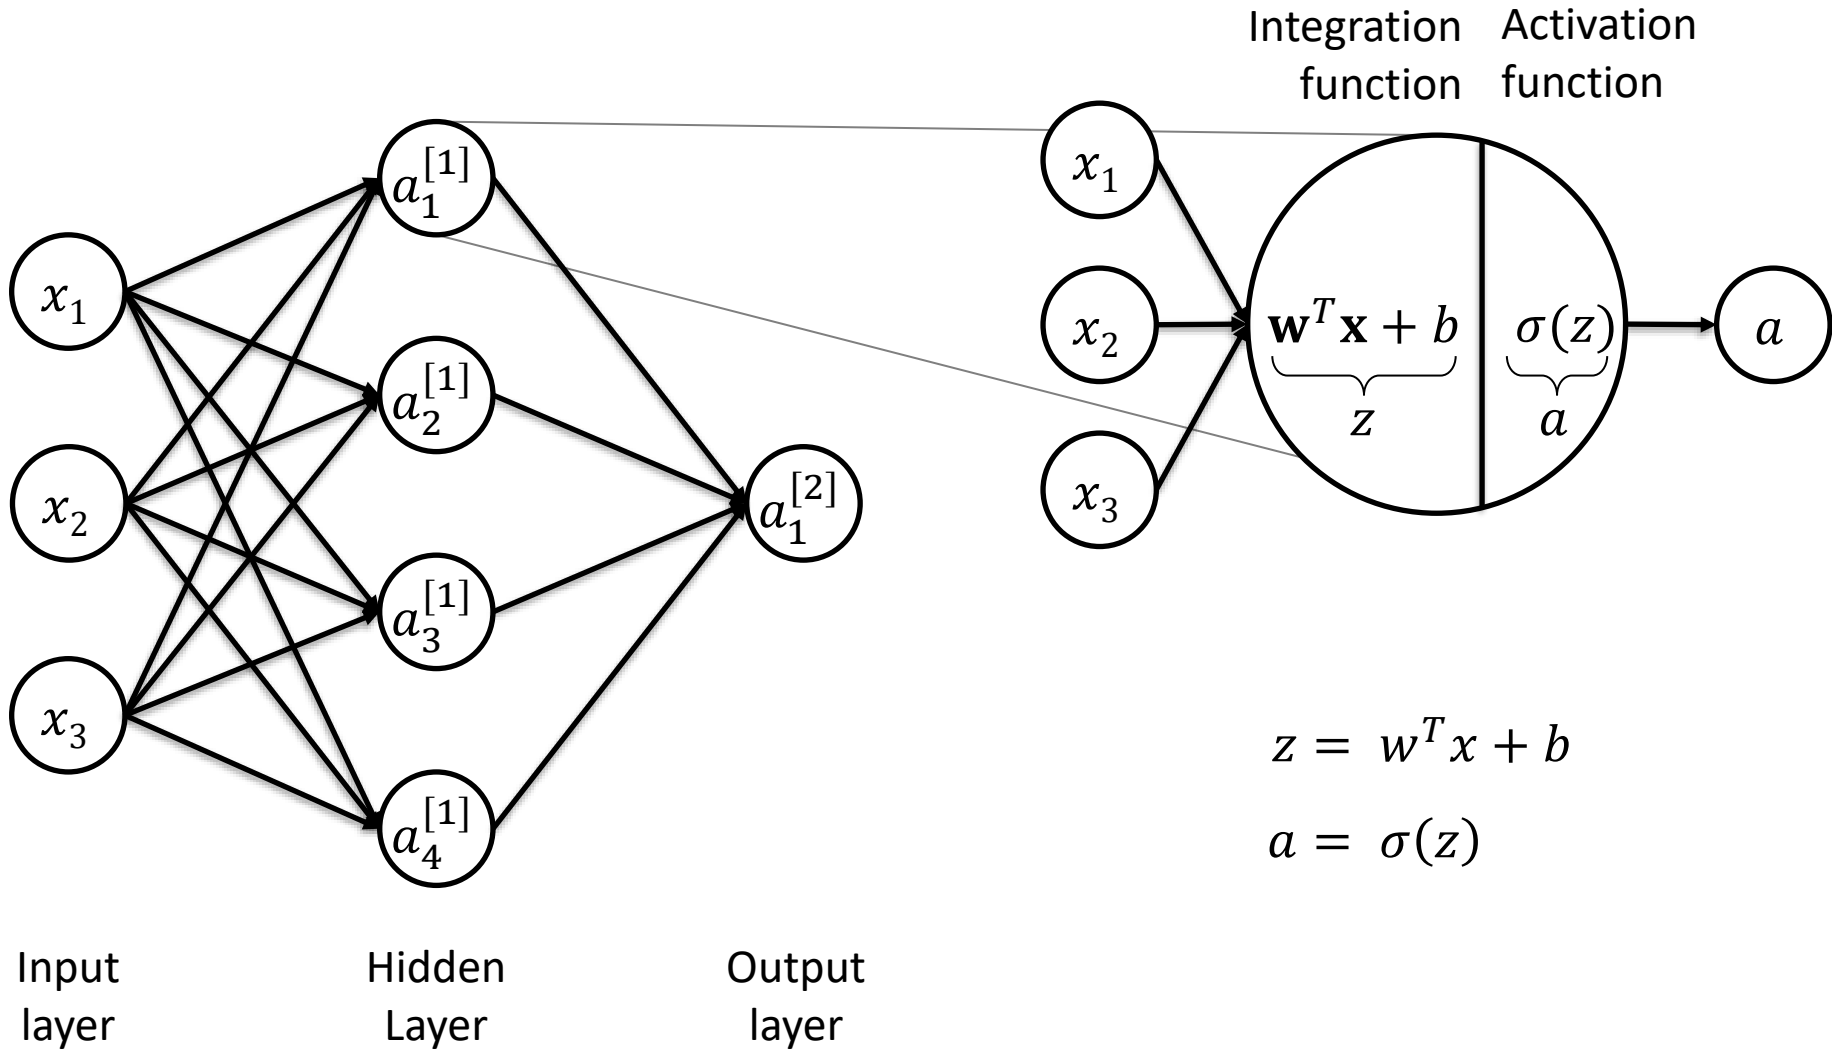


NEURAL NETWORKS NOTATION


Neural Network Notation



Neural Network Notation



Neural Network Notation




$$z_1^{[1]} = w_{1,1}^{[1]}x_1 + w_{1,2}^{[1]}x_2 + w_{1,3}^{[1]}x_3 + b_1^{[1]}$$

$$z_1^{[1]} = \mathbf{w}_1^{[1]T} \mathbf{x} + b_1^{[1]}$$

$$a_1^{[1]} = \sigma(z_1^{[1]})$$

Neural Network Notation




$$z_2^{[1]} = w_{2,1}^{[1]}x_1 + w_{2,2}^{[1]}x_2 + w_{2,3}^{[1]}x_3 + b_2^{[1]}$$

$$z_2^{[1]} = \mathbf{w}_2^{[1]T} \mathbf{x} + b_2^{[1]}$$

$$a_2^{[1]} = \sigma(z_2^{[1]})$$

Neural Network Notation




$$z_1^{[2]} = w_{1,1}^{[2]}a_1 + w_{1,2}^{[2]}a_2 + w_{1,3}^{[2]}a_3 + w_{1,4}^{[2]}a_4 + b_1^{[2]}$$

$$z_1^{[2]} = \mathbf{w}_1^{[2]T} \mathbf{a}^{[1]} + b_1^{[2]}$$

$$a_1^{[2]} = \sigma(z_1^{[2]})$$

Neural Network Notation




$$\begin{aligned} z_1^{[1]} &= \mathbf{w}_1^{[1]T} \mathbf{x} + b_1^{[1]}, & a_1^{[1]} &= \sigma(z_1^{[1]}) \\ z_2^{[1]} &= \mathbf{w}_2^{[1]T} \mathbf{x} + b_2^{[1]}, & a_2^{[1]} &= \sigma(z_2^{[1]}) \\ z_3^{[1]} &= \mathbf{w}_3^{[1]T} \mathbf{x} + b_3^{[1]}, & a_3^{[1]} &= \sigma(z_3^{[1]}) \\ z_4^{[1]} &= \mathbf{w}_4^{[1]T} \mathbf{x} + b_4^{[1]}, & a_4^{[1]} &= \sigma(z_4^{[1]}) \end{aligned}$$

$$\begin{bmatrix} z_1^{[1]} \\ z_2^{[1]} \\ z_3^{[1]} \\ z_4^{[1]} \end{bmatrix} = \begin{bmatrix} - & \mathbf{w}_1^{[1]T} & - \\ - & \mathbf{w}_2^{[1]T} & - \\ - & \mathbf{w}_3^{[1]T} & - \\ - & \mathbf{w}_4^{[1]T} & - \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} + \begin{bmatrix} b_1^{[1]} \\ b_2^{[1]} \\ b_3^{[1]} \\ b_4^{[1]} \end{bmatrix}$$

$$\begin{bmatrix} a_1^{[1]} \\ a_2^{[1]} \\ a_3^{[1]} \\ a_4^{[1]} \end{bmatrix} = \sigma \left(\begin{bmatrix} z_1^{[1]} \\ z_2^{[1]} \\ z_3^{[1]} \\ z_4^{[1]} \end{bmatrix} \right)$$

$$\mathbf{z}^{[1]} = \mathbf{W}^{[1]T} \mathbf{x} + \mathbf{b}^{[1]} \quad \mathbf{a}^{[1]} = \sigma(\mathbf{z}^{[1]})$$

Neural Network Notation



$$\begin{aligned} z_1^{[1]} &= \mathbf{w}_1^{[1]T} \mathbf{x} + b_1^{[1]}, & a_1^{[1]} &= \sigma(z_1^{[1]}) \\ z_2^{[1]} &= \mathbf{w}_2^{[1]T} \mathbf{x} + b_2^{[1]}, & a_2^{[1]} &= \sigma(z_2^{[1]}) \\ z_3^{[1]} &= \mathbf{w}_3^{[1]T} \mathbf{x} + b_3^{[1]}, & a_3^{[1]} &= \sigma(z_3^{[1]}) \\ z_4^{[1]} &= \mathbf{w}_4^{[1]T} \mathbf{x} + b_4^{[1]}, & a_4^{[1]} &= \sigma(z_4^{[1]}) \end{aligned}$$

Alternatively, in row notation:

$$\begin{bmatrix} z_1^{[1]} & z_2^{[1]} & z_3^{[1]} & z_4^{[1]} \end{bmatrix} = \begin{bmatrix} x_1 & x_2 & x_3 \end{bmatrix} \begin{bmatrix} \mathbf{w}_1^{[1]} \\ \mathbf{w}_2^{[1]} \\ \mathbf{w}_3^{[1]} \\ \mathbf{w}_4^{[1]} \end{bmatrix} + \begin{bmatrix} b_1^{[1]} & b_2^{[1]} & b_3^{[1]} & b_4^{[1]} \end{bmatrix}$$

$$\begin{bmatrix} a_1^{[1]} & a_2^{[1]} & a_3^{[1]} & a_4^{[1]} \end{bmatrix} = \sigma(\begin{bmatrix} z_1^{[1]} & z_2^{[1]} & z_3^{[1]} & z_4^{[1]} \end{bmatrix})$$

$$\mathbf{z}^{[1]T} = \mathbf{x}^T \mathbf{W}^{[1]} + \mathbf{b}^{[1]T} \quad \mathbf{a}^{[1]T} = \sigma(\mathbf{z}^{[1]T})$$

Neural Network Notation

With multiple points, for layer [1]:

Sample number

Broadcasted

$$\begin{matrix} (m \times 4) \\ \left[\begin{array}{ccc|c} - & \mathbf{z}^{(1)} & - & \\ \vdots & \vdots & \vdots & \\ - & \mathbf{z}^{(i)} & - & \\ \vdots & \vdots & \vdots & \\ - & \mathbf{z}^{(m)} & - & \end{array} \right]^{[1]} \end{matrix} = \begin{matrix} (m \times 3) \\ \left[\begin{array}{ccc|c} - & \mathbf{x}^{(1)} & - & \\ \vdots & \vdots & \vdots & \\ - & \mathbf{x}^{(i)} & - & \\ \vdots & \vdots & \vdots & \\ - & \mathbf{x}^{(m)} & - & \end{array} \right] \end{matrix} \begin{matrix} (3 \times 4) \\ \left[\begin{array}{cccc} | & | & | & | \\ \mathbf{w}_1 & \mathbf{w}_2 & \mathbf{w}_3 & \mathbf{w}_4 \\ | & | & | & | \end{array} \right]^{[1]} \end{matrix} + \begin{matrix} (m \times 4) \\ \left[\begin{array}{ccc|c} - & \mathbf{b} & - & \\ \vdots & \vdots & \vdots & \\ - & \mathbf{b} & - & \\ \vdots & \vdots & \vdots & \\ - & \mathbf{b} & - & \end{array} \right]^{[1]} \end{matrix}$$


$$\begin{matrix} (m \times 4) \\ \left[\begin{array}{ccc|c} - & \boldsymbol{\alpha}^{(1)} & - & \\ \vdots & \vdots & \vdots & \\ - & \boldsymbol{\alpha}^{(i)} & - & \\ \vdots & \vdots & \vdots & \\ - & \boldsymbol{\alpha}^{(m)} & - & \end{array} \right]^{[1]} \end{matrix} = \sigma \left(\begin{matrix} (m \times 4) \\ \left[\begin{array}{ccc|c} - & \mathbf{z}^{(1)} & - & \\ \vdots & \vdots & \vdots & \\ - & \mathbf{z}^{(i)} & - & \\ \vdots & \vdots & \vdots & \\ - & \mathbf{z}^{(m)} & - & \end{array} \right]^{[1]} \end{matrix} \right)$$

Design Matrix

(#samples × #features)

$$\mathbf{Z}^{[1]} = \mathbf{X} \mathbf{W}^{[1]} + \mathbf{B}^{[1]}$$

$$\mathbf{A}^{[1]} = \sigma(\mathbf{Z}^{[1]})$$



Neural Network Notation

In general, to go from layer $[L - 1]$ of k units to layer $[L]$ of n units, for a batch of m samples

$$\begin{array}{c} (m \times n) \\ \left[\begin{array}{ccc} - & \mathbf{z}^{(1)} & - \\ \vdots & \vdots & \vdots \\ - & \mathbf{z}^{(i)} & - \\ \vdots & \vdots & \vdots \\ - & \mathbf{z}^{(m)} & - \end{array} \right]^{[L]} \end{array} = \begin{array}{c} (m \times k) \\ \left[\begin{array}{ccc} - & \boldsymbol{\alpha}^{(1)} & - \\ \vdots & \vdots & \vdots \\ - & \boldsymbol{\alpha}^{(i)} & - \\ \vdots & \vdots & \vdots \\ - & \boldsymbol{\alpha}^{(m)} & - \end{array} \right]^{[L-1]} \end{array} \begin{array}{c} (k \times n) \\ \left[\begin{array}{ccc} | & \cdots & | \\ \mathbf{w}_1 & \cdots & \mathbf{w}_n \\ | & \cdots & | \end{array} \right]^{[L]} \end{array} + \begin{array}{c} (m \times n) \\ \left[\begin{array}{ccc} - & \mathbf{b} & - \\ \vdots & \vdots & \vdots \\ - & \mathbf{b} & - \\ \vdots & \vdots & \vdots \\ - & \mathbf{b} & - \end{array} \right]^{[L]} \end{array}$$

$$\begin{array}{c} (m \times n) \\ \left[\begin{array}{ccc} - & \boldsymbol{\alpha}^{(1)} & - \\ \vdots & \vdots & \vdots \\ - & \boldsymbol{\alpha}^{(i)} & - \\ \vdots & \vdots & \vdots \\ - & \boldsymbol{\alpha}^{(m)} & - \end{array} \right]^{[L]} \end{array} = \sigma \left(\begin{array}{c} (m \times n) \\ \left[\begin{array}{ccc} - & \mathbf{z}^{(1)} & - \\ \vdots & \vdots & \vdots \\ - & \mathbf{z}^{(i)} & - \\ \vdots & \vdots & \vdots \\ - & \mathbf{z}^{(m)} & - \end{array} \right]^{[L]} \end{array} \right)$$


$$\mathbf{Z}^{[L]} = \mathbf{A}^{[L-1]} \mathbf{W}^{[L]} + \mathbf{B}^{[L]}$$

$$\mathbf{A}^{[L]} = \sigma(\mathbf{Z}^{[L]})$$

Multi-layer neural networks

LEARNING WITH HIDDEN UNITS

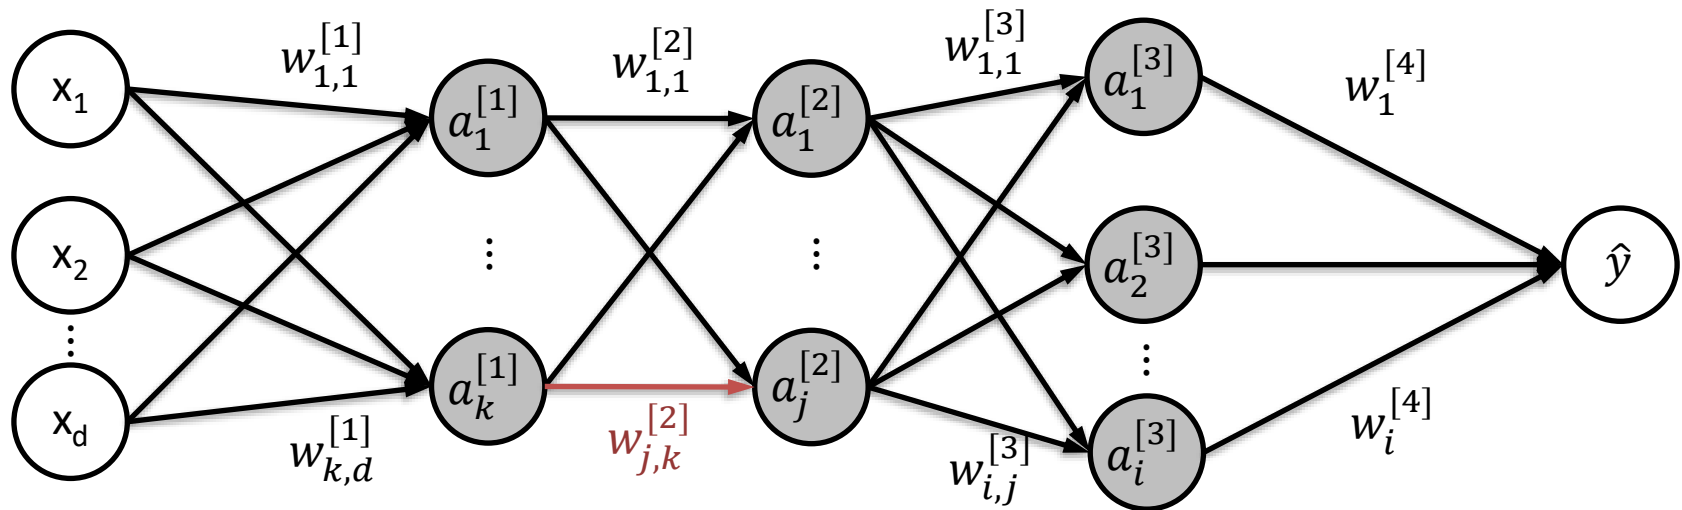
Backpropagation Algorithm



1. Receive new observation $\mathbf{x} = [x_1, x_2, \dots, x_d]$ and target output y
2. Feed-forward: let the network calculate its predicted output \hat{y}
3. Get the prediction \hat{y} and calculate the error (loss) e.g. $E = \frac{1}{2}(\hat{y} - y)^2$
4. **Back-propagate error:** calculate how each of the weights contributed to this error... HOW?

Backpropagation Algorithm


How should I change $w_{j,k}^{[2]}$?



1. Receive new observation $\mathbf{x} = [x_1, x_2, \dots, x_d]$ and target output y
2. Feed-forward: let the network calculate its predicted output \hat{y}
3. Get the prediction \hat{y} and calculate the error (loss) e.g. $E = \frac{1}{2}(\hat{y} - y)^2$
4. **Back-propagate error**: calculate how each of the weights contributed to this error... HOW?

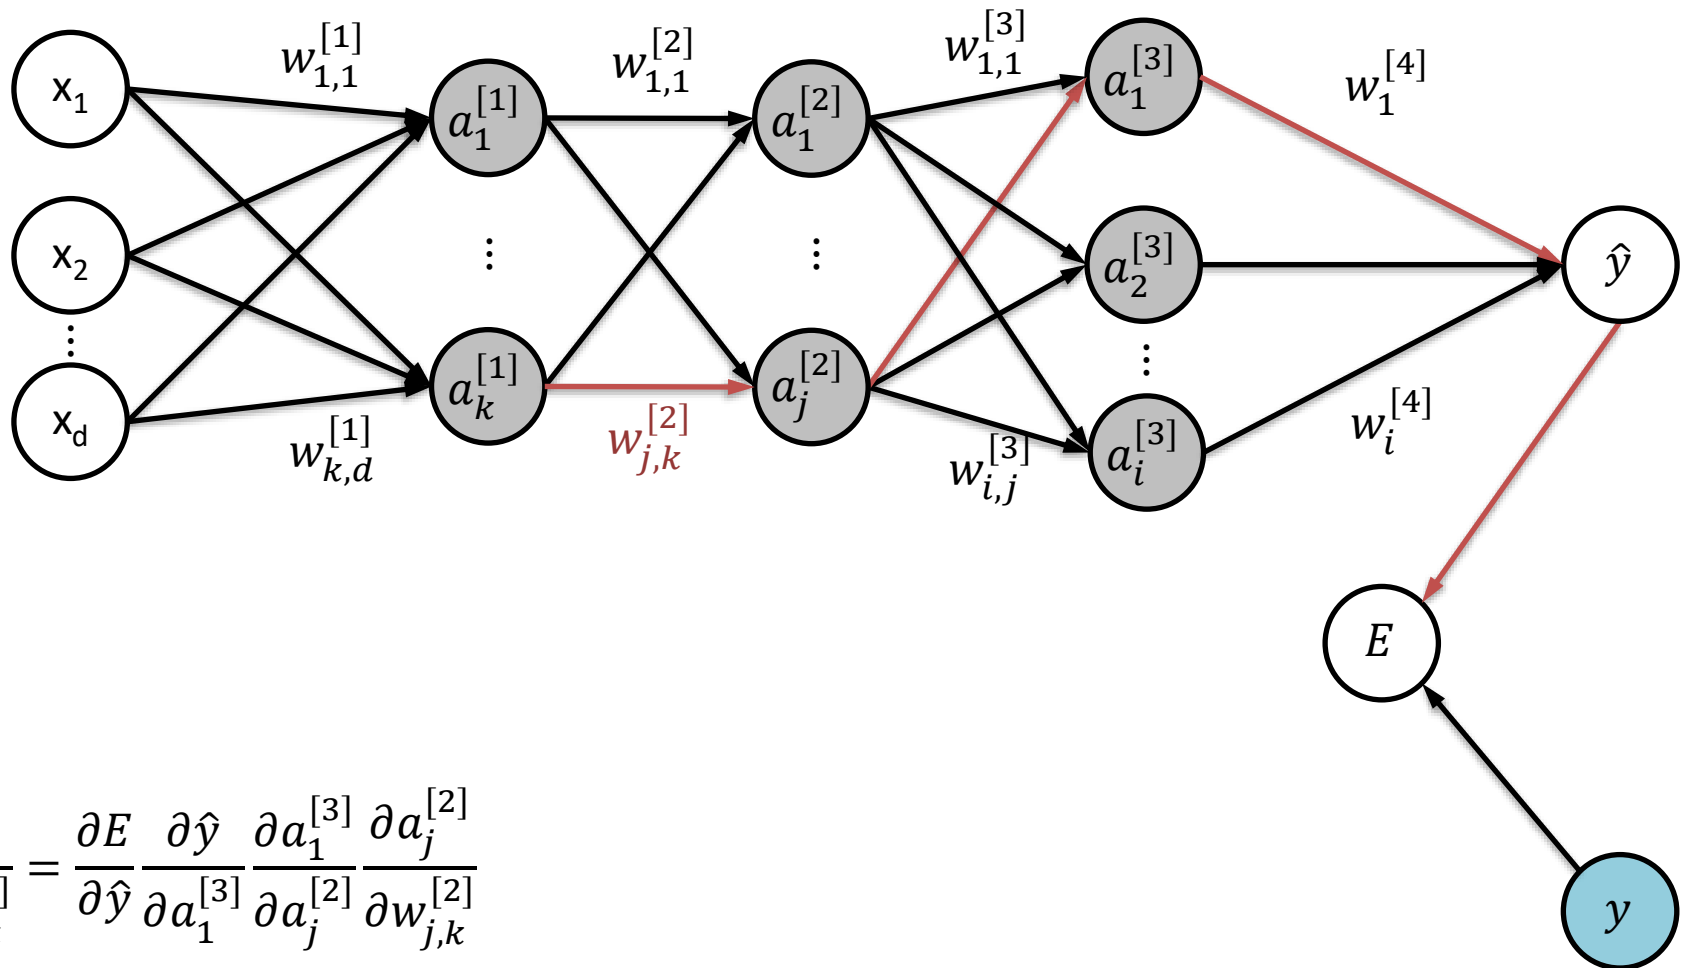
Backpropagation Algorithm

How should I change $w_{j,k}^{[2]}$?



Backpropagation Algorithm


How should I change $w_{j,k}^{[2]}$?



$$\frac{\partial E}{\partial w_{j,k}^{[2]}} = \frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial a_1^{[3]}} \frac{\partial a_1^{[3]}}{\partial a_j^{[2]}} \frac{\partial a_j^{[2]}}{\partial w_{j,k}^{[2]}}$$

Backpropagation Algorithm


How should I change $w_{j,k}^{[2]}$?



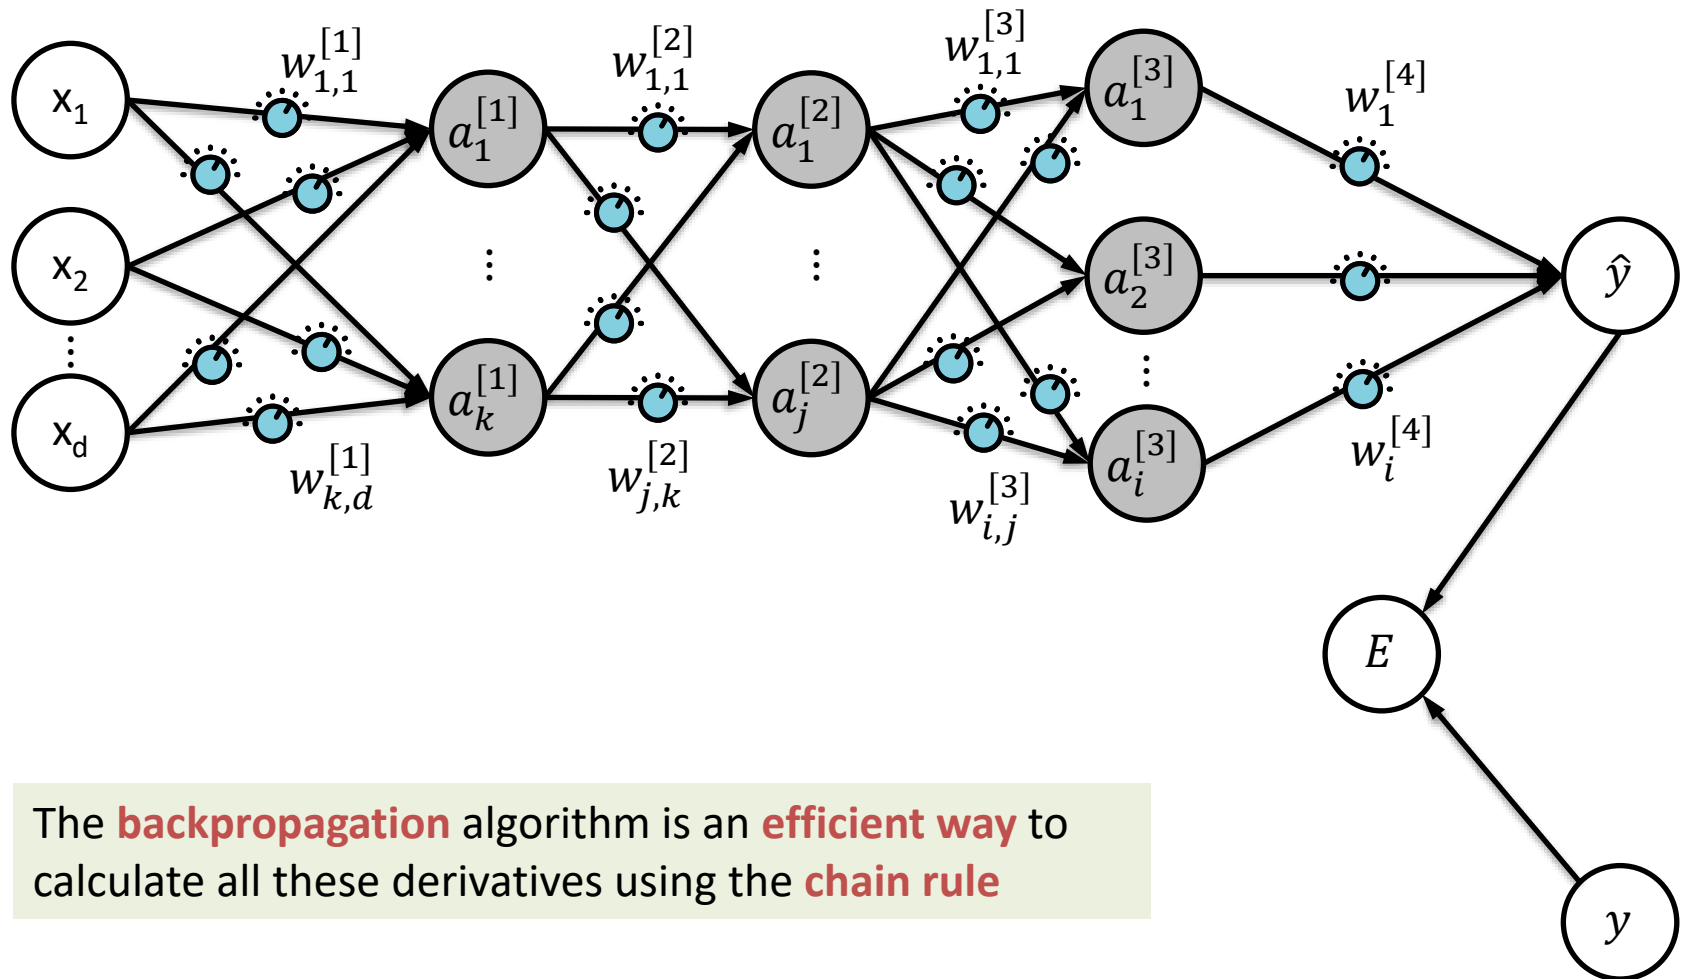
$$\frac{\partial E}{\partial w_{j,k}^{[2]}} = \frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial a_1^{[3]}} \frac{\partial a_1^{[3]}}{\partial a_j^{[2]}} \frac{\partial a_j^{[2]}}{\partial w_{j,k}^{[2]}} + \frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial a_i^{[3]}} \frac{\partial a_i^{[3]}}{\partial a_j^{[2]}} \frac{\partial a_j^{[2]}}{\partial w_{j,k}^{[2]}}$$

Backpropagation Algorithm

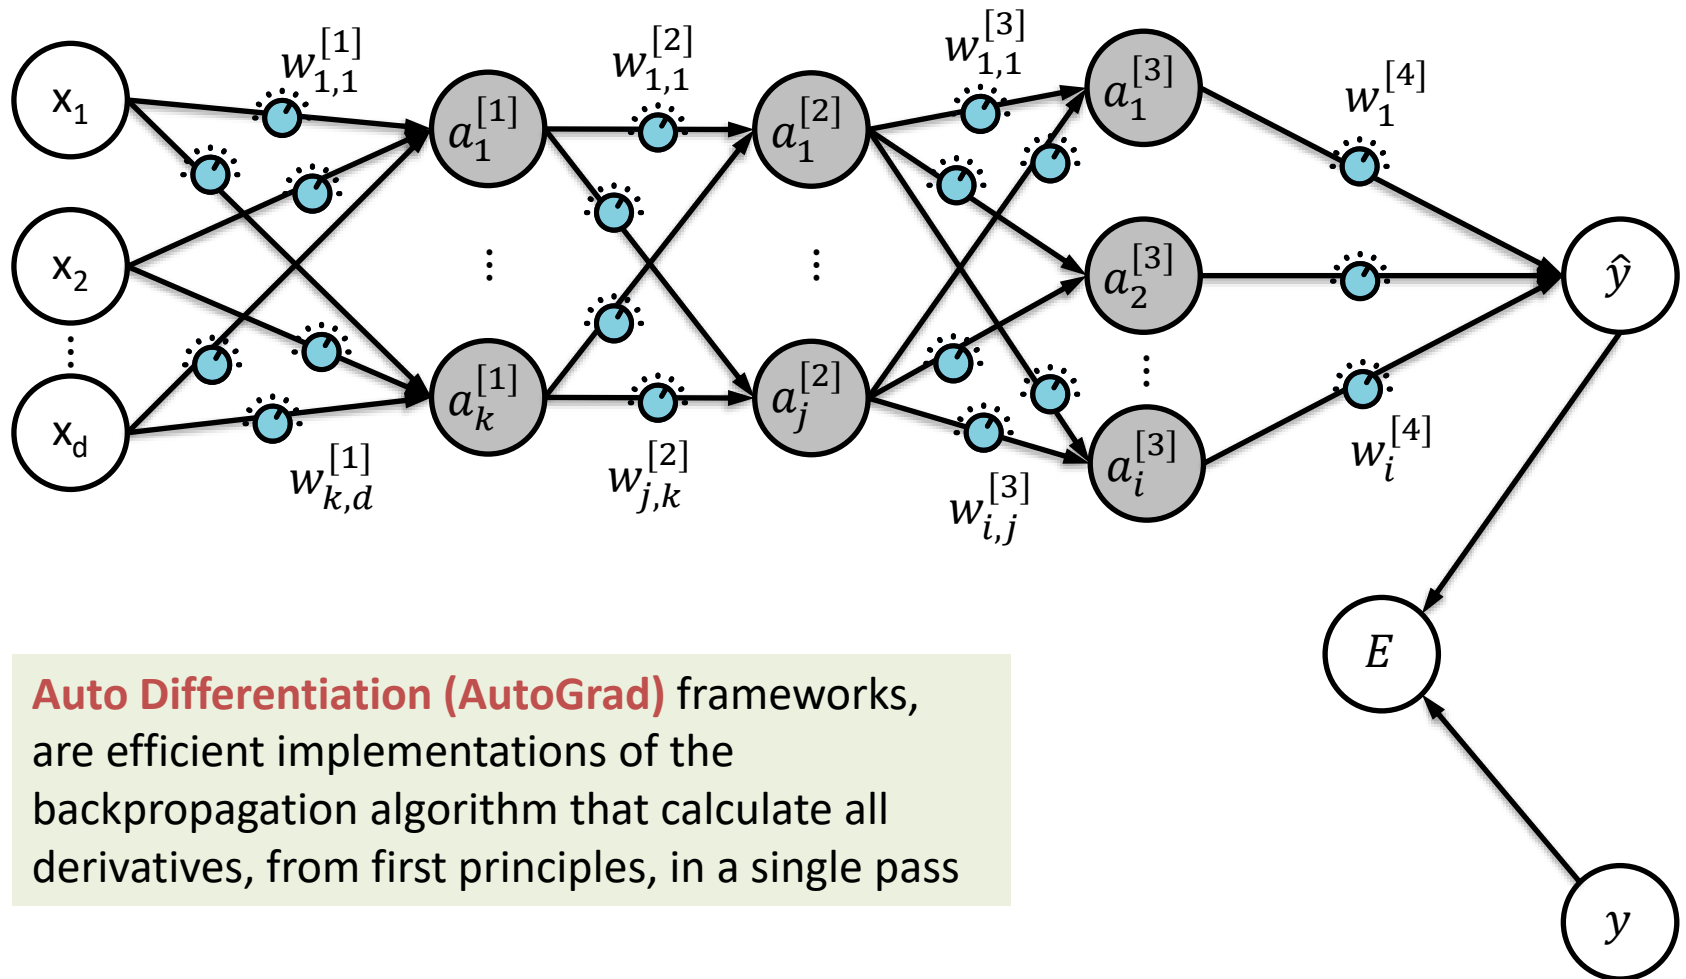
How should I change $w_{j,k}^{[2]}$?



Backpropagation Algorithm



Backpropagation Algorithm



Auto Differentiation (AutoGrad) frameworks, are efficient implementations of the backpropagation algorithm that calculate all derivatives, from first principles, in a single pass

Calculating Derivatives of Composite Functions

AUTO DIFFERENTIATION

Example

```
a = 4
b = 3
c = a + b # = 4 + 3 = 7
d = a * c # = 4 * 7 = 28
```

What is the derivative of d with respect to a : $\frac{\partial d}{\partial a}$?

$$d = a * c$$



Solving the traditional way

Example

a = 4
b = 3
c = a + b # = 4 + 3 = 7
d = a * c # = 4 * 7 = 28

What is the derivative of d with respect to a : $\frac{\partial d}{\partial a}$?

$$\begin{aligned}d &= a * c \\ \frac{\partial d}{\partial a} &= \frac{\partial a}{\partial a} * c + a * \frac{\partial c}{\partial a} \\ &= c + a * \frac{\partial c}{\partial a} \\ &= (a + b) + a * \frac{\partial(a + b)}{\partial a} \\ &= a + b + a * \left(\frac{\partial a}{\partial a} + \frac{\partial b}{\partial a} \right) \\ &= a + b + a * (1 + 0) \\ &= a + b + a = 2a + b \\ &= 2 * 4 + 3 = 11\end{aligned}$$


Solving the traditional way

Now, what is the derivative of d with respect to b : $\frac{\partial d}{\partial b}$?

You would have to carry out the whole process again...



Computational Graph

```
a = 4  
b = 3  
c = a + b # = 4 + 3 = 7  
d = a * c # = 4 * 7 = 28
```





Any expression can be broken down into a series of **simple operations** that are **applied sequentially**

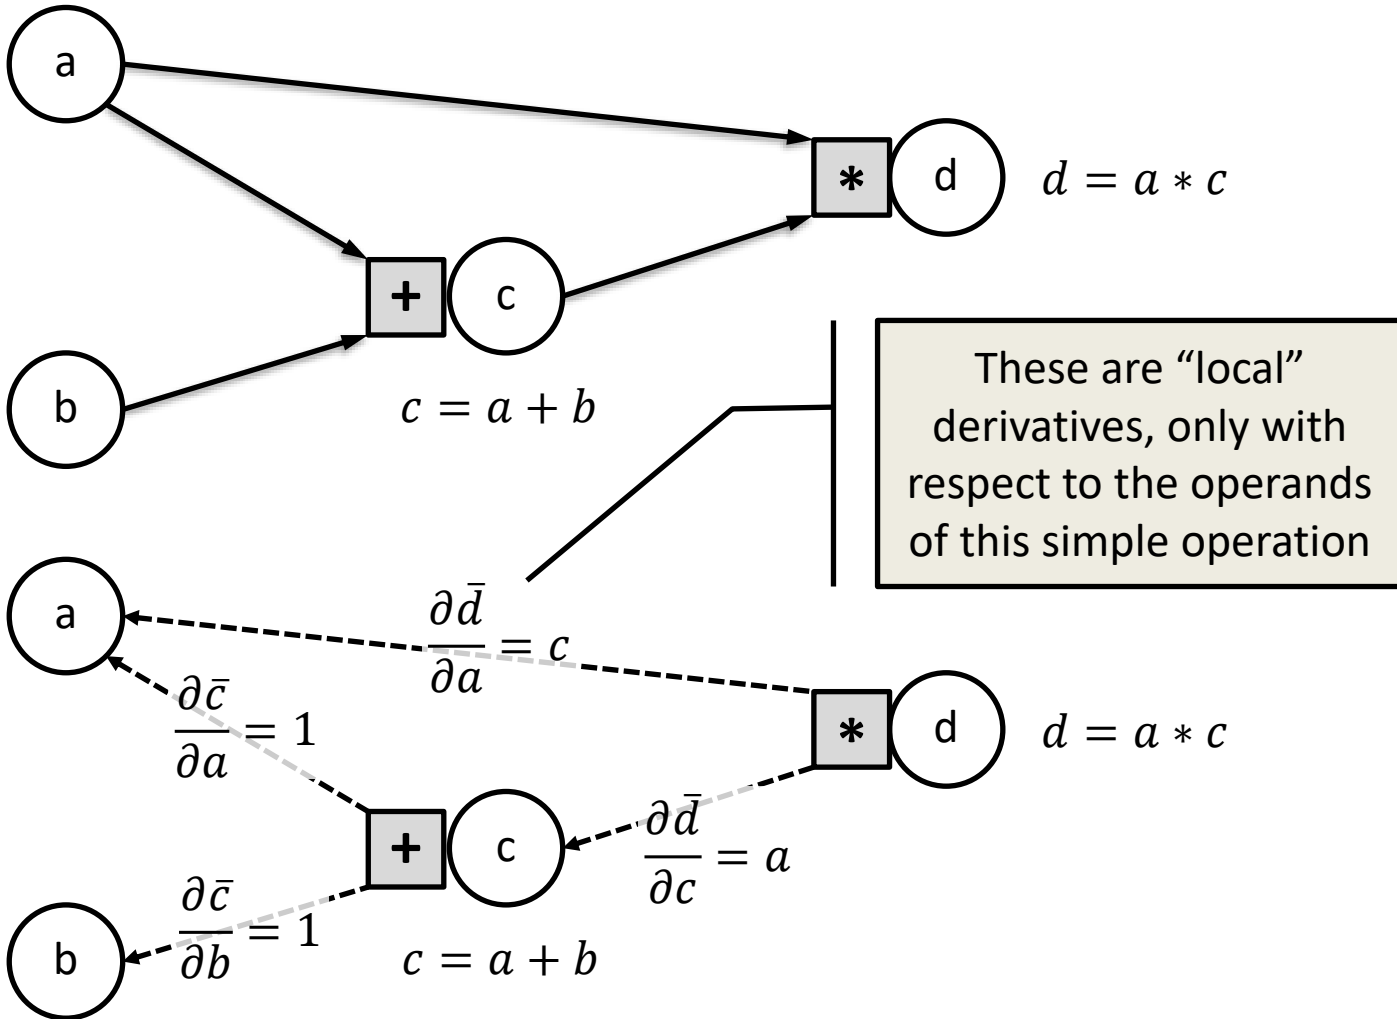
Local Derivatives



Local Derivatives




Local Derivatives



Automatic Differentiation (AutoGrad)



$$\frac{\partial d}{\partial a} = \frac{\partial \bar{d}}{\partial a} + \frac{\partial \bar{d}}{\partial c} * \frac{\partial \bar{c}}{\partial a}$$




Automatic Differentiation (AutoGrad)



$$\begin{aligned}\frac{\partial d}{\partial a} &= \frac{\partial \bar{d}}{\partial a} + \frac{\partial \bar{d}}{\partial c} * \frac{\partial \bar{c}}{\partial a} \\ &= c + a * 1 \\ &= a + b + a \\ &= 2a + b \\ &= 11\end{aligned}$$

To calculate **any** derivative using the computational graph:

- **Multiply** the edges of a route
- **Add** together the different routes that lead from the quantity to derive to the node of interest




Automatic Differentiation (AutoGrad)



$$\begin{aligned}
 \frac{\partial d}{\partial a} &= \frac{\partial \bar{d}}{\partial a} + \frac{\partial \bar{d}}{\partial c} * \frac{\partial \bar{c}}{\partial a} \\
 &= c + a * 1 \\
 &= a + b + a \\
 &= 2a + b \\
 &= 11
 \end{aligned}$$

To calculate **any** derivative using the computational graph:

- **Multiply** the edges of a route
- **Add** together the different routes that lead from the quantity to derive to the node of interest



$$d = a * c$$

$$\begin{aligned}
 \frac{\partial d}{\partial a} &= \frac{\partial a}{\partial a} * c + a * \frac{\partial c}{\partial a} \\
 &= c + a * \frac{\partial c}{\partial a} \\
 &= (a + b) + a * \frac{\partial (a + b)}{\partial a} \\
 &= a + b + a * \left(\frac{\partial a}{\partial a} + \frac{\partial b}{\partial a} \right) \\
 &= a + b + a * (1 + 0) \\
 &= a + b + a = 2a + b \\
 &= 2 * 4 + 3 = 11
 \end{aligned}$$

Solving the traditional way

Automatic Differentiation (AutoGrad)


Q1: What is the derivative of d with respect to b ?

Q2: What is the derivative of d with respect to c ?

Q3: What is the derivative of c with respect to a ?

Remember:

- **Multiply** the edges of a route
- **Add** together the different routes that lead to the node




The Backwards Pass

The backwards pass is the algorithmic implementation of the backpropagation process, that calculates the derivative of d with respect to each node in the graph in a single pass

Remember:

- **Multiply** the edges of a route
- **Add** together the different routes that lead to a node

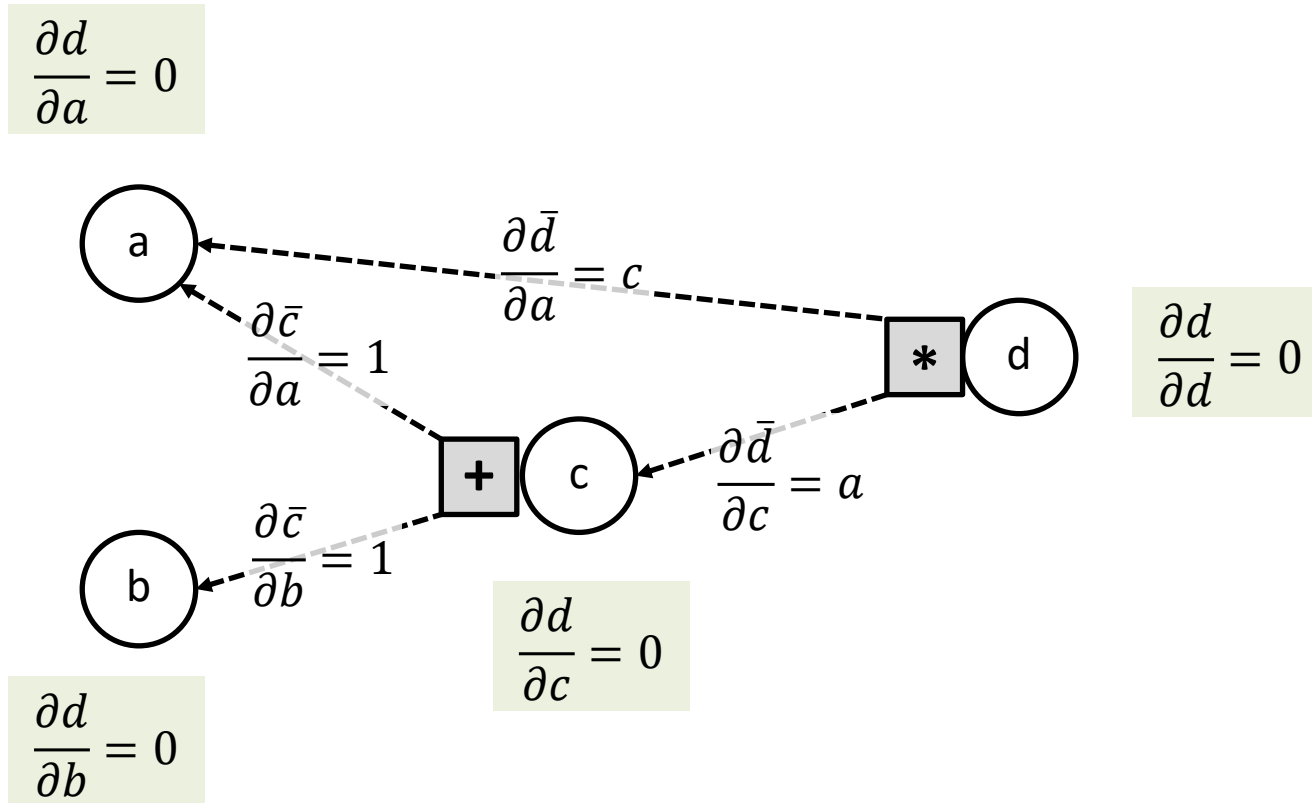


The Backwards Pass

The backwards pass is the algorithmic implementation of the backpropagation process, that calculates the derivative of d with respect to each node in the graph in a single pass

Remember:

- **Multiply** the edges of a route
- **Add** together the different routes that lead to a node

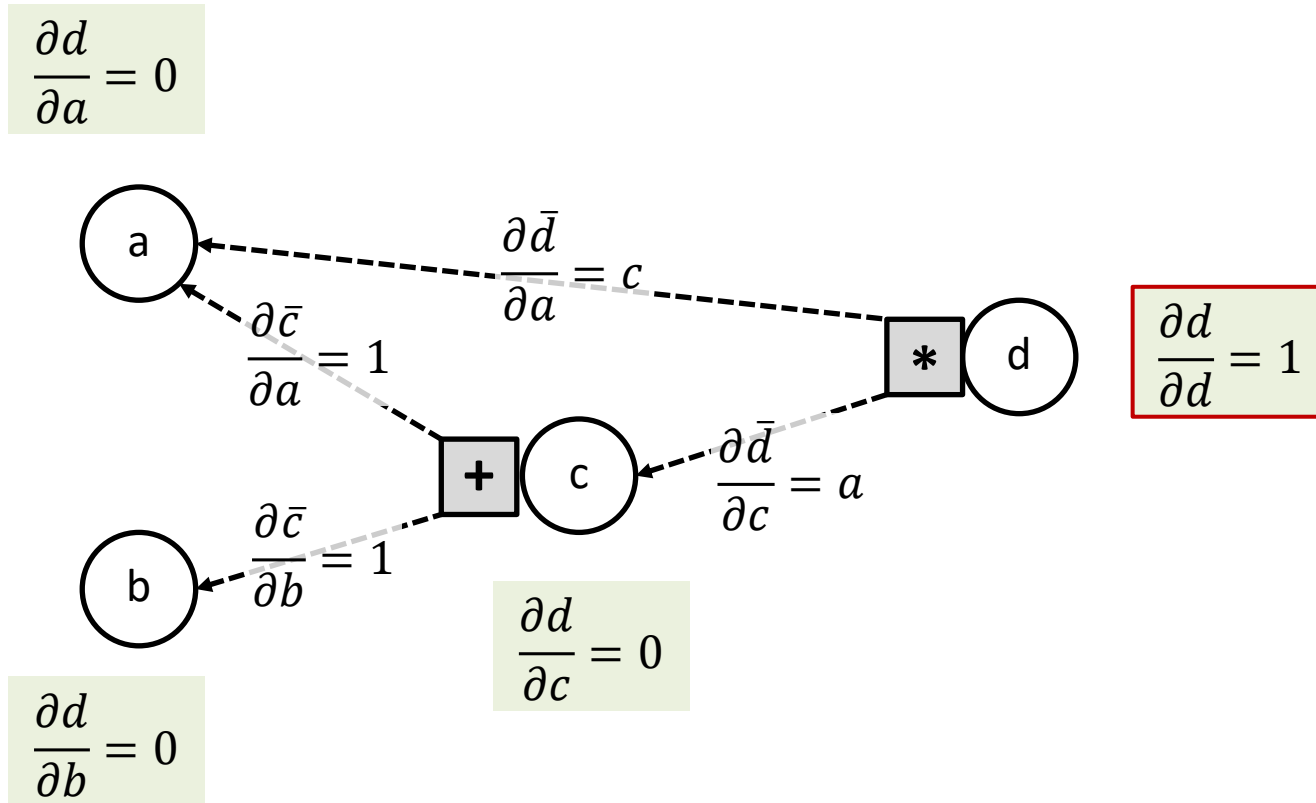


The Backwards Pass

The backwards pass is the algorithmic implementation of the backpropagation process, that calculates the derivative of d with respect to each node in the graph in a single pass

Remember:

- **Multiply** the edges of a route
- **Add** together the different routes that lead to a node




The Backwards Pass

The backwards pass is the algorithmic implementation of the backpropagation process, that calculates the derivative of d with respect to each node in the graph in a single pass

Remember:

- **Multiply** the edges of a route
- **Add** together the different routes that lead to a node

$$\frac{\partial d}{\partial a} = 0 + 1 * c$$




The Backwards Pass

The backwards pass is the algorithmic implementation of the backpropagation process, that calculates the derivative of d with respect to each node in the graph in a single pass

Remember:

- **Multiply** the edges of a route
- **Add** together the different routes that lead to a node

$$\frac{\partial d}{\partial a} = 0 + 1 * c$$



$$\frac{\partial d}{\partial b} = 0$$


The Backwards Pass

The backwards pass is the algorithmic implementation of the backpropagation process, that calculates the derivative of d with respect to each node in the graph in a single pass

Remember:

- **Multiply** the edges of a route
- **Add** together the different routes that lead to a node

$$\frac{\partial d}{\partial a} = 0 + 1 * c + (1 * a) * 1$$




The Backwards Pass

The backwards pass is the algorithmic implementation of the backpropagation process, that calculates the derivative of d with respect to each node in the graph in a single pass

Remember:

- **Multiply** the edges of a route
- **Add** together the different routes that lead to a node

$$\frac{\partial d}{\partial a} = 0 + 1 * c + (1 * a) * 1$$



$$\frac{\partial d}{\partial b} = 0 + (1 * a) * 1$$


The Backwards Pass

The backwards pass is the algorithmic implementation of the backpropagation process, that calculates the derivative of d with respect to each node in the graph in a single pass

Remember:

- **Multiply** the edges of a route
- **Add** together the different routes that lead to a node

$$\frac{\partial d}{\partial a} = c + a$$



$$\frac{\partial d}{\partial d} = 1$$



$$\frac{\partial d}{\partial c} = a$$

$$\frac{\partial d}{\partial b} = a$$

Using our AutoGrad framework

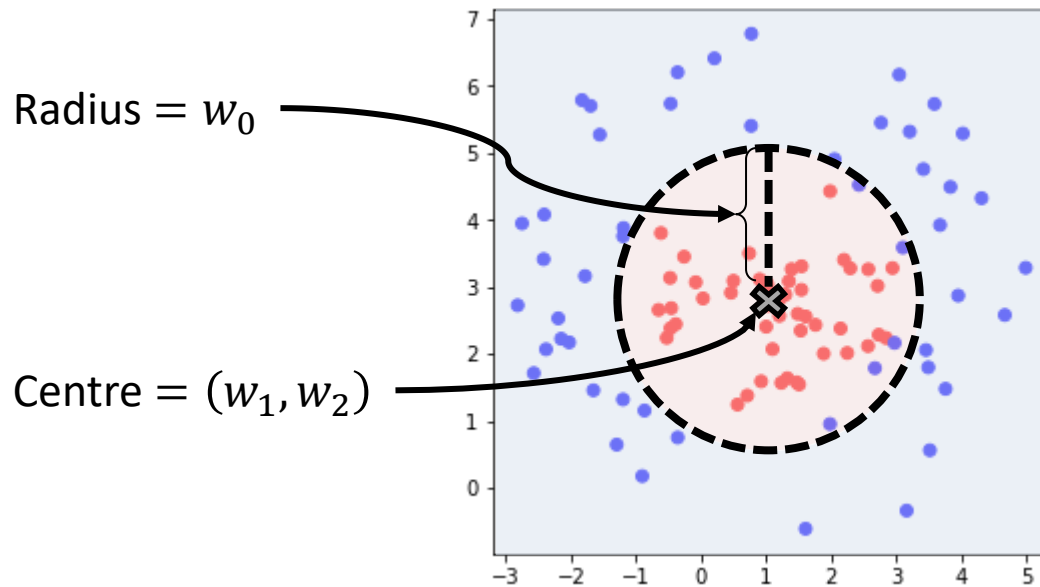
LEARNING THE PARAMETERS OF COMPOSITE FUNCTIONS

Non linear decision boundaries



$$(x_1, x_2) \rightarrow (x_1, x_2, x_3 = x_1^2 + x_2^2)$$

Example



We have some good intuition that we are looking for a closed decision boundary. We could try with a circle – but we have no prior knowledge of where the centre is, nor the radius. These are the parameters we are looking for.

$$z = (x_1 - w_1)^2 + (x_2 - w_2)^2 - w_0^2$$

Gradient Descent

Gradient descent works as usual, but in this case, you would have to calculate a complicated derivative, including the derivative of $\partial z / \partial w_i$

$$z = (x_1 - w_1)^2 + (x_2 - w_2)^2 - w_0^2$$

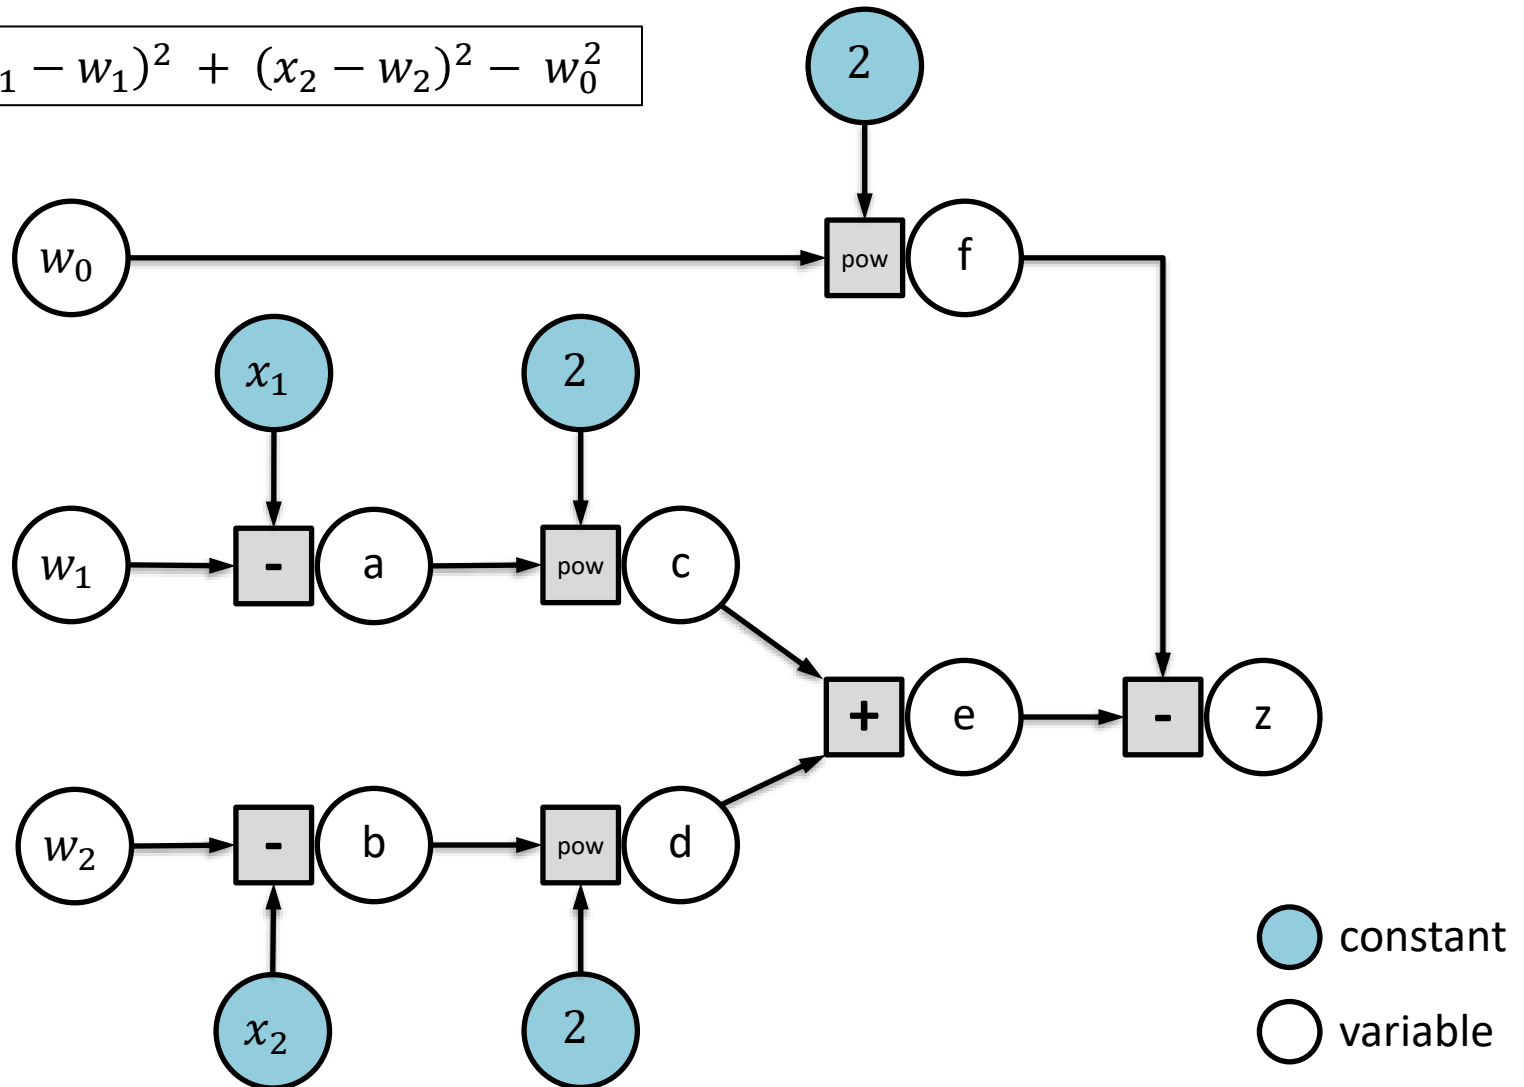
$$\frac{\partial z}{\partial w_0} = ? \quad \frac{\partial z}{\partial w_1} = ? \quad \frac{\partial z}{\partial w_2} = ?$$

AutoGrad can calculate all these derivatives for us. Then we can use normal gradient descent:


$$w_i \leftarrow w_i - \alpha \frac{\partial E}{\partial w_i}$$

Computational Graph

$$z = (x_1 - w_1)^2 + (x_2 - w_2)^2 - w_0^2$$




Activation Function (Sigmoid)




$$z = (x_1 - w_1)^2 + (x_2 - w_2)^2 - w_0^2$$

We would like wherever $z > 0.5$ to classify as **class 1**, and wherever $z < 0.5$ to classify as **class 0**. Hence, we apply a sigmoid on z .

$$\hat{y} = g(z) = \frac{1}{1 + e^{-z}}$$



Loss (binary cross-entropy loss)




$$z = (x_1 - w_1)^2 + (x_2 - w_2)^2 - w_0^2$$

$$\hat{y} = g(z) = \frac{1}{1 + e^{-z}}$$

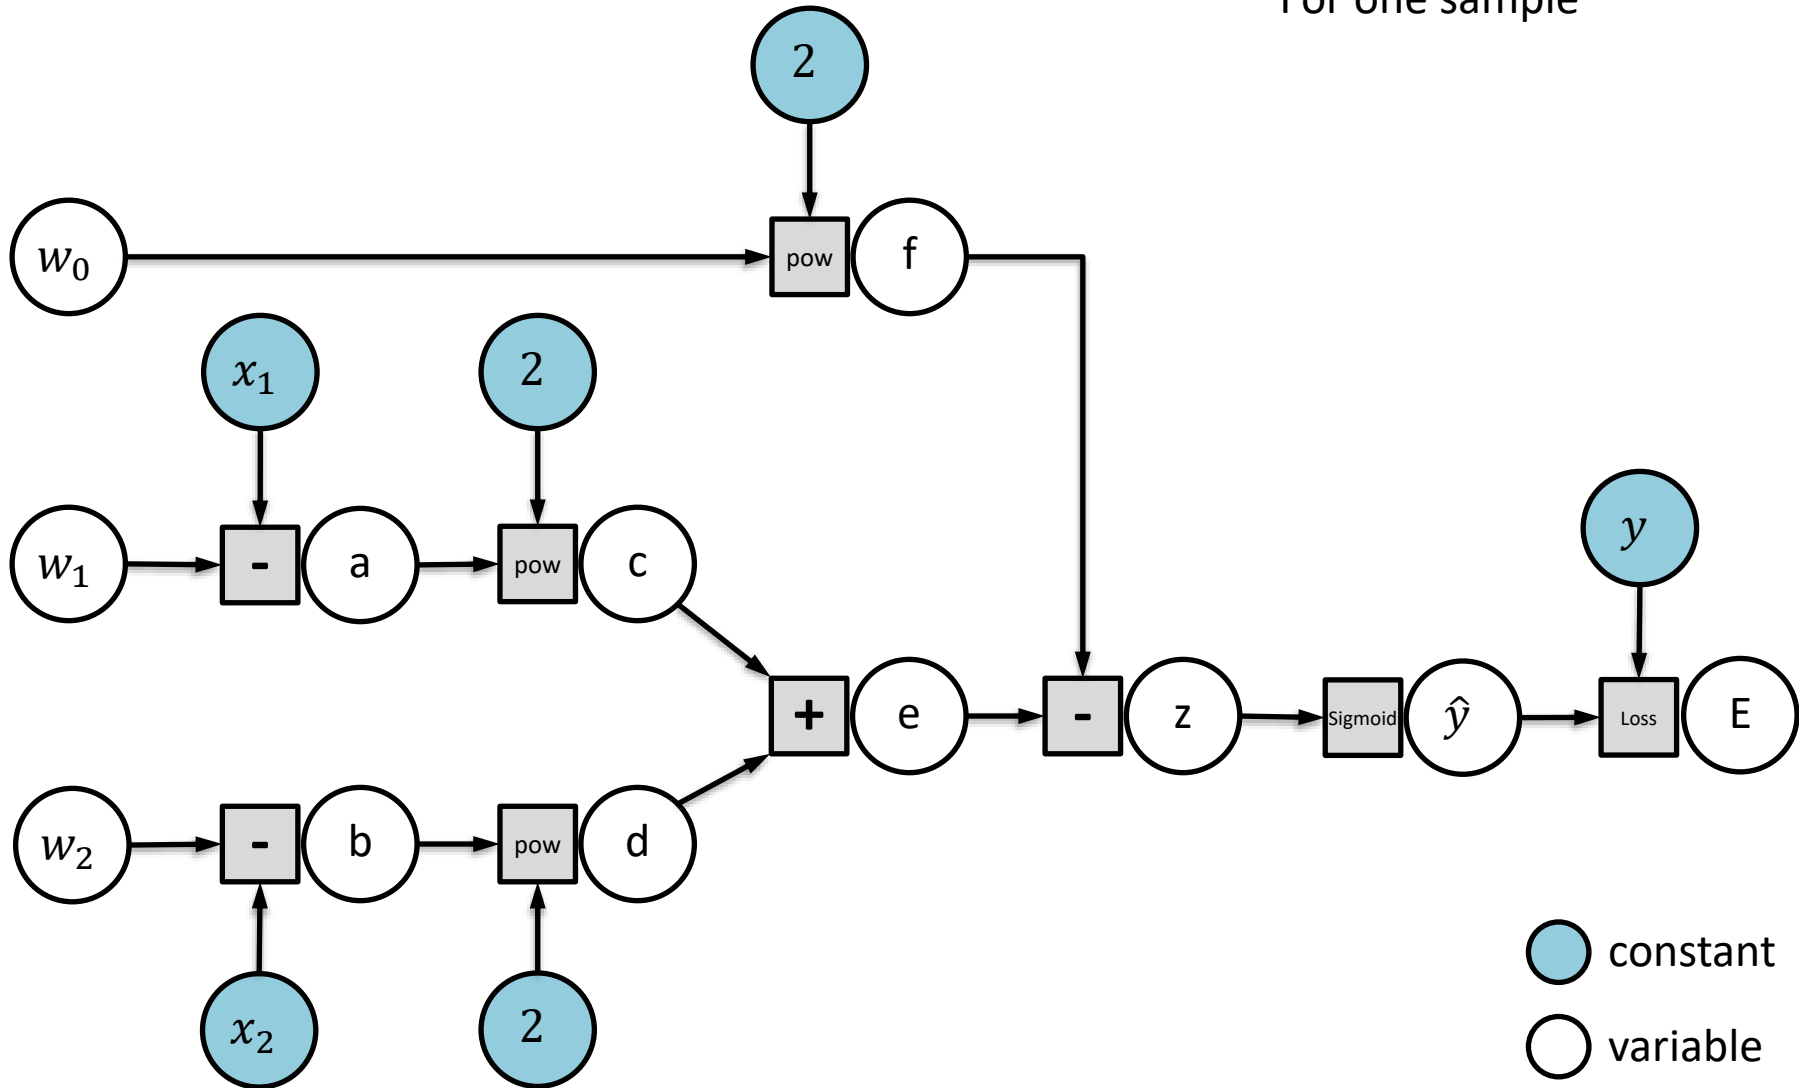
Finally, we would compare the output to the correct class using the cross-entropy loss we saw before:

$$Loss = -y \log(g(z)) - (1 - y) \log(1 - g(z))$$



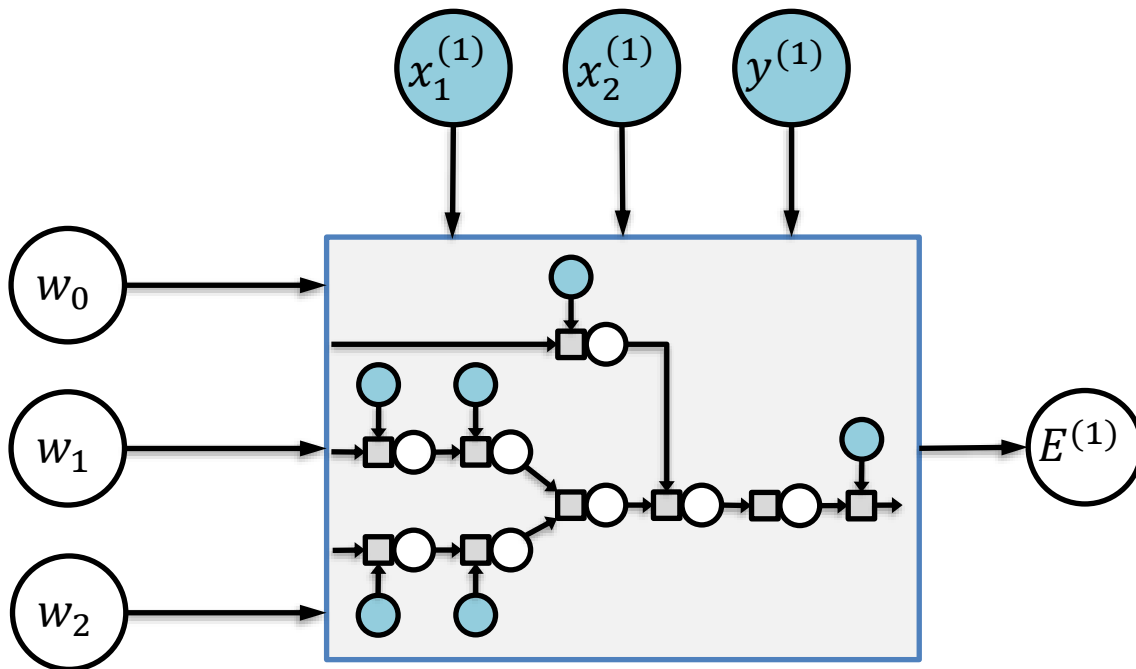
Complete Computational Graph

For one sample



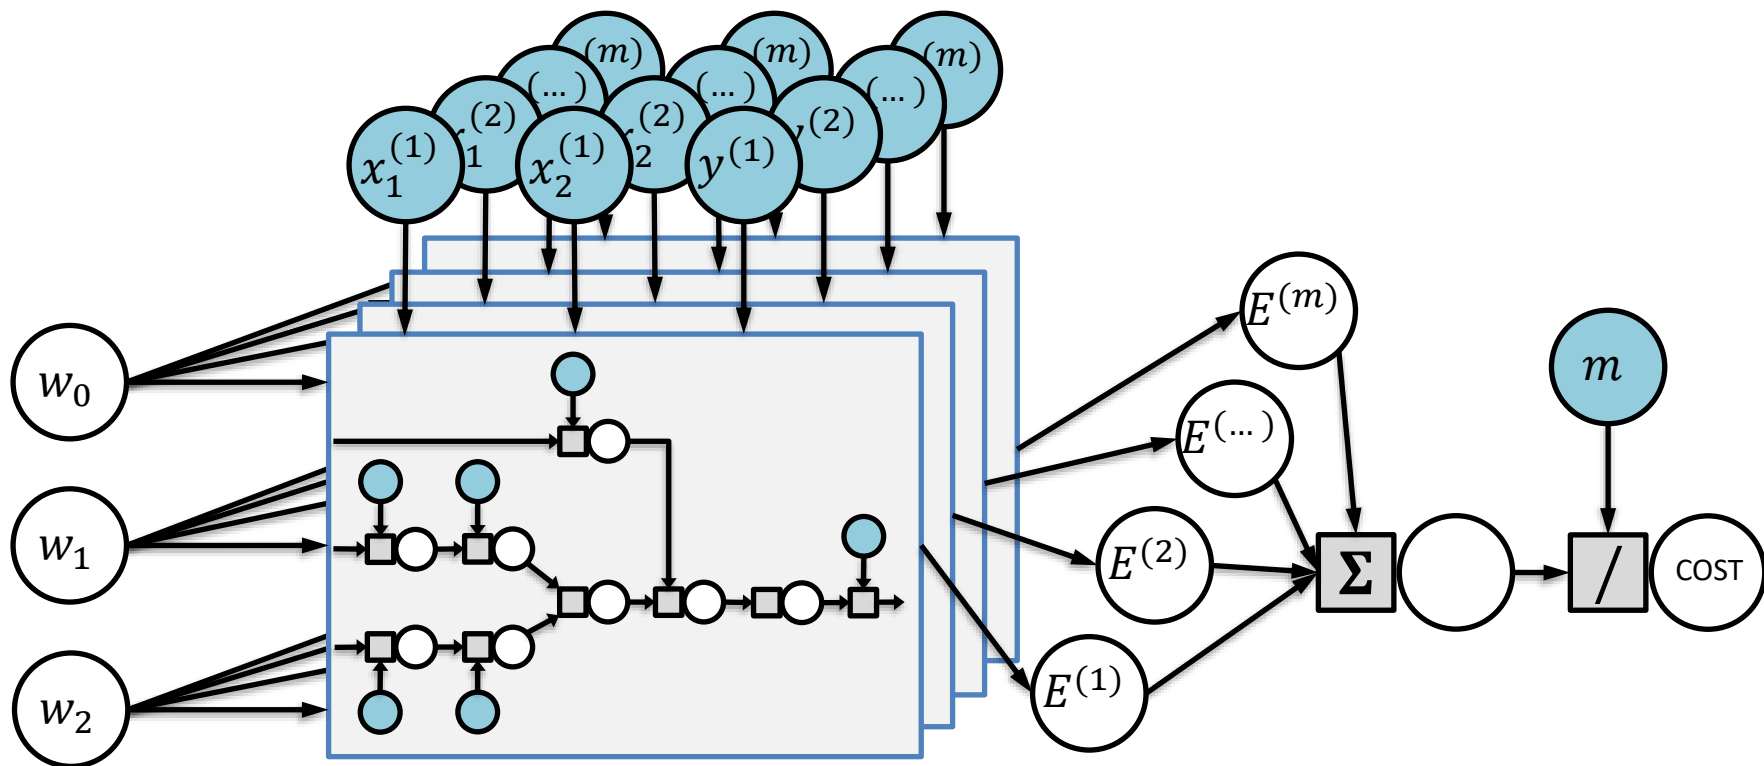
Complete Computational Graph

For one sample



Complete Computational Graph

Full batch of m samples



Gaining Efficiency

The derivative of the sum, equals the sum of derivatives... We can backpropagate errors for each sample individually, and accumulate the derivatives

Initialise gradients to zero


```
for epoch in range(1, 1000):  
    w0.zeroGradient()  
    w1.zeroGradient()  
    w2.zeroGradient()
```

The backpropagated error for every point is accumulated to the derivative of each weight

```
for x, y in trainingSamples:  
    # Do forward pass  
    # backpropagate error
```


Remember to divide by the number of samples when applying gradient descent

```
w0 = w0 + learningRate * w0.grad/m  
w1 = w1 + learningRate * w1.grad/m  
w2 = w2 + learningRate * w2.grad/m
```



Gaining Efficiency


Vectorise data, and take advantage of the SIMD (Single Instruction, Multiple Data) capabilities of CPUs and GPUs



BACKPROPAGATION – TAKE TWO

Backpropagation Algorithm


How should I change $w_{j,k}^{[2]}$?



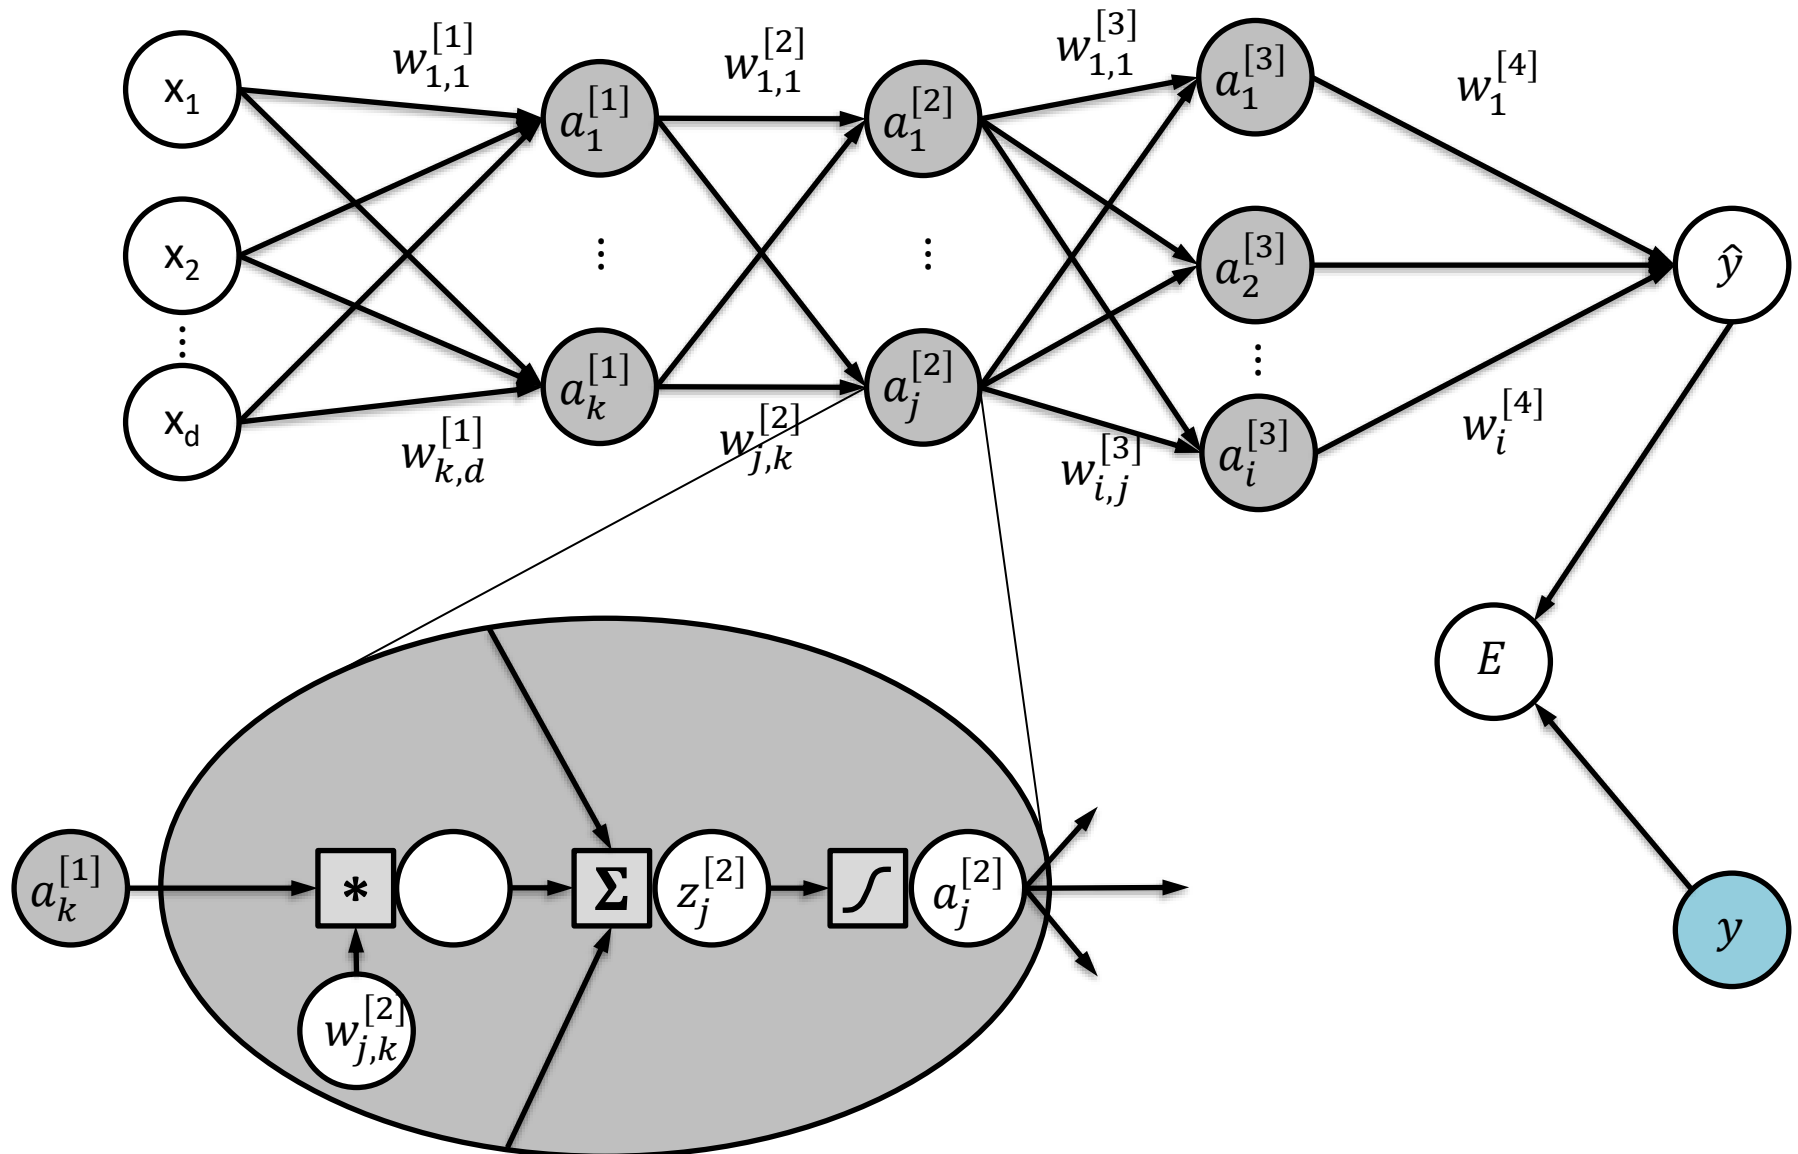
1. Receive new observation $\mathbf{x} = [x_1, x_2, \dots, x_d]$ and target output y
2. Feed-forward: let the network calculate its output \hat{y}
3. Get the prediction \hat{y} and calculate the error (loss) e.g. $E = \frac{1}{2}(\hat{y} - y)^2$
4. **Back-propagate error**: calculate how each of the weights contributed to this error

Backpropagation Algorithm

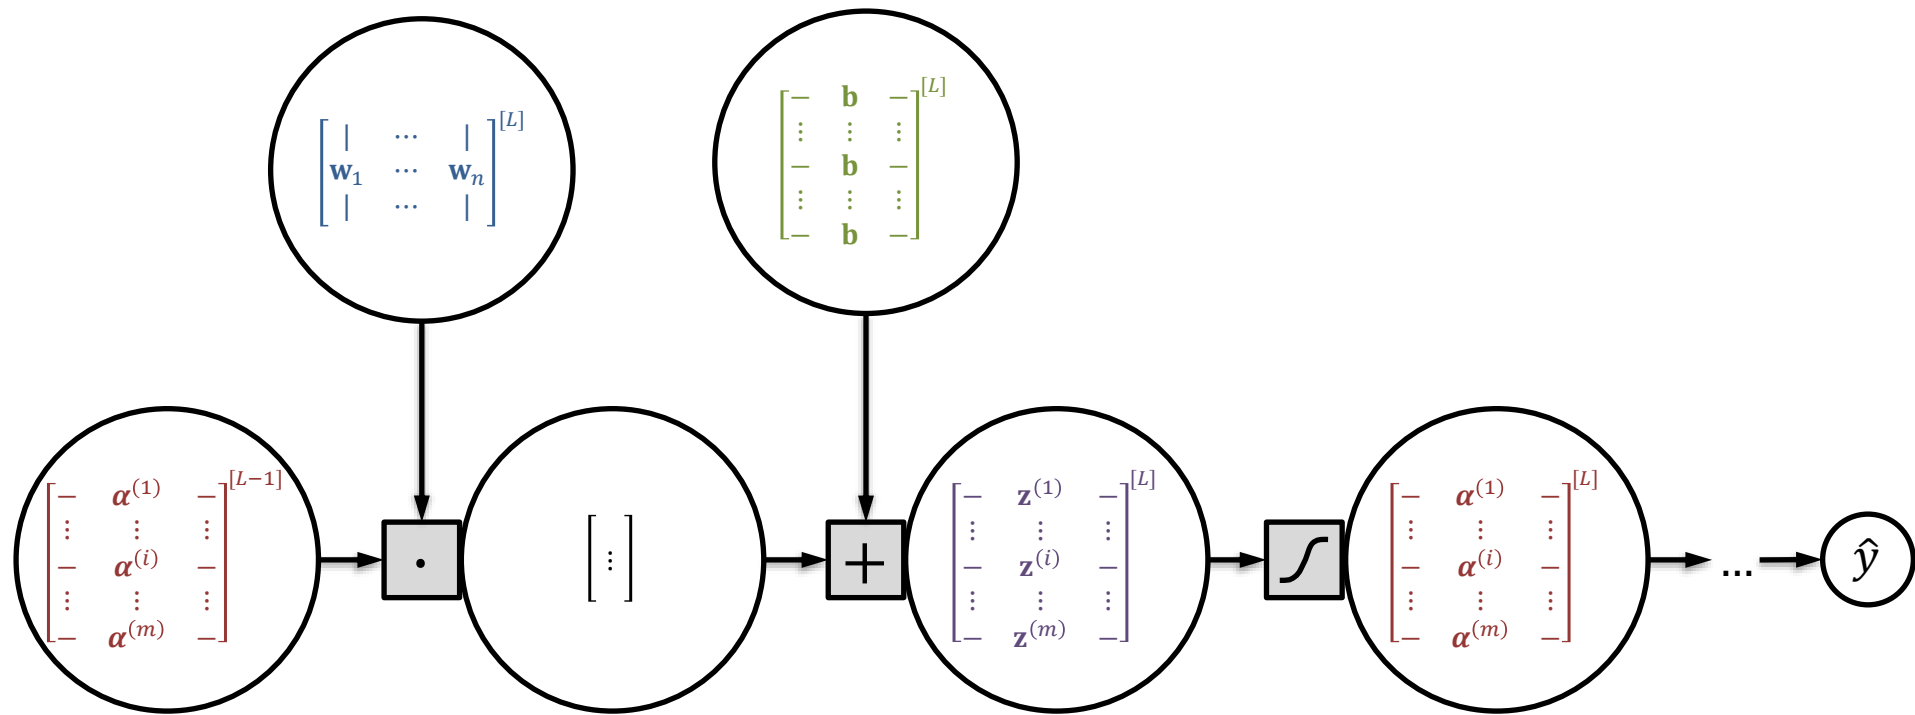
How should I change $w_{j,k}^{[2]}$?



Computation Graph of a NN



Computation Graph of a NN



In practice, all operations are vectorised and highly optimised to take advantage of the SIMD (Single Instruction, Multiple Data) capabilities of CPUs and GPUs

MATRIX DERIVATIVES

Derivatives with respect to a vector

Many times we need to calculate all the partial derivatives of a function whose input and output are both vectors.

For example, imagine the function $f: \mathbb{R}^3 \rightarrow \mathbb{R}^4$

$$\mathbf{y} = f(\mathbf{x}) = \mathbf{W}\mathbf{x}$$

$$\begin{matrix} (4 \times 1) & & (4 \times 3) & & (3 \times 1) \\ \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \end{bmatrix} & = & \begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} \\ w_{2,1} & w_{2,2} & w_{2,3} \\ w_{3,1} & w_{3,2} & w_{3,3} \\ w_{4,1} & w_{4,2} & w_{4,3} \end{bmatrix} & \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} \end{matrix}$$

$$\frac{d\mathbf{y}}{d\mathbf{x}} = ?$$

A full characterisation of the derivative of \mathbf{y} with respect to \mathbf{x} requires the partial derivative of **each component of \mathbf{y}** with respect to **each component of \mathbf{x}**

Derivatives with respect to a vector

A full characterisation of the derivative of \mathbf{y} with respect to \mathbf{x} requires the partial derivative of **each component of \mathbf{y}** with respect to **each component of \mathbf{x}**

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \end{bmatrix} = \begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} \\ w_{2,1} & w_{2,2} & w_{2,3} \\ w_{3,1} & w_{3,2} & w_{3,3} \\ w_{4,1} & w_{4,2} & w_{4,3} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$$

Let's compute one of these, e.g. the derivative of y_2 to x_3

$$y_2 = \sum_{j=1}^3 w_{2,j} x_j$$

$$y_2 = w_{2,1}x_1 + w_{2,2}x_2 + w_{2,3}x_3$$

$$\frac{\partial y_2}{\partial x_3} = \frac{\partial}{\partial x_3} [w_{2,1}x_1 + w_{2,2}x_2 + w_{2,3}x_3]$$

$$\frac{\partial y_2}{\partial x_3} = 0 + 0 + \frac{\partial}{\partial x_3} [w_{2,3}x_3]$$

$$\frac{\partial y_2}{\partial x_3} = w_{2,3}$$

In general: $\frac{\partial y_i}{\partial x_j} = w_{i,j}$

Jacobian Matrix

We can organise all these partial derivatives into a new matrix called the **Jacobian matrix**.

$$\mathbf{y} = \mathbf{W}\mathbf{x}$$

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \end{bmatrix} = \begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} \\ w_{2,1} & w_{2,2} & w_{2,3} \\ w_{3,1} & w_{3,2} & w_{3,3} \\ w_{4,1} & w_{4,2} & w_{4,3} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$$

In this case the Jacobian matrix would be:

$$J = \begin{bmatrix} \frac{\partial y_1}{\partial x_1} & \frac{\partial y_1}{\partial x_2} & \frac{\partial y_1}{\partial x_3} \\ \frac{\partial y_2}{\partial x_1} & \frac{\partial y_2}{\partial x_2} & \frac{\partial y_2}{\partial x_3} \\ \frac{\partial y_3}{\partial x_1} & \frac{\partial y_3}{\partial x_2} & \frac{\partial y_3}{\partial x_3} \\ \frac{\partial y_4}{\partial x_1} & \frac{\partial y_4}{\partial x_2} & \frac{\partial y_4}{\partial x_3} \end{bmatrix} = \begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} \\ w_{2,1} & w_{2,2} & w_{2,3} \\ w_{3,1} & w_{3,2} & w_{3,3} \\ w_{4,1} & w_{4,2} & w_{4,3} \end{bmatrix}$$

Note: The diagram shows a curved arrow from the general expression $\frac{\partial y_i}{\partial x_j} = w_{i,j}$ to the corresponding element in the Jacobian matrix.

$$\frac{d\mathbf{y}}{d\mathbf{x}} = \mathbf{J}_y(\mathbf{x}) = \mathbf{W}$$

Jacobian Matrix

In general, for every function $\mathbf{f}: \mathbb{R}^m \rightarrow \mathbb{R}^n$, the Jacobian matrix $\mathbf{J} \in \mathbb{R}^{n \times m}$ of \mathbf{f} is defined such that $J_{i,j} = \frac{\partial f(\mathbf{x})_i}{\partial x_j}$

For a function $\mathbf{f}: \mathbb{R}^3 \rightarrow \mathbb{R}^4$

The Jacobian matrix would be

$$\mathbf{J}_f(\mathbf{x}) = \begin{bmatrix} \frac{\partial f(\mathbf{x})_1}{\partial x_1} & \frac{\partial f(\mathbf{x})_1}{\partial x_2} & \frac{\partial f(\mathbf{x})_1}{\partial x_3} \\ \frac{\partial f(\mathbf{x})_2}{\partial x_1} & \frac{\partial f(\mathbf{x})_2}{\partial x_2} & \frac{\partial f(\mathbf{x})_2}{\partial x_3} \\ \frac{\partial f(\mathbf{x})_3}{\partial x_1} & \frac{\partial f(\mathbf{x})_3}{\partial x_2} & \frac{\partial f(\mathbf{x})_3}{\partial x_3} \\ \frac{\partial f(\mathbf{x})_4}{\partial x_1} & \frac{\partial f(\mathbf{x})_4}{\partial x_2} & \frac{\partial f(\mathbf{x})_4}{\partial x_3} \end{bmatrix}$$

Always be careful with the numerator and denominator layout notation when doing matrix calculus!

What about row vectors?

$$\cancel{y = f(\mathbf{x}) = \mathbf{W}\mathbf{x}}$$

$$\cancel{y = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \end{bmatrix} = \begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} \\ w_{2,1} & w_{2,2} & w_{2,3} \\ w_{3,1} & w_{3,2} & w_{3,3} \\ w_{4,1} & w_{4,2} & w_{4,3} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}}$$

$$\mathbf{y} = \mathbf{x}\mathbf{W}$$

$$(1 \times 4)$$

$$(1 \times 3)$$

$$(3 \times 4)$$

$$[y_1 \quad y_2 \quad y_3 \quad y_4] = [x_1 \quad x_2 \quad x_3] \begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} & w_{1,4} \\ w_{2,1} & w_{2,2} & w_{2,3} & w_{2,4} \\ w_{3,1} & w_{3,2} & w_{3,3} & w_{3,4} \end{bmatrix}$$

$$\frac{d\mathbf{y}}{d\mathbf{x}} = ?$$

Work this out at home. You should be able to show that the derivative (Jacobian) in this case is also equal to \mathbf{W}

Dealing with more than two dimensions

Convert to the row equivalent

Let's consider now the problem of computing the derivative with respect to the matrix \mathbf{W}

$$\mathbf{y} = \mathbf{W}\mathbf{x}$$

$$\begin{matrix} (4 \times 1) & & (4 \times 3) & & (3 \times 1) \\ \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \end{bmatrix} = \begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} \\ w_{2,1} & w_{2,2} & w_{2,3} \\ w_{3,1} & w_{3,2} & w_{3,3} \\ w_{4,1} & w_{4,2} & w_{4,3} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} \end{matrix}$$

$$\frac{d\mathbf{y}}{d\mathbf{W}} = ?$$

Let's define a 3D tensor \mathbf{T} , with elements: $t_{i,j,k} = \frac{\partial y_i}{\partial w_{k,j}}$

A full characterisation of the derivative of \mathbf{y} with respect to \mathbf{W} requires the partial derivative of **each component of \mathbf{y}** with respect to **each component of \mathbf{W}**

$$\begin{bmatrix} \frac{\partial y_1}{\partial w_{1,1}} & \frac{\partial y_1}{\partial w_{1,2}} & \frac{\partial y_1}{\partial w_{1,3}} \\ \frac{\partial y_1}{\partial w_{2,1}} & \frac{\partial y_1}{\partial w_{2,2}} & \frac{\partial y_1}{\partial w_{2,3}} \\ \frac{\partial y_1}{\partial w_{3,1}} & \frac{\partial y_1}{\partial w_{3,2}} & \frac{\partial y_1}{\partial w_{3,3}} \\ \frac{\partial y_1}{\partial w_{4,1}} & \frac{\partial y_1}{\partial w_{4,2}} & \frac{\partial y_1}{\partial w_{4,3}} \\ \frac{\partial y_2}{\partial w_{1,1}} & \frac{\partial y_2}{\partial w_{1,2}} & \frac{\partial y_2}{\partial w_{1,3}} \\ \frac{\partial y_2}{\partial w_{2,1}} & \frac{\partial y_2}{\partial w_{2,2}} & \frac{\partial y_2}{\partial w_{2,3}} \\ \frac{\partial y_2}{\partial w_{3,1}} & \frac{\partial y_2}{\partial w_{3,2}} & \frac{\partial y_2}{\partial w_{3,3}} \\ \frac{\partial y_2}{\partial w_{4,1}} & \frac{\partial y_2}{\partial w_{4,2}} & \frac{\partial y_2}{\partial w_{4,3}} \\ \frac{\partial y_3}{\partial w_{1,1}} & \frac{\partial y_3}{\partial w_{1,2}} & \frac{\partial y_3}{\partial w_{1,3}} \\ \frac{\partial y_3}{\partial w_{2,1}} & \frac{\partial y_3}{\partial w_{2,2}} & \frac{\partial y_3}{\partial w_{2,3}} \\ \frac{\partial y_3}{\partial w_{3,1}} & \frac{\partial y_3}{\partial w_{3,2}} & \frac{\partial y_3}{\partial w_{3,3}} \\ \frac{\partial y_3}{\partial w_{4,1}} & \frac{\partial y_3}{\partial w_{4,2}} & \frac{\partial y_3}{\partial w_{4,3}} \\ \frac{\partial y_4}{\partial w_{1,1}} & \frac{\partial y_4}{\partial w_{1,2}} & \frac{\partial y_4}{\partial w_{1,3}} \\ \frac{\partial y_4}{\partial w_{2,1}} & \frac{\partial y_4}{\partial w_{2,2}} & \frac{\partial y_4}{\partial w_{2,3}} \\ \frac{\partial y_4}{\partial w_{3,1}} & \frac{\partial y_4}{\partial w_{3,2}} & \frac{\partial y_4}{\partial w_{3,3}} \\ \frac{\partial y_4}{\partial w_{4,1}} & \frac{\partial y_4}{\partial w_{4,2}} & \frac{\partial y_4}{\partial w_{4,3}} \end{bmatrix} \quad (4 \times 3 \times 4)$$

Dealing with more than two dimensions

Convert to
the row
equivalent

Same as before, let's compute just one of these components, e.g. the derivative of y_2 to $w_{1,3}$

$$y_2 = \sum_{j=1}^3 w_{2,j} x_j$$

$$y_2 = w_{2,1}x_1 + w_{2,2}x_2 + w_{2,3}x_3$$

$$\frac{\partial y_2}{\partial w_{1,3}} = \frac{\partial}{\partial w_{1,3}} [w_{2,1}x_1 + w_{2,2}x_2 + w_{2,3}x_3]$$

$$\frac{\partial y_2}{\partial w_{1,3}} = 0$$

The only derivatives y_2 that are non-zero are the ones involving the second column of \mathbf{W} , the elements: $w_{i,2}$

For example:

$$\frac{\partial y_2}{\partial w_{2,3}} = \frac{\partial}{\partial w_{2,3}} [w_{2,1}x_1 + w_{2,2}x_2 + w_{2,3}x_3] = x_3$$

In general:

$$\frac{\partial y_i}{\partial w_{i,j}} = x_j$$

Dealing with more than two dimensions

Convert to the row equivalent

Most elements will be zero, except for the elements for which $i = k$.

$$t_{i,j,k} = \begin{cases} x_j & , \text{if } i = k \\ 0 & , \text{otherwise} \end{cases}$$

$$\begin{bmatrix} \frac{\partial y_1}{\partial w_{1,1}} & \frac{\partial y_1}{\partial w_{1,2}} & \frac{\partial y_1}{\partial w_{1,3}} & \frac{\partial y_1}{\partial w_{2,1}} & \frac{\partial y_1}{\partial w_{2,2}} & \frac{\partial y_1}{\partial w_{2,3}} & \frac{\partial y_1}{\partial w_{3,1}} & \frac{\partial y_1}{\partial w_{3,2}} & \frac{\partial y_1}{\partial w_{3,3}} \\ \frac{\partial y_2}{\partial w_{1,1}} & \frac{\partial y_2}{\partial w_{1,2}} & \frac{\partial y_2}{\partial w_{1,3}} & \frac{\partial y_2}{\partial w_{2,1}} & \frac{\partial y_2}{\partial w_{2,2}} & \frac{\partial y_2}{\partial w_{2,3}} & \frac{\partial y_2}{\partial w_{3,1}} & \frac{\partial y_2}{\partial w_{3,2}} & \frac{\partial y_2}{\partial w_{3,3}} \\ \frac{\partial y_3}{\partial w_{1,1}} & \frac{\partial y_3}{\partial w_{1,2}} & \frac{\partial y_3}{\partial w_{1,3}} & \frac{\partial y_3}{\partial w_{2,1}} & \frac{\partial y_3}{\partial w_{2,2}} & \frac{\partial y_3}{\partial w_{2,3}} & \frac{\partial y_3}{\partial w_{3,1}} & \frac{\partial y_3}{\partial w_{3,2}} & \frac{\partial y_3}{\partial w_{3,3}} \\ \frac{\partial y_4}{\partial w_{1,1}} & \frac{\partial y_4}{\partial w_{1,2}} & \frac{\partial y_4}{\partial w_{1,3}} & \frac{\partial y_4}{\partial w_{2,1}} & \frac{\partial y_4}{\partial w_{2,2}} & \frac{\partial y_4}{\partial w_{2,3}} & \frac{\partial y_4}{\partial w_{3,1}} & \frac{\partial y_4}{\partial w_{3,2}} & \frac{\partial y_4}{\partial w_{3,3}} \end{bmatrix} = \begin{bmatrix} 0 & 0 & 0 \\ x_1 & x_2 & x_3 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

If $y_{i,:}$ is the i^{th} element of \mathbf{y} and $W_{i,:}$ is the i^{th} row of \mathbf{W} then $\frac{\partial y_i}{\partial W_{i,:}} = \mathbf{x}$

All the non-trivial portion of this tensor can be stored in a compact way in a 2D matrix

Multiple data points

Let's now use multiple row-vector samples $x^{(i)}$, stacked together to form a matrix \mathbf{X} .

$$\mathbf{Y} = \mathbf{XW}$$

If $Y_{i,:}$ is the i^{th} row of \mathbf{Y} and $X_{i,:}$ is the i^{th} row of \mathbf{X} it is easy to show that

$$\frac{\partial Y_{i,:}}{\partial X_{i,:}} = \mathbf{W}$$

Work this out at home

$$\text{and } \frac{\partial Y_{i,:}}{\partial X_{j,:}} = ? \quad \text{if } i \neq j$$

The chain rule

$$\mathbf{y} = \mathbf{V}\mathbf{W}\mathbf{x} \qquad \frac{d\mathbf{y}}{d\mathbf{x}} = \mathbf{V}\mathbf{W}$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} \qquad \frac{d\mathbf{z}}{d\mathbf{x}} = \mathbf{W}$$

$$\mathbf{y} = \mathbf{V}\mathbf{z} \qquad \frac{d\mathbf{y}}{d\mathbf{z}} = \mathbf{V}$$

$$\frac{d\mathbf{y}}{d\mathbf{x}} = \frac{d\mathbf{y}}{d\mathbf{z}} \frac{d\mathbf{z}}{d\mathbf{x}} = \mathbf{V}\mathbf{W}$$

“Vector, Matrix, and Tensor Derivatives”,

Erik Learned-Miller

<http://cs231n.stanford.edu/vecDerivs.pdf>

Pytorch made easy

```
import torch
dtype = torch.float
device = torch.device("cpu")
N, D_in, H, D_out = 64, 1000, 100, 10      # N is batch size; D_in is input dimension; H is hidden dimension; D_out is output dimension.
# Create random Tensors to hold input and outputs.
# Setting requires_grad=False indicates that we do not need to compute gradients with respect to these Tensors during the backward pass.
x = torch.randn(N, D_in, device=device, dtype= torch.float)
y = torch.randn(N, D_out, device=device, dtype= torch.float)
# Create random Tensors for weights. Setting requires_grad=True indicates that we want to compute gradients with respect to these Tensors during the backward pass.
w1 = torch.randn(D_in, H, device=device, dtype= torch.float, requires_grad=True)
w2 = torch.randn(H, D_out, device=device, dtype= torch.float, requires_grad=True)

learning_rate = 1e-6
for t in (500):
    # Forward pass: compute predicted y using operations on Tensors; these are exactly the same operations we used to compute the forward pass using Tensors, but
    # we do not need to keep references to intermediate values since we are not implementing the backward pass by hand.
    y_pred = x.mm(w1).clamp(min=0).mm(w2)

    # Compute and print loss using operations on Tensors. Now loss is a Tensor of shape (1,), loss.item() gets the scalar value held in the loss.
    loss = (y_pred - y).pow(2).sum()

    # Use autograd to compute the backward pass. This call will compute the gradient of loss with respect to all Tensors with requires_grad=True. After this call
    # w1.grad and w2.grad will be Tensors holding the gradient of the loss with respect to w1 and w2 respectively.
    loss.backward()

    # Manually update weights using gradient descent. Wrap in torch.no_grad() because weights have requires_grad=True, but we don't need to track this in autograd.
    with torch.no_grad():
        w1 -= learning_rate * w1.grad
        w2 -= learning_rate * w2.grad
    # Manually zero the gradients after updating weights
    w1.grad.zero_()
    w2.grad.zero_()
```

Summary

- Backpropagation computes the **gradient** of the **loss function** with respect to the **weights** of the network for a single input–output example
- Auto Differentiation (AutoGrad) is at the core of modern deep learning frameworks, and enables efficient backpropagation schemes
 - Single pass process to calculate all needed derivatives
 - The key is that the process is always local (children nodes to parent nodes)
 - Scalable: we only need to compute stuff once, for all variables
 - Flexible: we can define new models easily (usually transparent from the end user)
 - Vectorizable: use matrix calculus

Still Not A Learning algorithm

- We know how to compute error derivatives for every weight on a single training point
- We got an idea about how to extend this for a whole batch of points
- We still need to see
 - **Loss functions**: How to measure our error? This depends on the task we want to solve.
 - **Activation functions**: What kind of neurons are there (neurons are defined by their integration and activation functions)?
 - **Architectures**: How to combine neurons together to build meaningful models?
 - **Optimisation**: Is batch gradient descent the best way to use these error derivatives to discover a good set of weights?
 - **Regularisation**: How do we make sure we do not overfit?
 - **Initialisation**: Where do we start our search?

More Information

- Some material on these slides has been adapted from various sources including the following highly recommended ones:
 - Andrew Ng's *Machine Learning Course*, Coursera
<https://www.coursera.org/course/ml>
 - Andrew Ng's *Deep Learning Specialization*, Coursera
<https://www.coursera.org/specializations/deep-learning>
 - Victor Lavrenko's *Machine Learning Course*
<https://www.youtube.com/channel/UCs7aIOMRnxhkfKAJ4JjZ7Wg>
 - Fei Fei Li and Andrej Karpathy's *Convolutional Neural Networks for Visual Recognition*
<http://cs231n.stanford.edu/>
 - Geoff Hinton's *Neural Networks for Machine Learning*, Coursera
<https://www.coursera.org/learn/neural-networks>
 - Luis Serrano's introductory videos
<https://www.youtube.com/channel/UCgBncpylJ1kiVaPyP-PZauQ>