


Neural Networks and Deep Learning

Architectures for Image Classification



What's in a scene



Type of environment	Indoor
Scene categories	martial_arts_gym (0.204)
Scene attributes	no horizon, enclosed area, man-made, cloth, natural light, competing, exercise, sports, glossy
Informative región for predicting the top category	

What's in a scene





Type of environment	Indoor	Outdoor
Scene categories	martial_arts_gym (0.204)	apartment_building/outdoor (0.195), office_building (0.140), hospital (0.131), parking_garage/outdoor (0.126)
Scene attributes	no horizon, enclosed area, man-made, cloth, natural light, competing, exercise, sports, glossy	natural light, man-made, open area, no horizon, sunny, glass, vegetation, vertical components, leaves
Informative región for predicting the top category		

What's in a scene

Binary
classification





Type of environment	Indoor	Outdoor
Scene categories	martial_arts_gym (0.204)	apartment_building/outdoor (0.195), office_building (0.140), hospital (0.131), parking_garage/outdoor (0.126)
Scene attributes	no horizon, enclosed area, man-made, cloth, natural light, competing, exercise, sports, glossy	natural light, man-made, open area, no horizon, sunny, glass, vegetation, vertical components, leaves
Informative región for predicting the top category		

What's in a scene

Multi-class
classification





Type of environment	Indoor	Outdoor
Scene categories	martial_arts_gym (0.204)	apartment_building/outdoor (0.195), office_building (0.140), hospital (0.131), parking_garage/outdoor (0.126)
Scene attributes	no horizon, enclosed area, man-made, cloth, natural light, competing, exercise, sports, glossy	natural light, man-made, open area, no horizon, sunny, glass, vegetation, vertical components, leaves
Informative región for predicting the top category		

What's in a scene

Multi-label
classification



Type of environment	Indoor	Outdoor
Scene categories	martial_arts_gym (0.204)	apartment_building/outdoor (0.195), office_building (0.140), hospital (0.131), parking_garage/outdoor (0.126)
Scene attributes	no horizon, enclosed area, man-made, cloth, natural light, competing, exercise, sports, glossy	natural light, man-made, open area, no horizon, sunny, glass, vegetation, vertical components, leaves
Informative region for predicting the top category		

What's in a scene

Understanding
/ visualising
CNNs





Type of environment	Indoor	Outdoor
Scene categories	martial_arts_gym (0.204)	apartment_building/outdoor (0.195), office_building (0.140), hospital (0.131), parking_garage/outdoor (0.126)
Scene attributes	no horizon, enclosed area, man-made, cloth, natural light, competing, exercise, sports, glossy	natural light, man-made, open area, no horizon, sunny, glass, vegetation, vertical components, leaves
Informative region for predicting the top category		

Image Classification

PROBLEM METRICS AND DATASETS

Evaluation Metrics

$$\text{Accuracy} = \frac{\# \text{ correct predictions}}{\# \text{ total predictions}}$$

$$\text{Misclassification Error} = 1 - \text{Accuracy}$$



Scene
categories

apartment_building/outdoor (0.195),
office_building (0.140),
hospital (0.131),
parking_garage/outdoor (0.126)

Evaluation Metrics

$$\text{Accuracy} = \frac{\# \text{ correct predictions}}{\# \text{ total predictions}}$$

$$\text{Misclassification Error} = 1 - \text{Accuracy}$$

$$\text{Top-k Accuracy} = \frac{\# \text{ correct prediction in the top } k}{\# \text{ total predictions}}$$

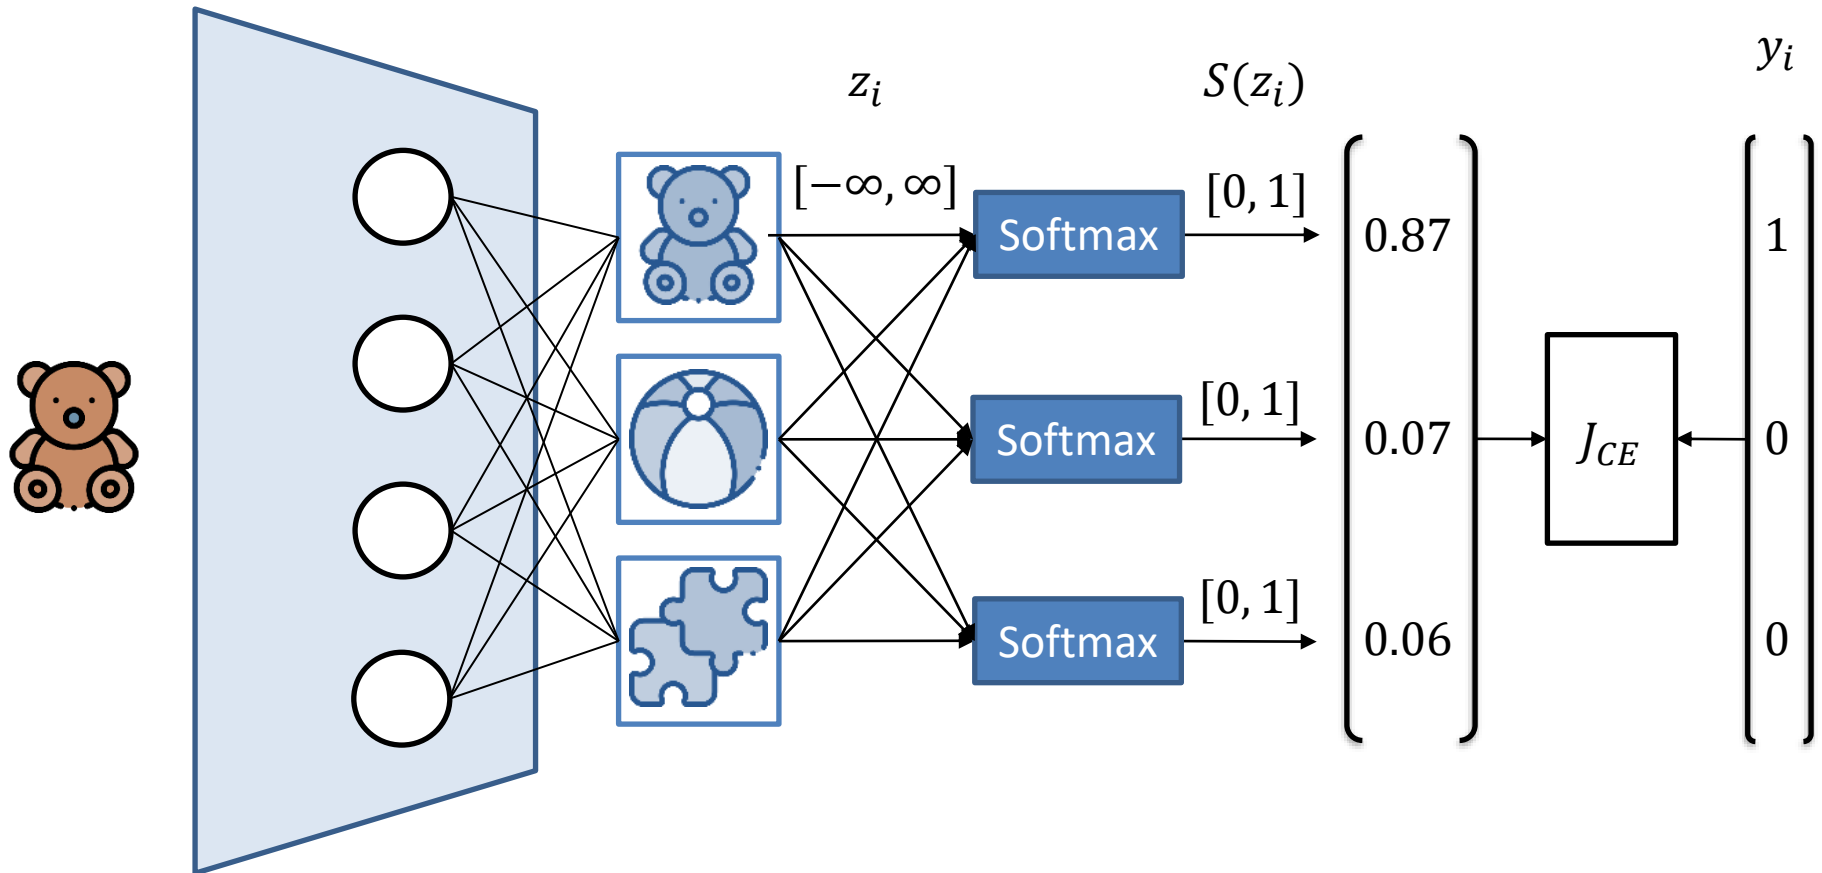
$$\text{Top-k Classification Error} = 1 - \text{Top-k Accuracy}$$



Scene
categories

apartment_building/outdoor (0.195),
office_building (0.140),
hospital (0.131),
parking_garage/outdoor (0.126)

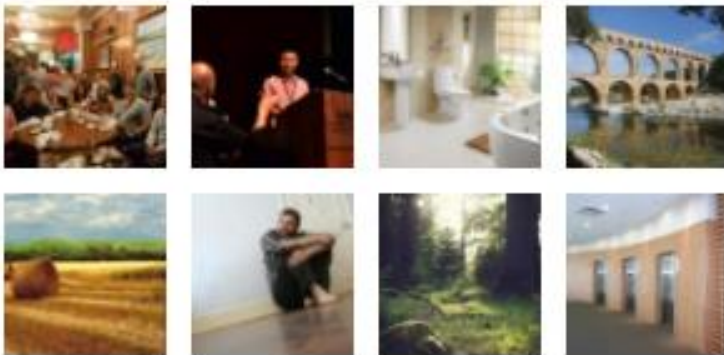
Multi-class classification



Datasets



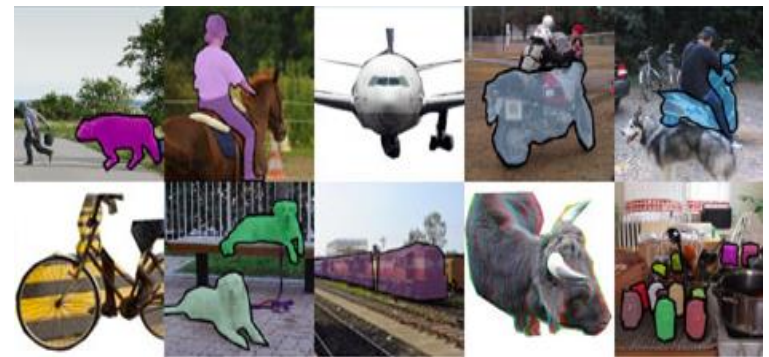
- Scene photographs from image search engines, non-uniform distribution of images per category.
- 8M train, 36K val, 328K test, 1000 classes.



from ILSVRC 2016, Places Dataset



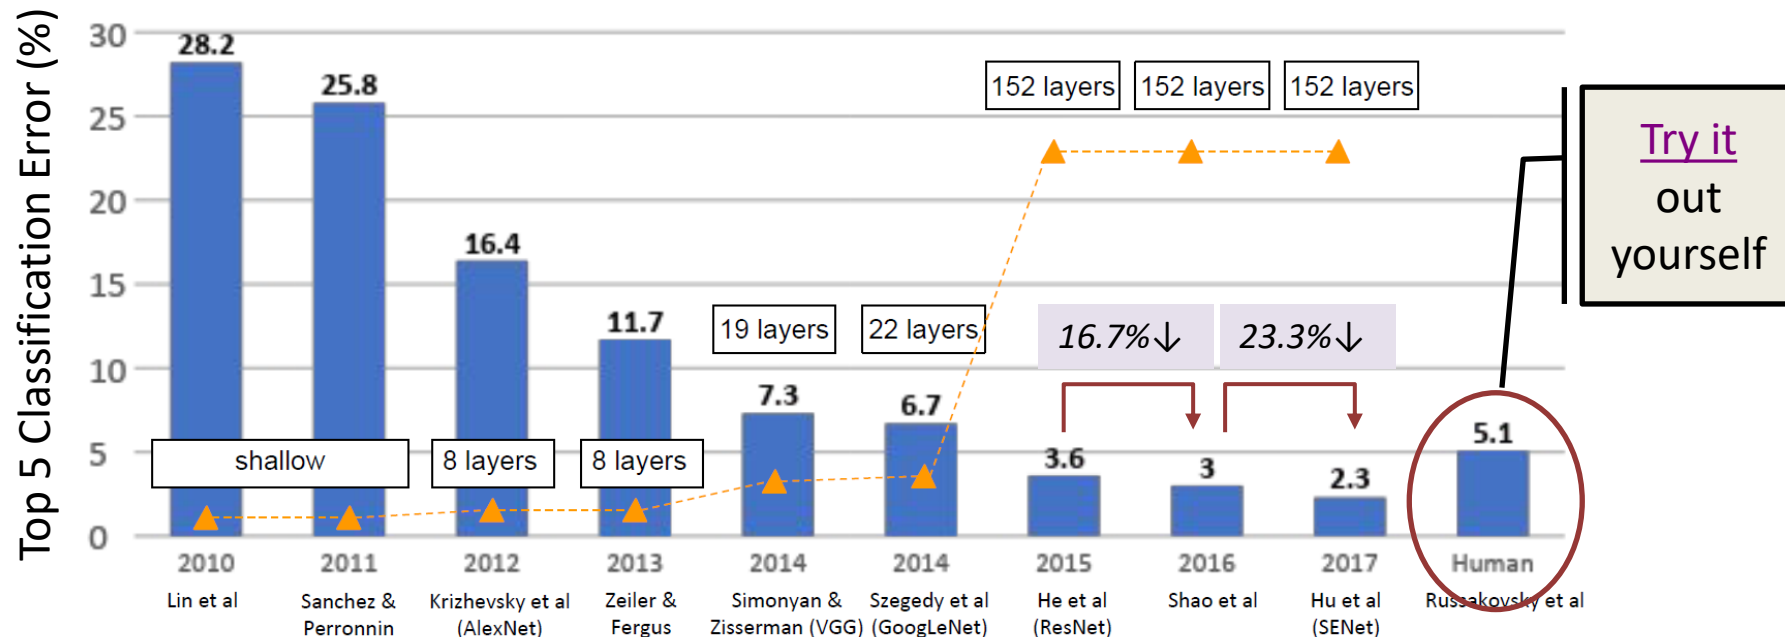
- Photos of 91 object types in the context of the broader question of scene understanding that would be easily recognizable by a 4 year old.
- 2.5 million labeled instances in 328k images.



from the COCO Dataset

ImageNet Large Scale Visual Recognition Challenge

- Standard benchmark for image classification
- Standardised evaluation (top-k classification error)
- Evolution over time



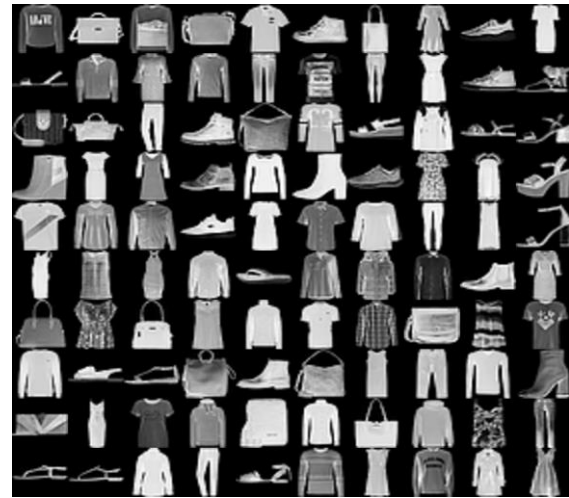
Prototype Datasets

MNIST



- 60000 images of 10 handwritten digits of 28x28 pixels
- Accuracy: 99.9 %
- Usage: test if your model runs correctly

Fashion-Mnist



- 50000 images of 10 different clothes from Zalando of 28x28 pixels
- Accuracy: 96 %
- Usage: test if your model runs correctly

Prototype Datasets

CIFAR

airplane



automobile



bird



cat



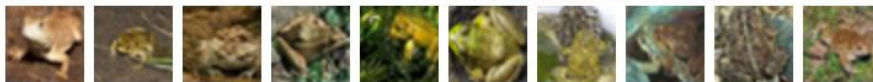
deer



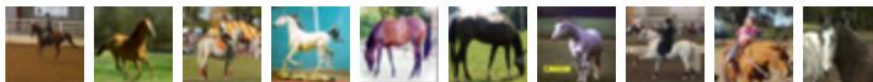
dog



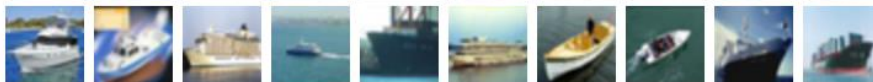
frog



horse



ship



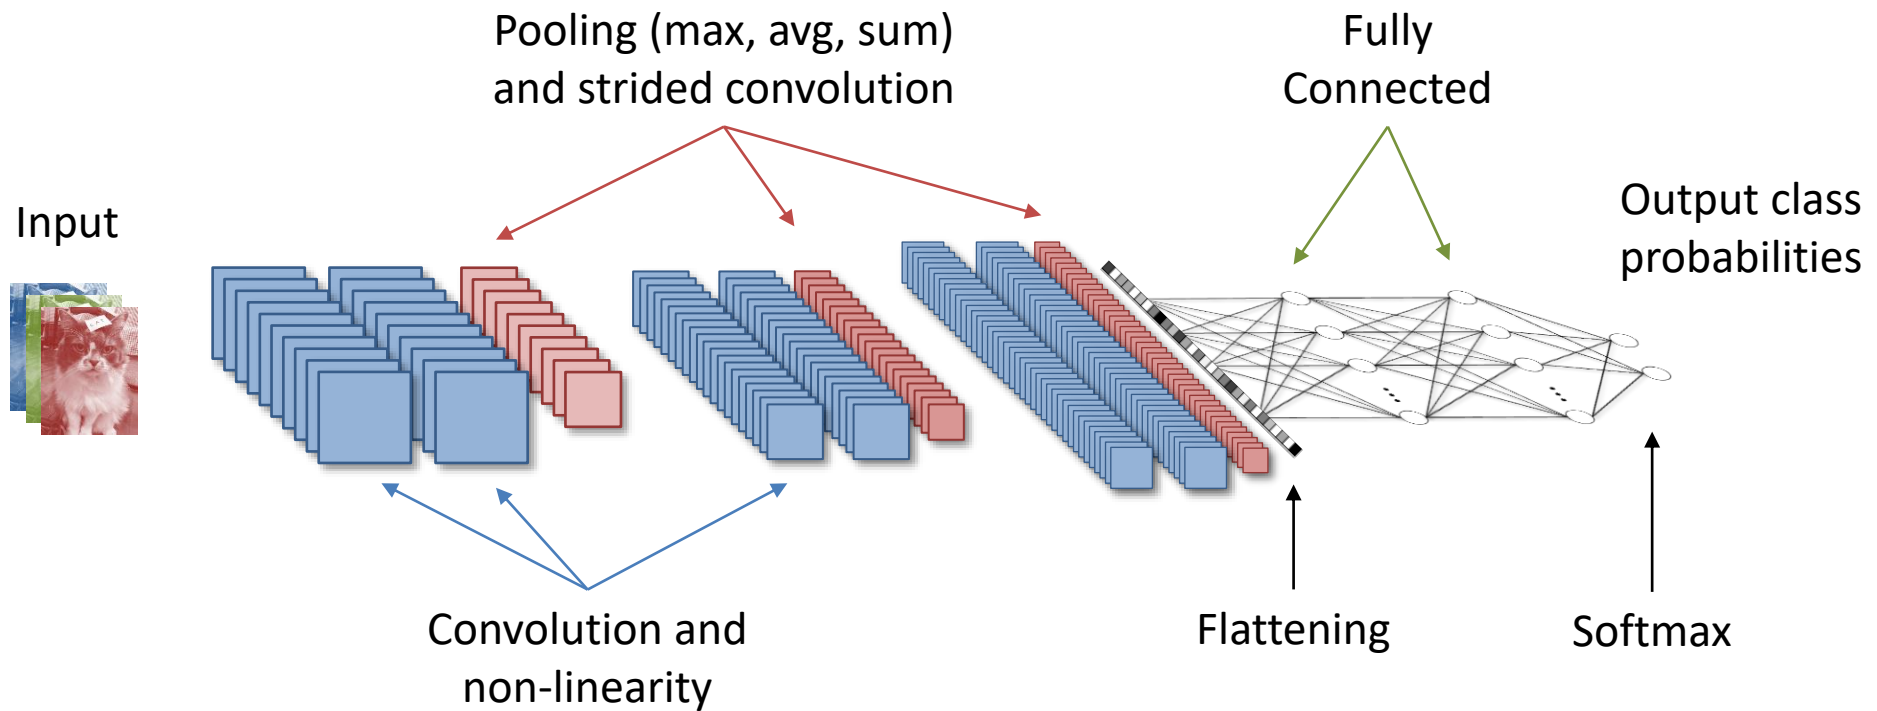
truck



- ◆ 60000 images of 32x32 color pixels
- ◆ Organized in 10 or 100 classes (CIFAR10 and CIFAR100)

BASIC CNN MODELS

Typical CNN architecture



Different CNN Diagrams

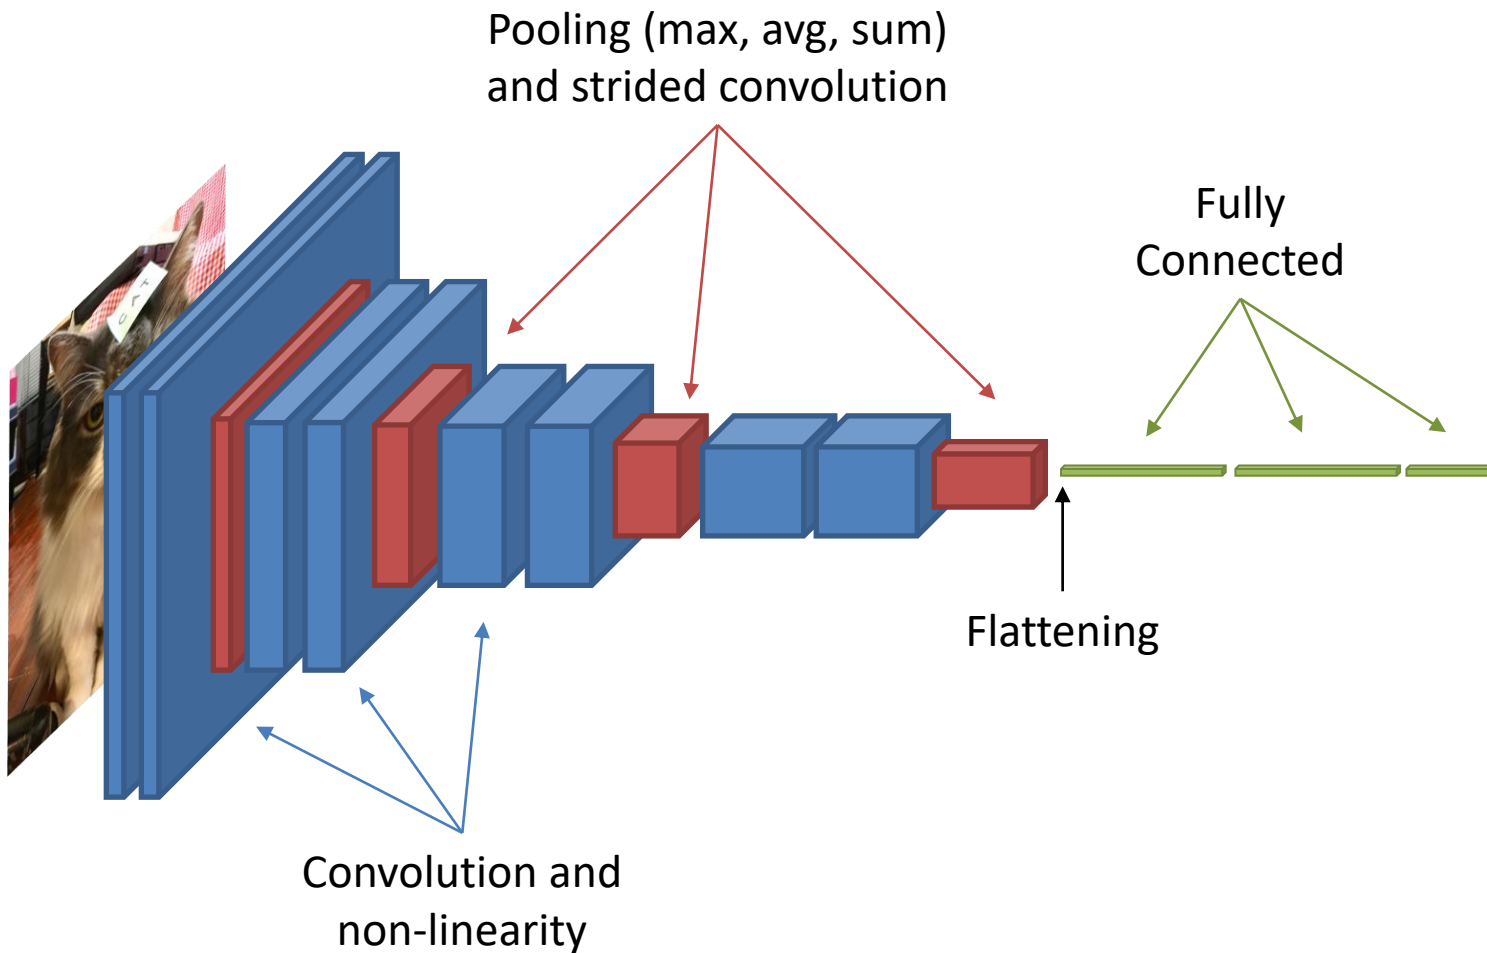
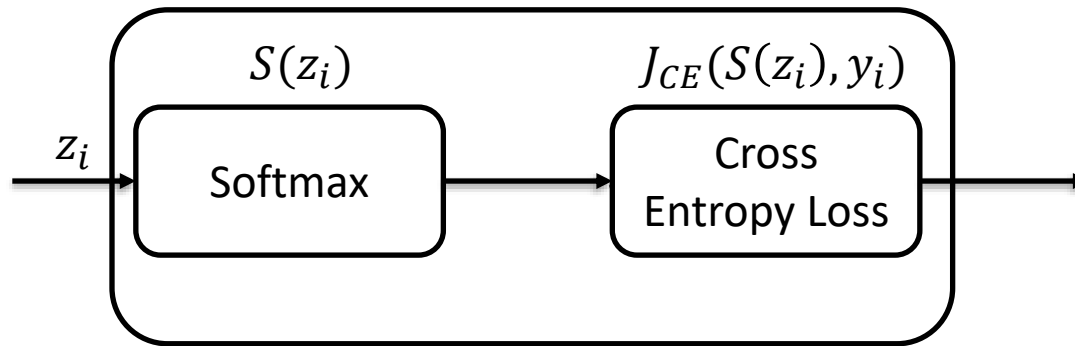
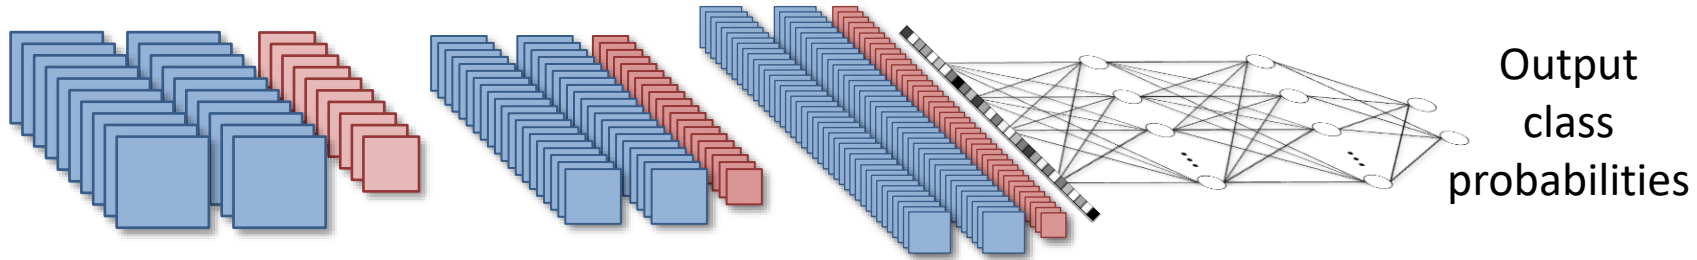


Image
Conv-64
Conv-64
maxpool
Conv-128
Conv-128
maxpool
Conv-256
Conv-256
maxpool
Conv-512
Conv-512
maxpool
FC-4096
FC-4096
FC-1000
Softmax

CNN architecture

Input

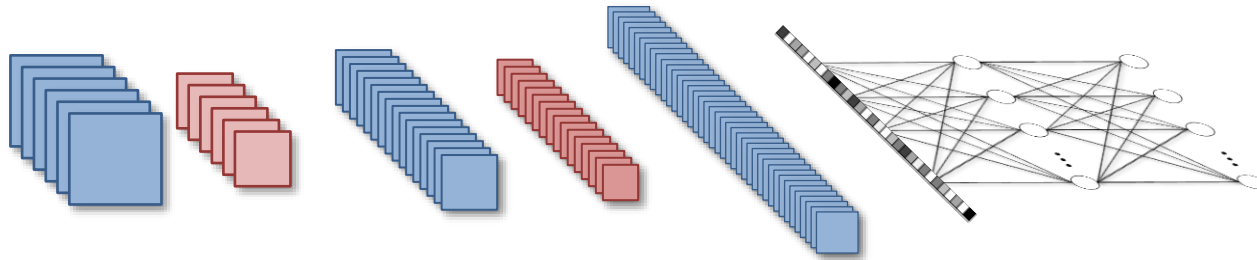


$$S(z)_i = \frac{e^{z_i}}{\sum_j^C e^{z_j}}$$

$$J_{CCE} = - \sum_i^C y_c \log(S(z)_i)$$

LeNet (1998)

Input

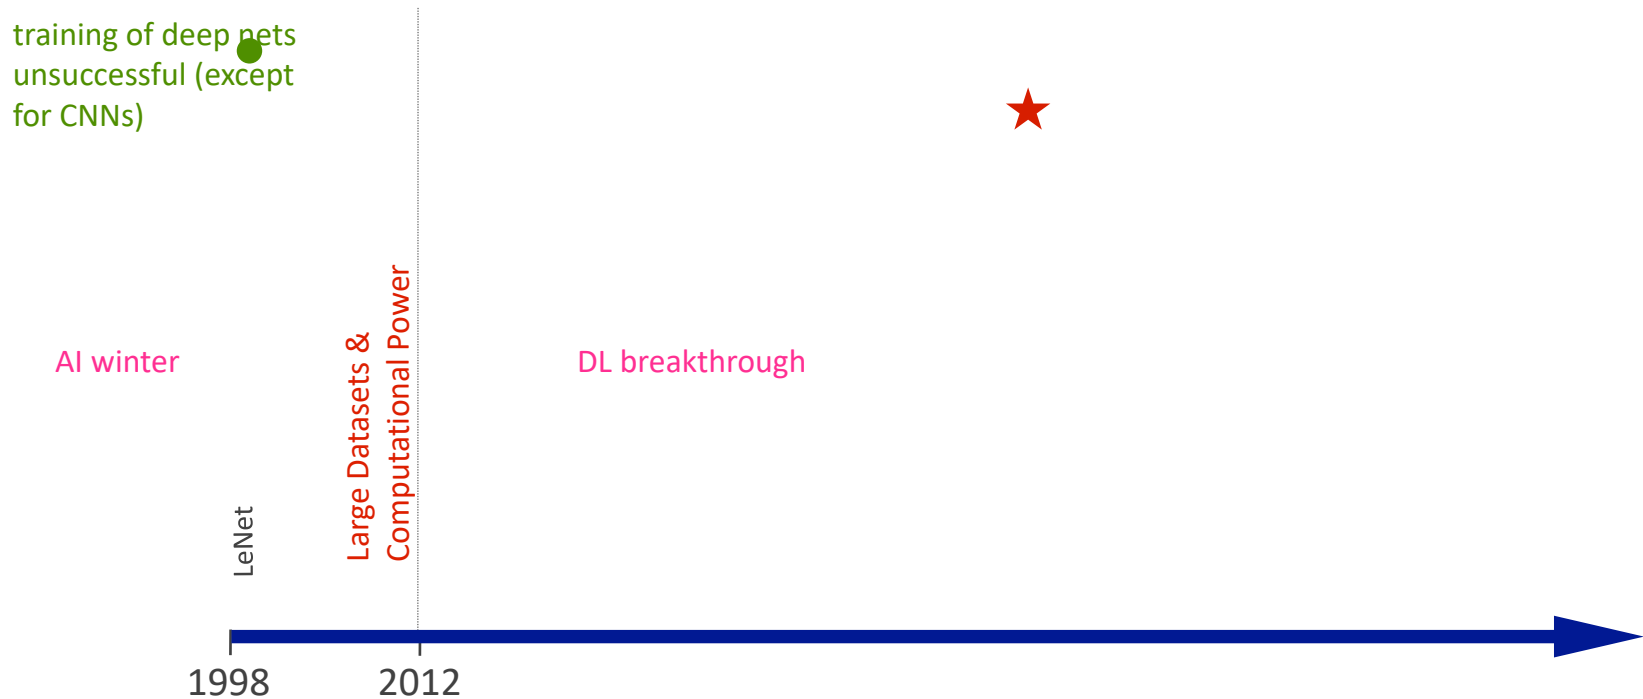


Output
class
probabilities

Conv. 1	Sub. 1	Conv. 2	Sub. 2	Conv. 2	FC 1	Output
(5x5, 6)	(2x2)	(5x5, 16)	(2x2)	(5x5, 120)	(84)	(10)

Nonlinearity used: *tanh*

Image Classification Evolution

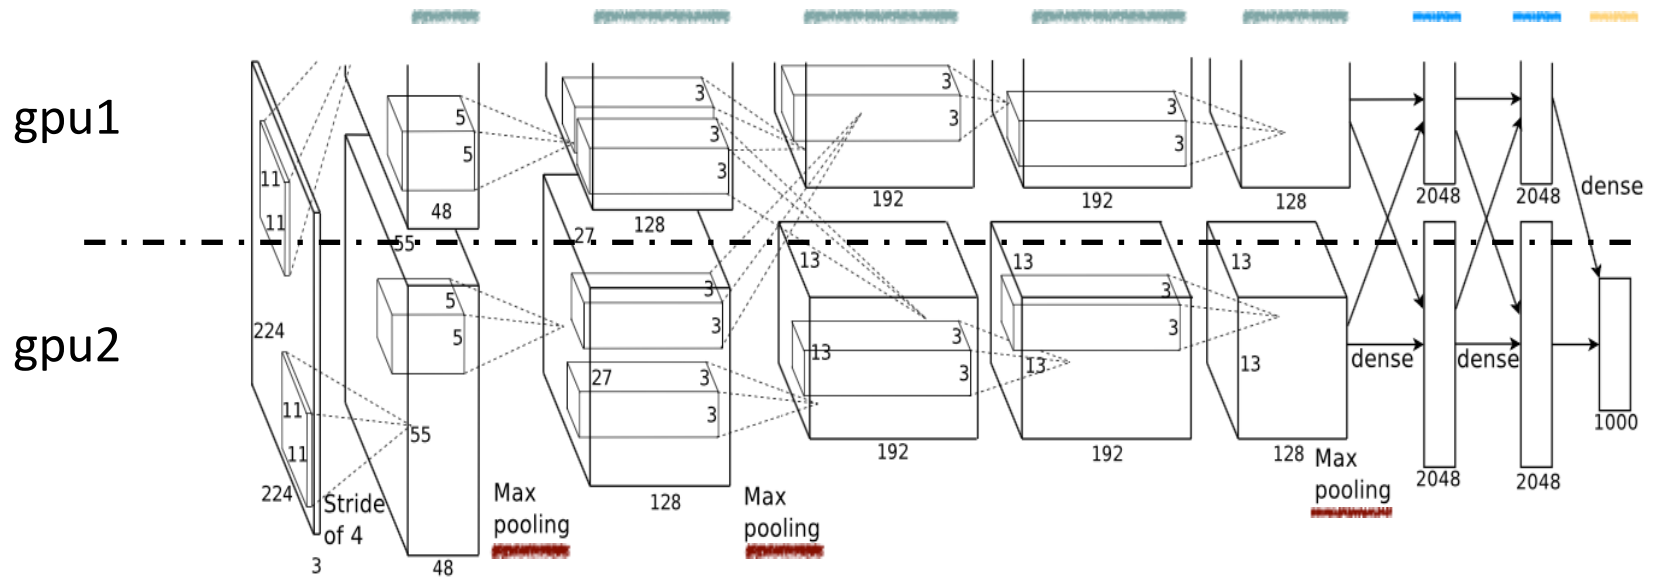


AlexNet, VGG, GoogLeNet

2012 - 2014

AlexNet (2012)

8 convolutional (11x11, 5x5, 3x3, 3x3, 3x3) and 3 fully-connected layers.
Took ~6 days on two GTX 580 3GB GPUs (own implementation)



Key concepts introduced to mainstream CNN methodology:

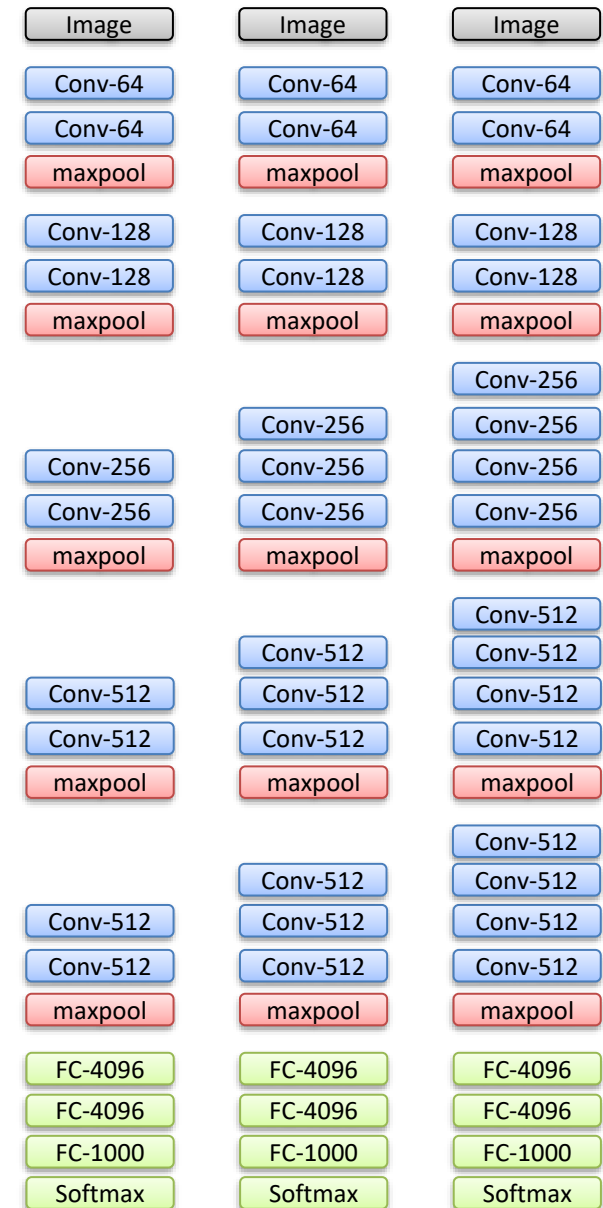
- **ReLU non-linearity** [see [here](#) for details]
- **Dropout** [see [here](#) for details]
- **Data Augmentation**

Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). Imagenet classification with deep convolutional neural networks. Advances in neural information processing systems, 25, 1097-1105.

VGG (2014)

“Our main contribution is a thorough evaluation of **networks of increasing depth** using an architecture with **very small (3×3) convolution** filters [...] significant improvement can be achieved by **pushing the depth to 16–19** weight layers.”

- “Very” deep network (up to 19 layers)
- **3x3 filters**, stride 1, padding 1
- **Stacked** convolutions
- 2x2 non-overlapping max-pooling
- # features increases as we go deeper
- ReLU non-linearity
- Data Augmentation
- Dropout

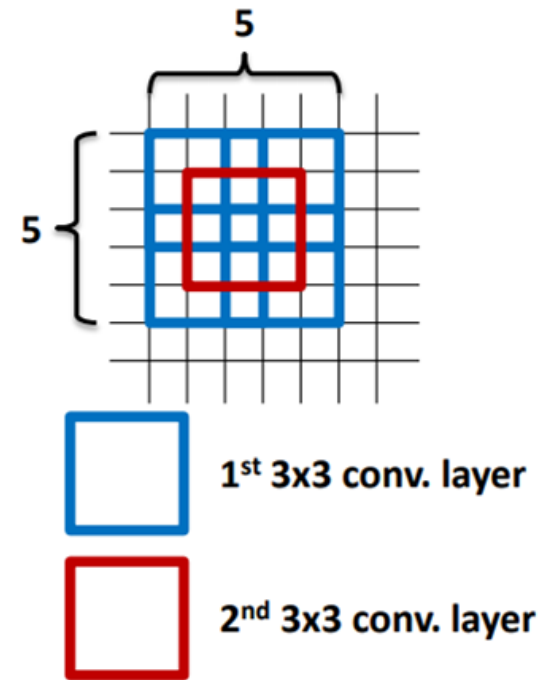


Simonyan, K., & Zisserman, A. (2014). Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556*

VGG (2014)

Why 3x3 layers?

What is the receptive field if we stack two 3×3 convolutions?



VGG (2014)

Why 3x3 layers?

How many parameters are in a layer if we use (a) 3×3 filters? How many if we use (b) 5×5 filters?



(a) $3 \times 3 \times C \times C + C \cong 590\text{k}$ (for $C = 256$)

(b) $5 \times 5 \times C \times C + C \cong 1.6\text{M}$ (for $C = 256$)

VGG (2014)

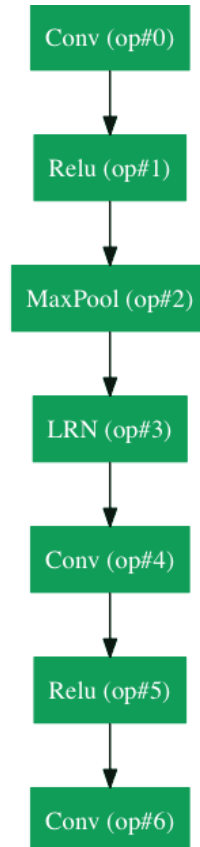
Why 3x3 layers?

- Stacked conv. layers have a large receptive field
 - two 3x3 layers: *5x5 receptive field*
 - three 3x3 layers: *7x7 receptive field*
- More non-linearity
- Less parameters to learn (compared to using large kernels)
 - ~140M per net

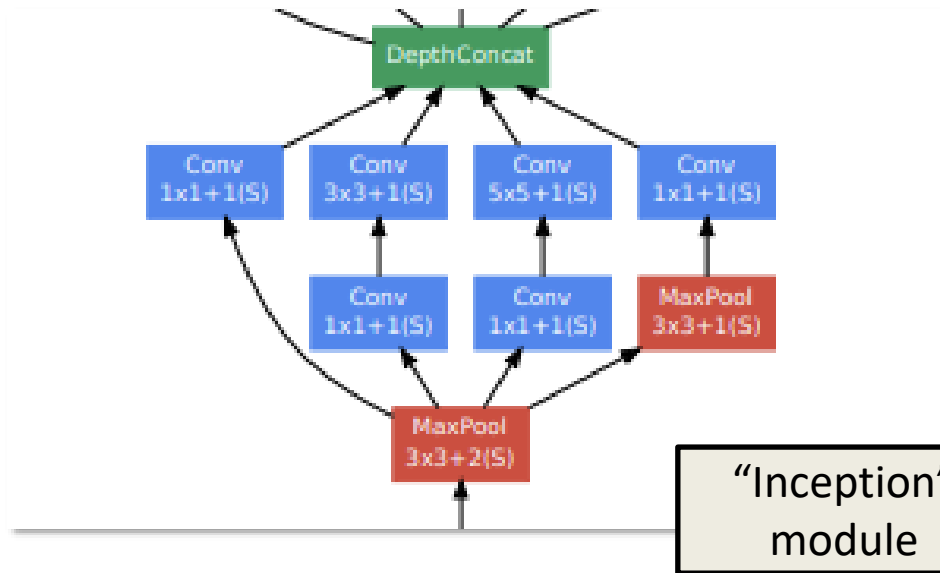
GoogLeNet (2015)



GoogLeNet (2015)

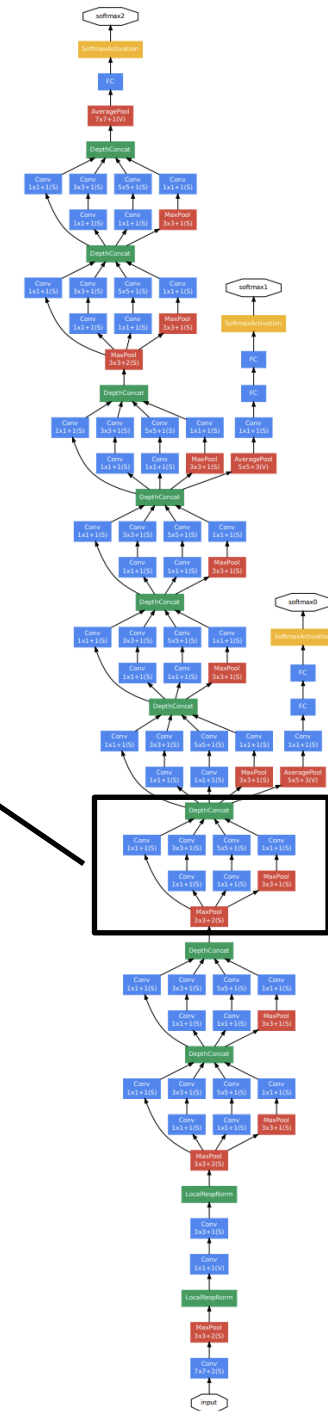


GoogLeNet (2015)

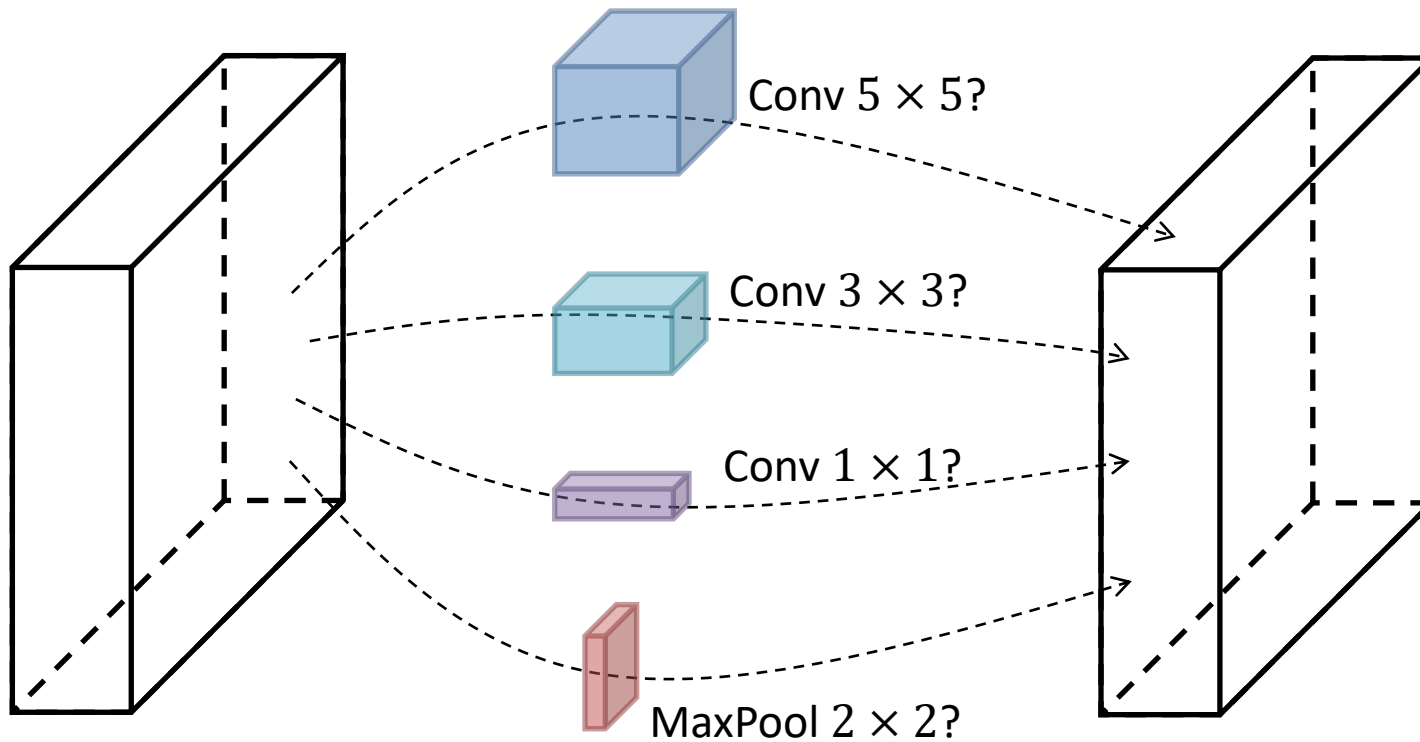


“Inception”
module

- Convolution / FC
- Max / Avg Pooling
- Softmax
- Concatenation

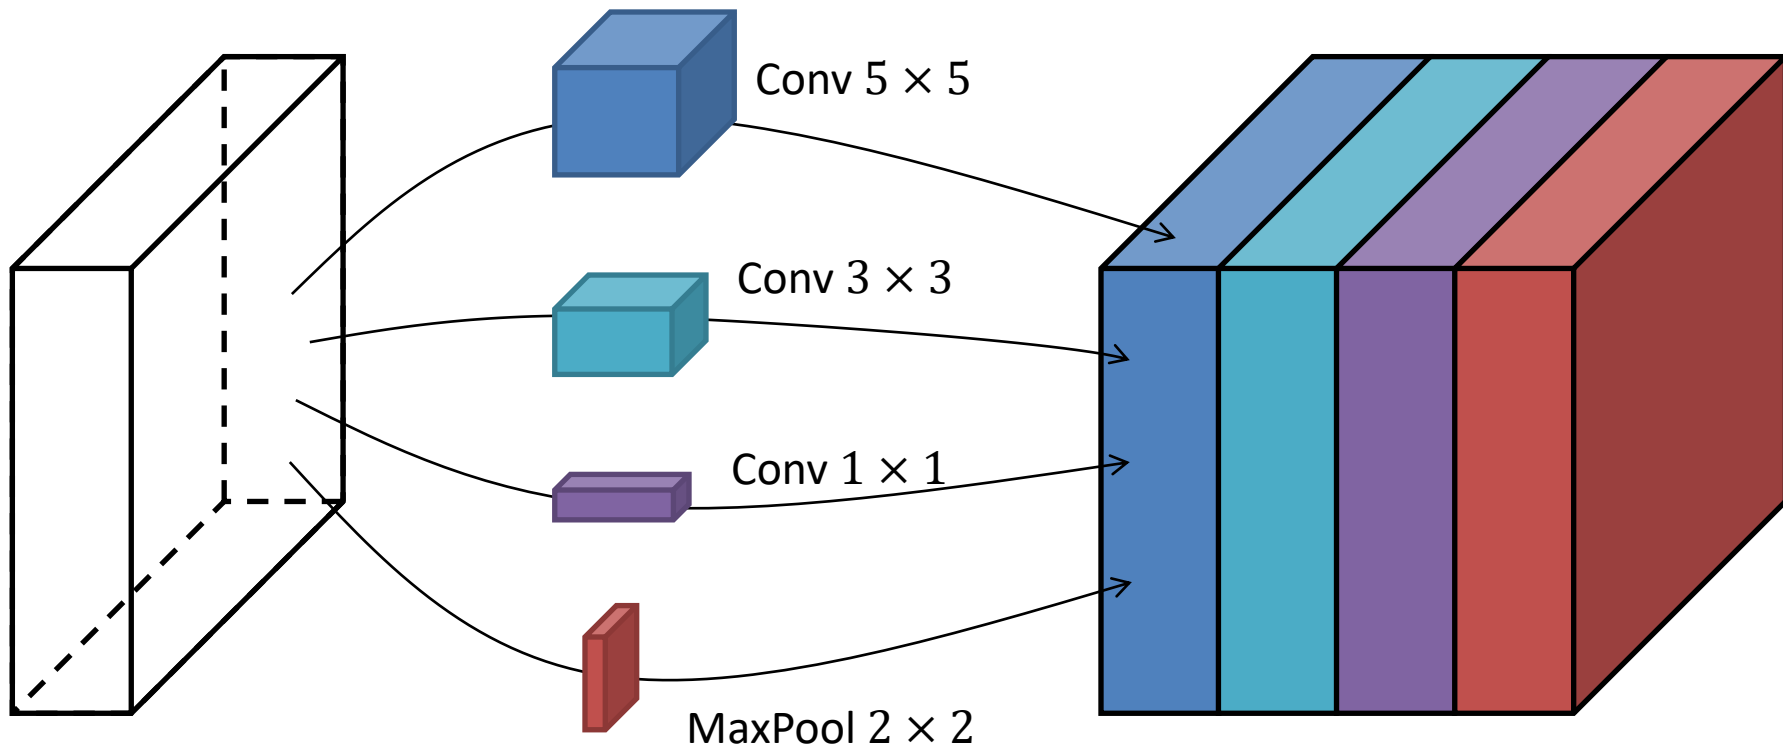


GoogLeNet (2015) – Inception Module



Which operation should we choose?

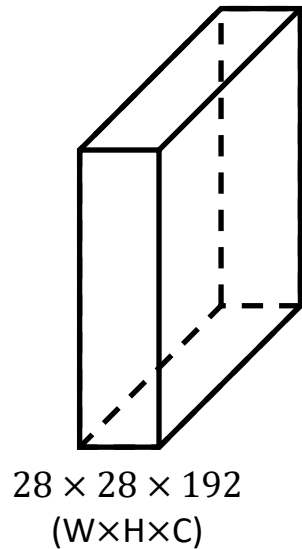
GoogLeNet (2015) – Inception Module



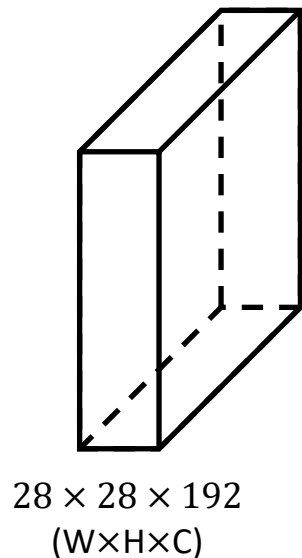
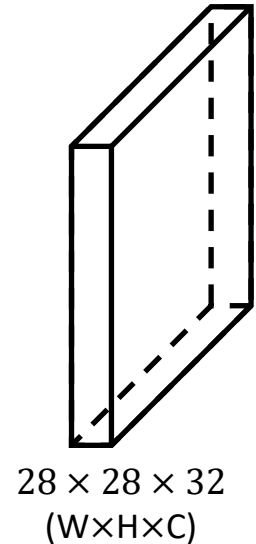
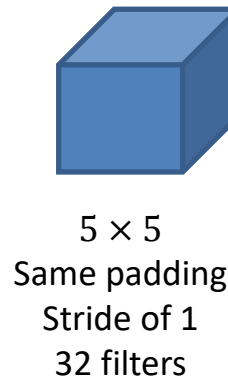
Let's do them all at once! Multi-scale

But the size explodes...

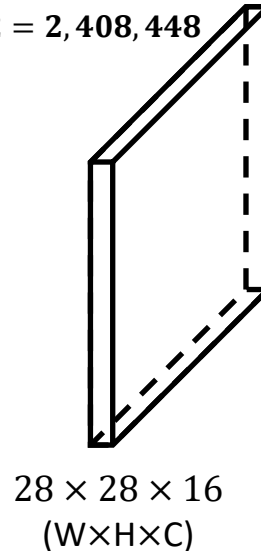
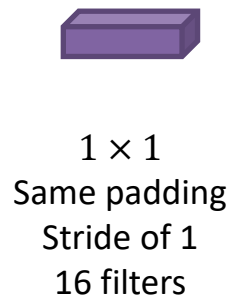
Reducing the Computational Cost



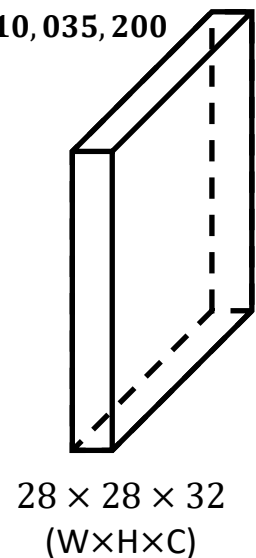
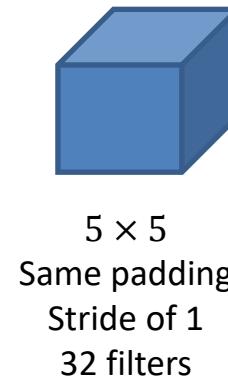
Number of multiplications:
 $28 \times 28 \times 32 \times 5 \times 5 \times 192 = 120,422,400$



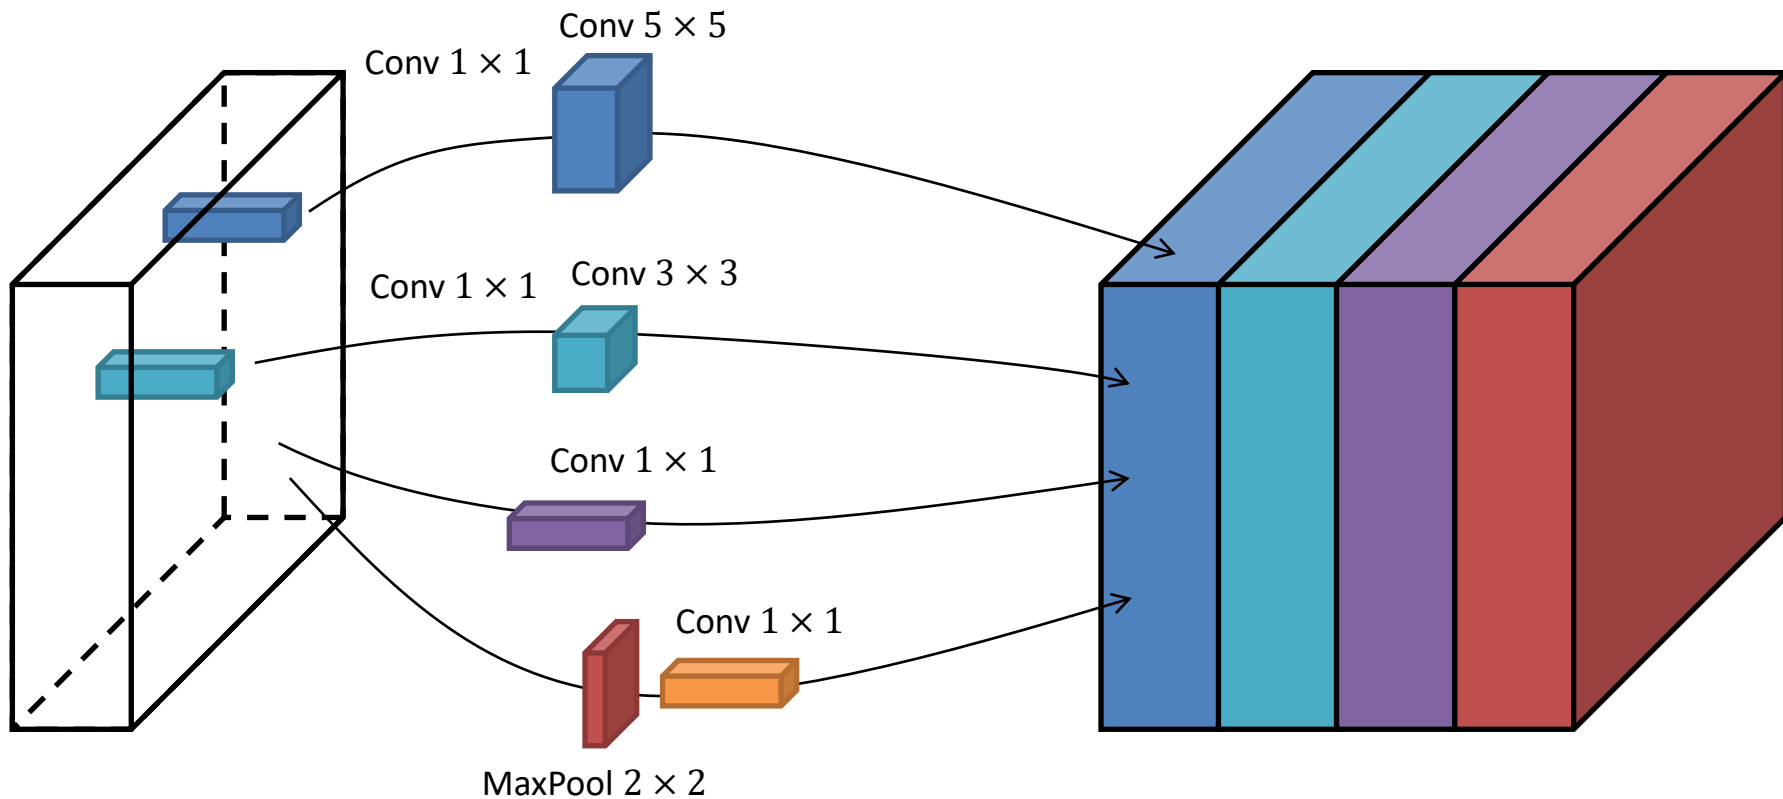
$28 \times 28 \times 16 \times 1 \times 1 \times 192 = 2,408,448$



$28 \times 28 \times 32 \times 5 \times 5 \times 16 = 10,035,200$



GoogLeNet (2015) – Inception Module

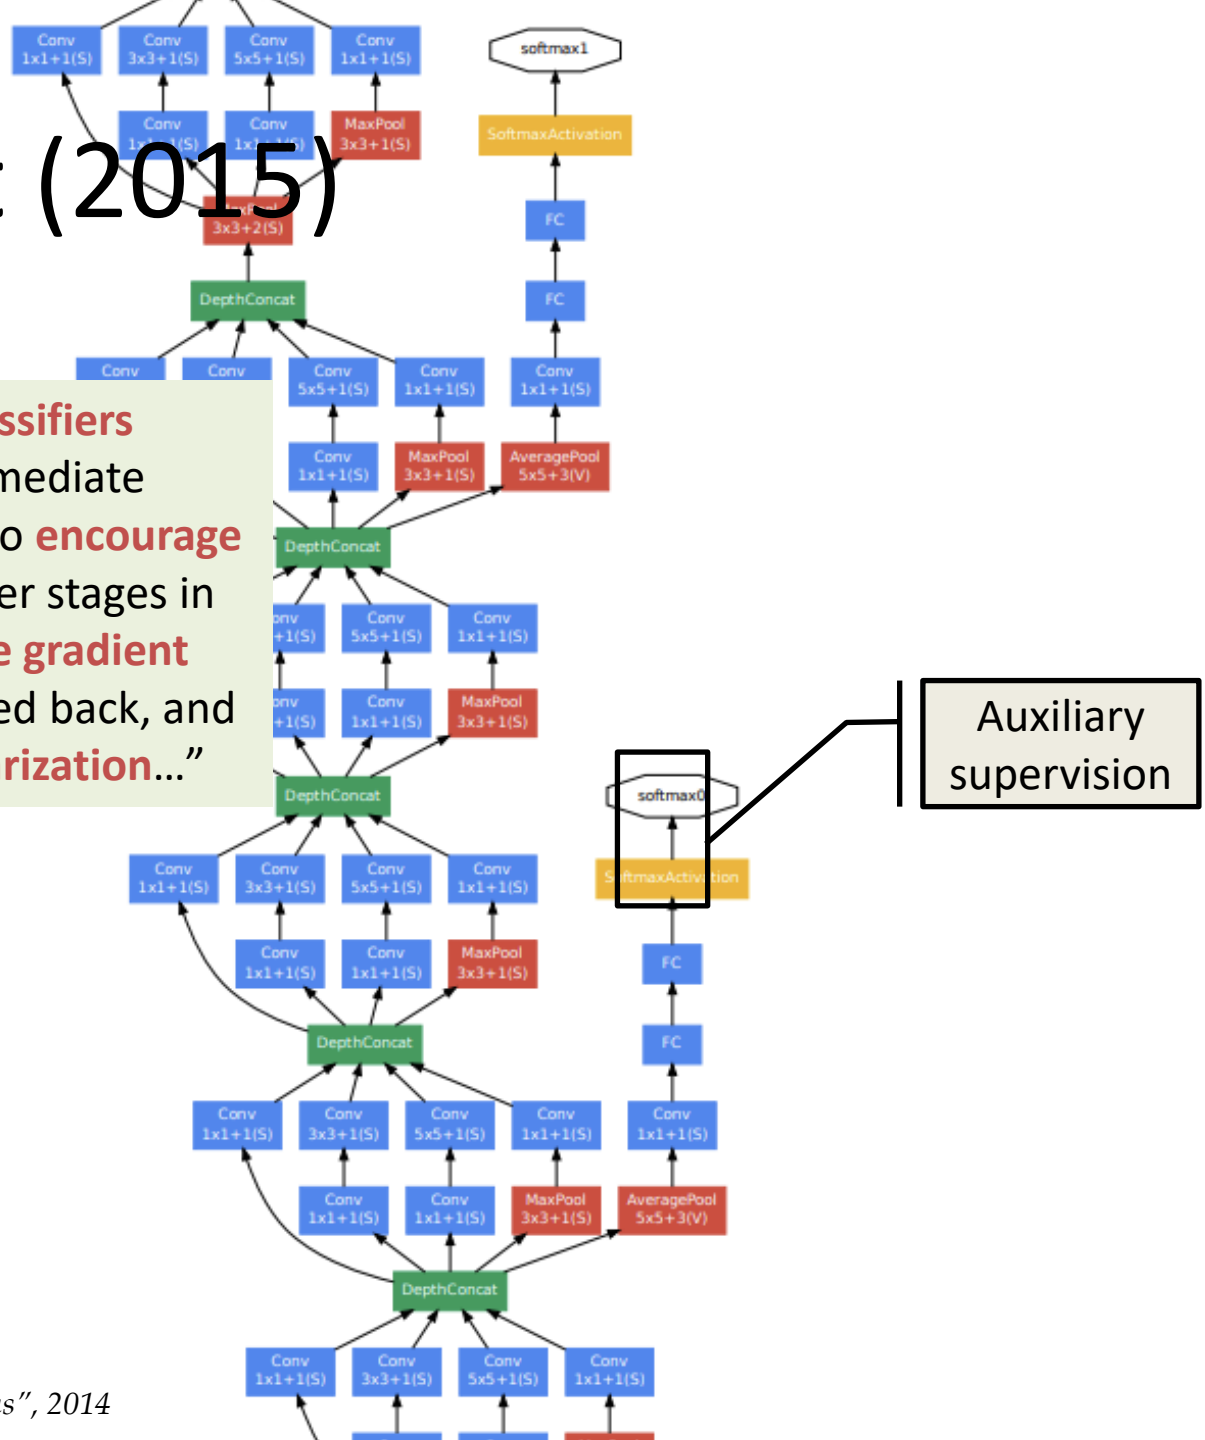


Reduce the number of channels, using 1×1 convolutions

GoogLeNet (2015)

“...By adding **auxiliary classifiers** connected to these intermediate layers, we would expect to **encourage discrimination** in the lower stages in the classifier, **increase the gradient signal** that gets propagated back, and provide **additional regularization**...”

- Convolution / FC
- Max / Avg Pooling
- Softmax
- Concatenation



GoogLeNet (2015)

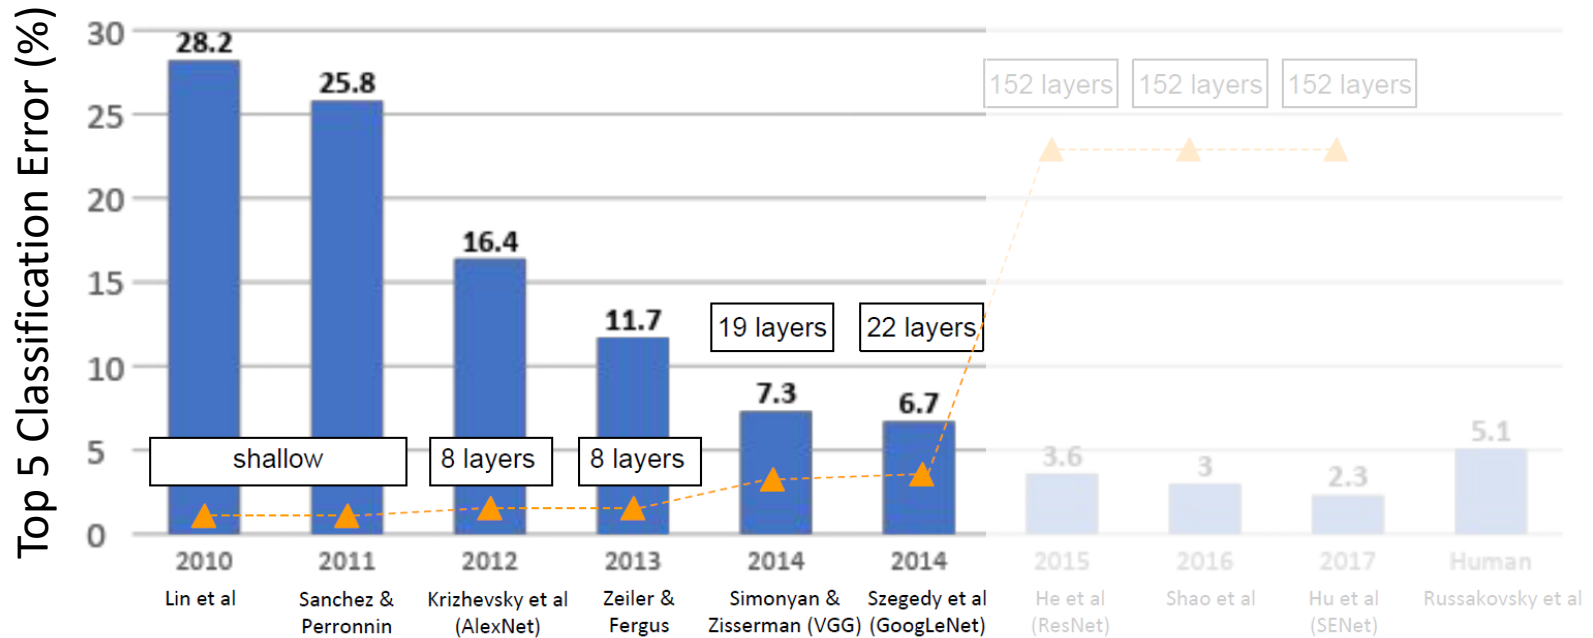
22-layer deep CNN, but reduced the number of parameters **from 60 million (AlexNet) to 4 million (12x)**

Low use of fully connected layers, they use average pooling instead

- Convolution / FC
- Max / Avg Pooling
- Softmax
- Concatenation



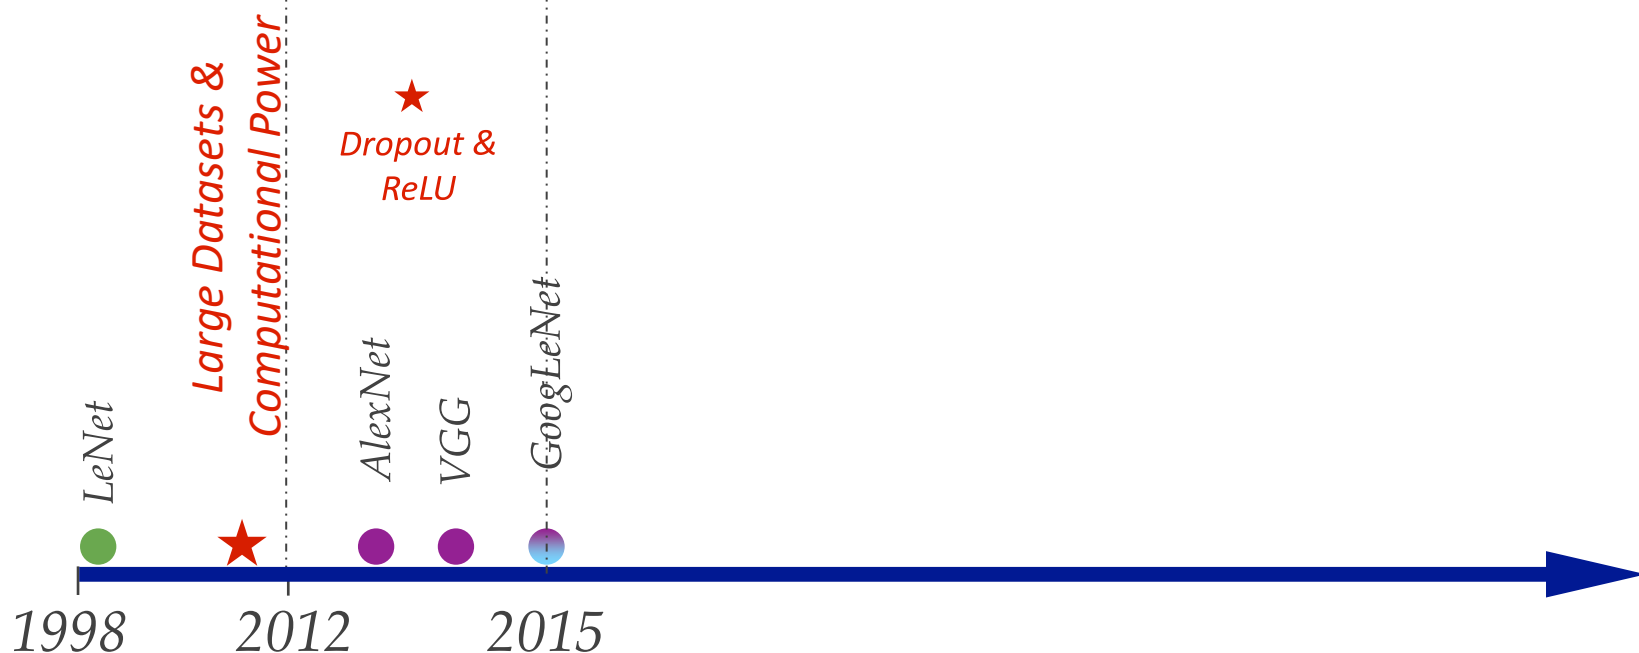
ImageNet Classification



Wrap Up

training of deep nets
unsuccessful (except
for CNNs)

time consuming,
high memory
demanding

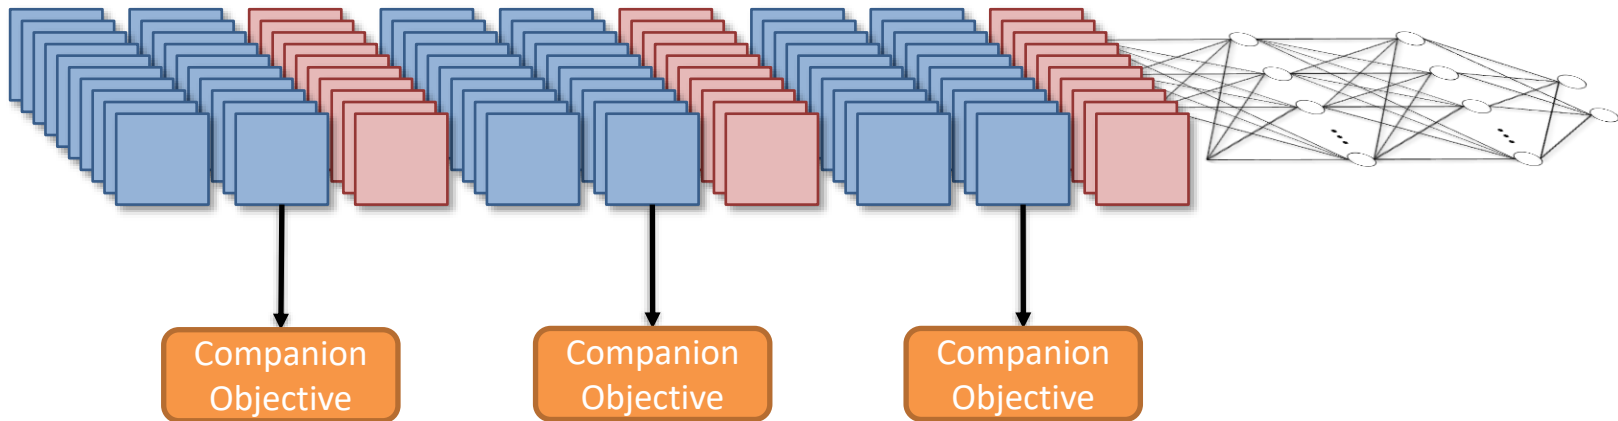


DSNs, FitNets, Highway Nets and ResNets

2015 - 2016

Deeply Supervised Networks

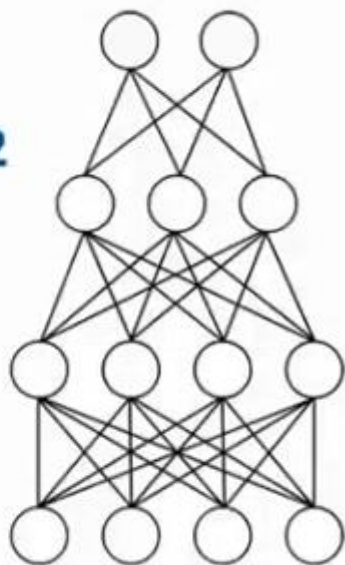
Adding intermediate supervision (discriminative loss)



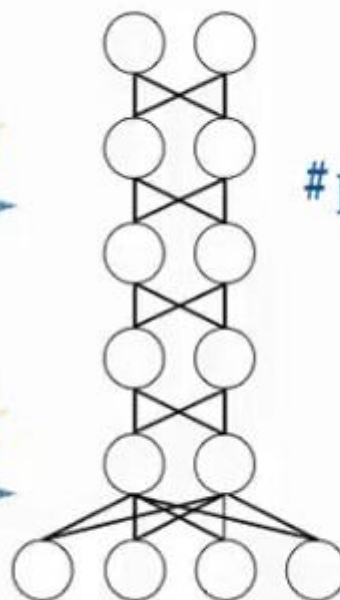


FitNets

Teacher
#parameters = 32



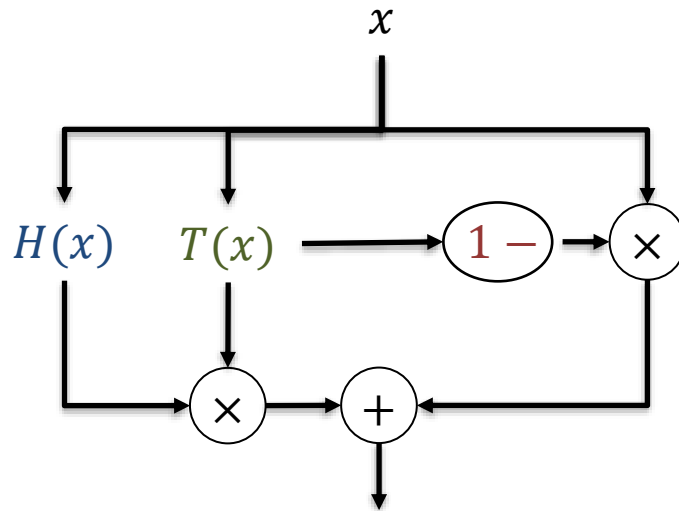
Student
#parameters = 24



Knowledge transfer: trade width for depth to reduce the number of parameters

Key intuition: bigger is not better, just easier to optimise! Over-parameterisation helps optimisation.

Highway Networks



$$\text{output} = H(x)T(x) + x(1 - T(x))$$

x : Input

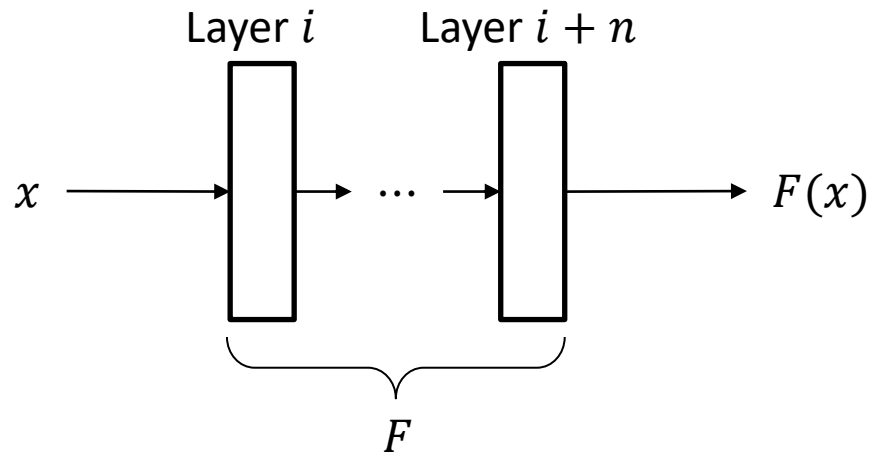
$H(x)$: Layer transform
(a block of operations)

$T(x)$: Transform gate
(how much we transform the input)

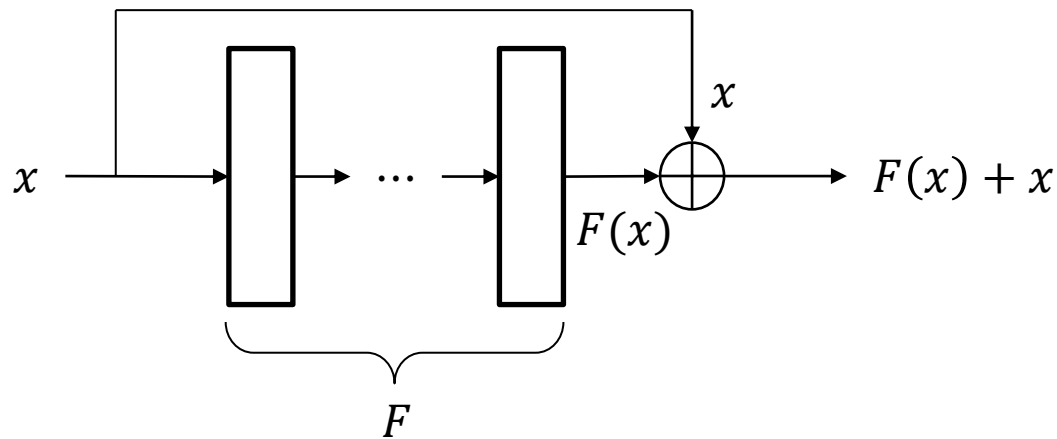
$1 - T(x)$: Carry gate
(how much we preserve the input)

Key idea: provide an alternative, direct path for the gradient to flow back

Residual connections



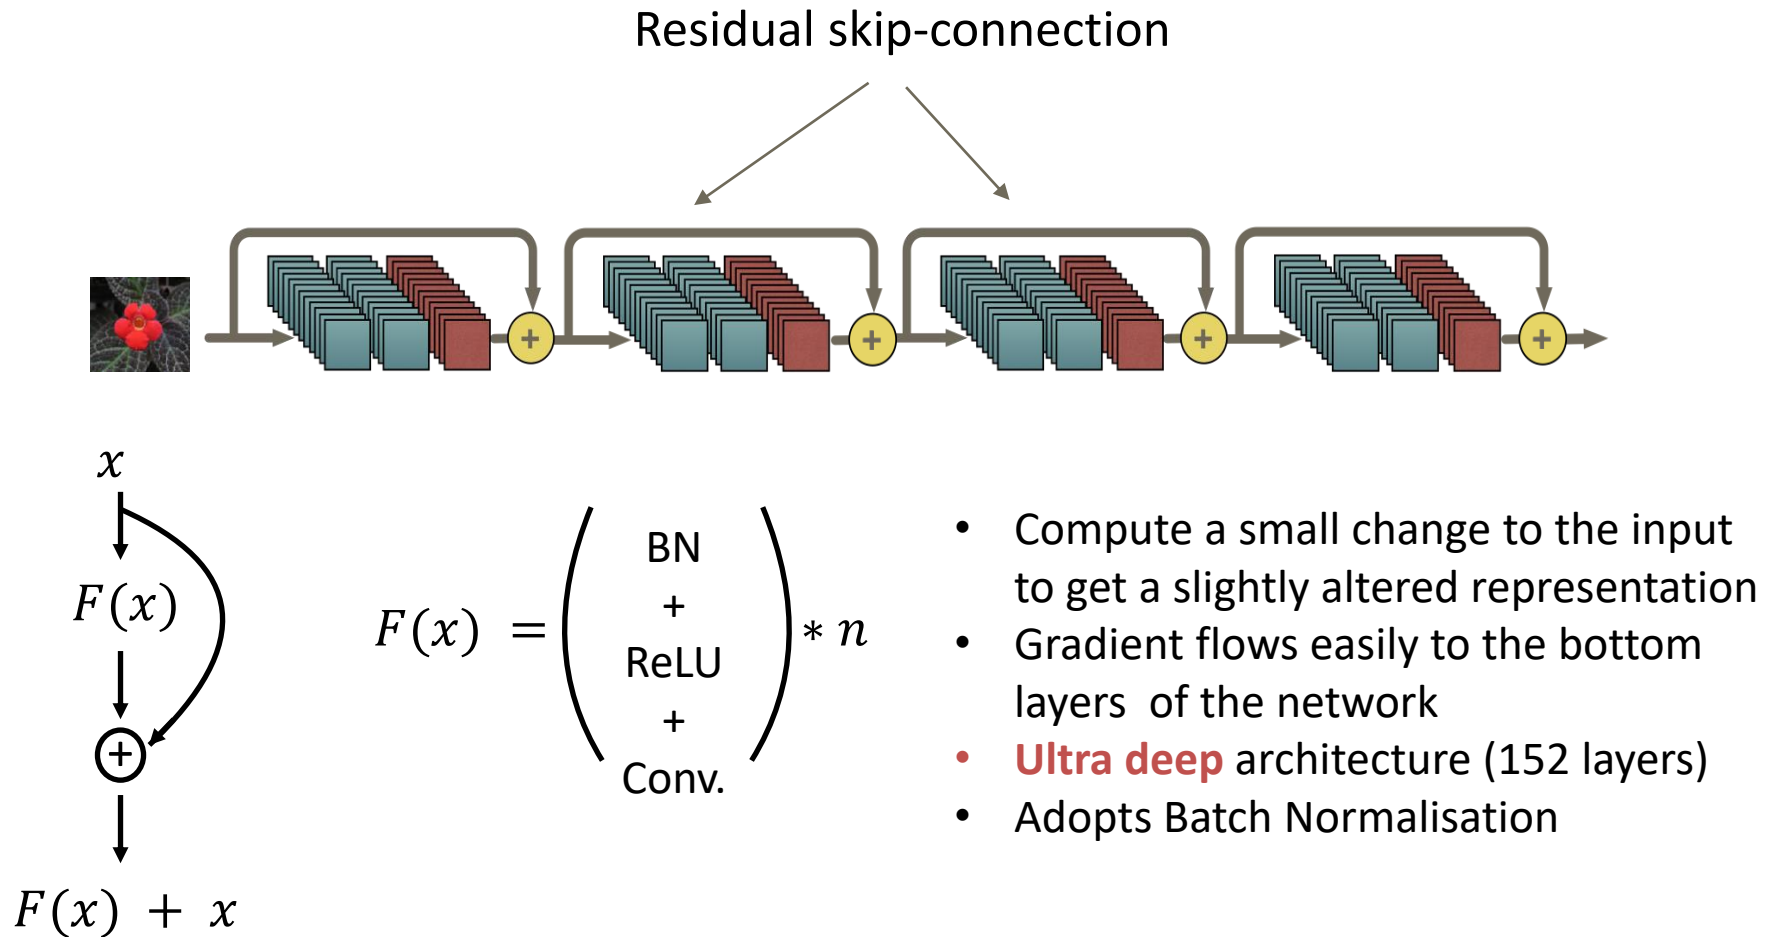
Traditional feedforward



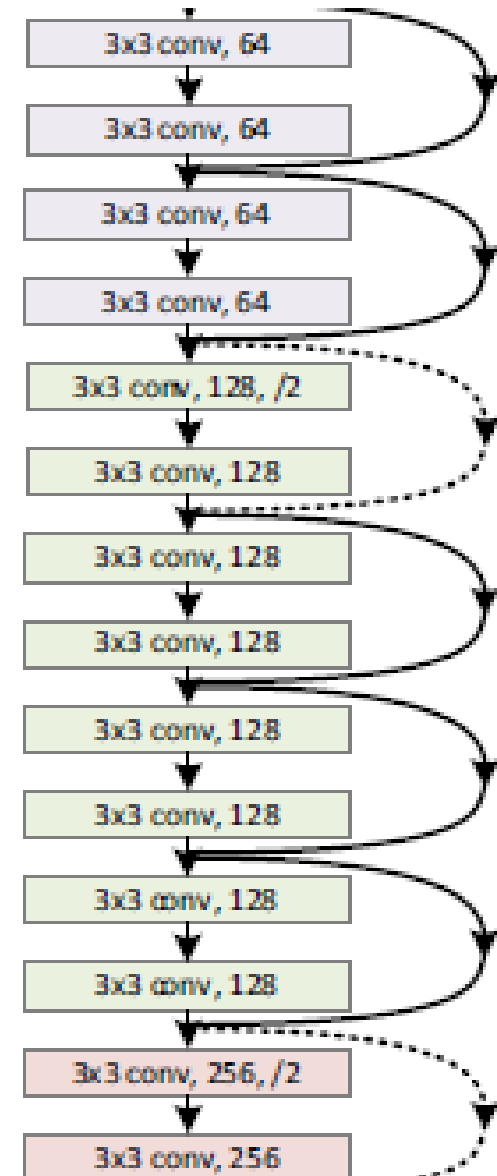
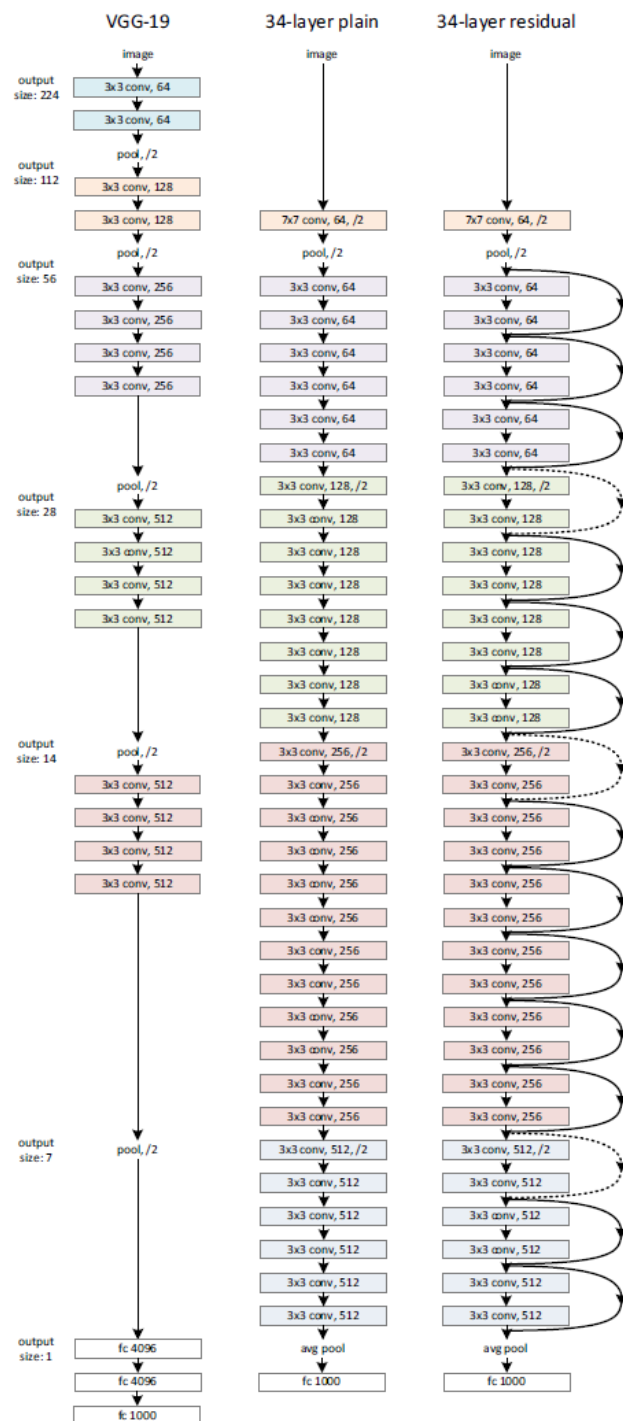
Feedforward with
residual connection

One way to avoid vanishing gradient is to provide alternative routes for your gradients to flow through, a.k.a. residual connections

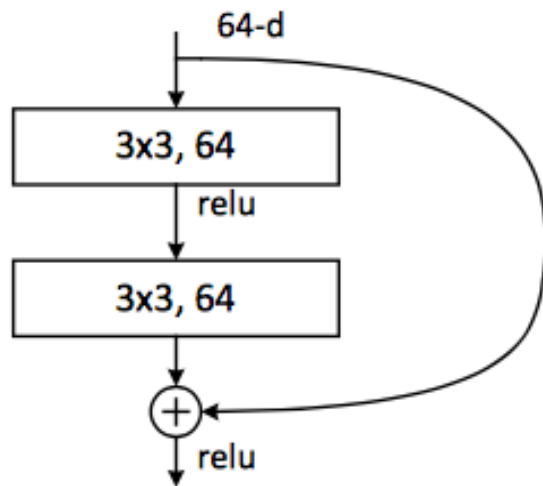
Residual Networks



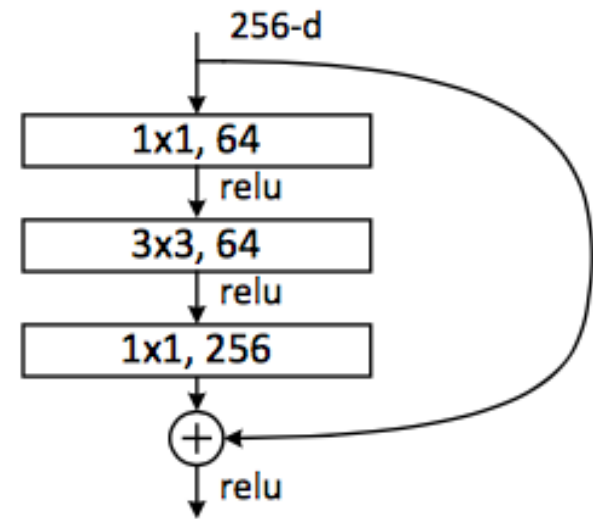
Residual Networks



Residual Blocks



Basic Block



Bottleneck block

Residual Networks

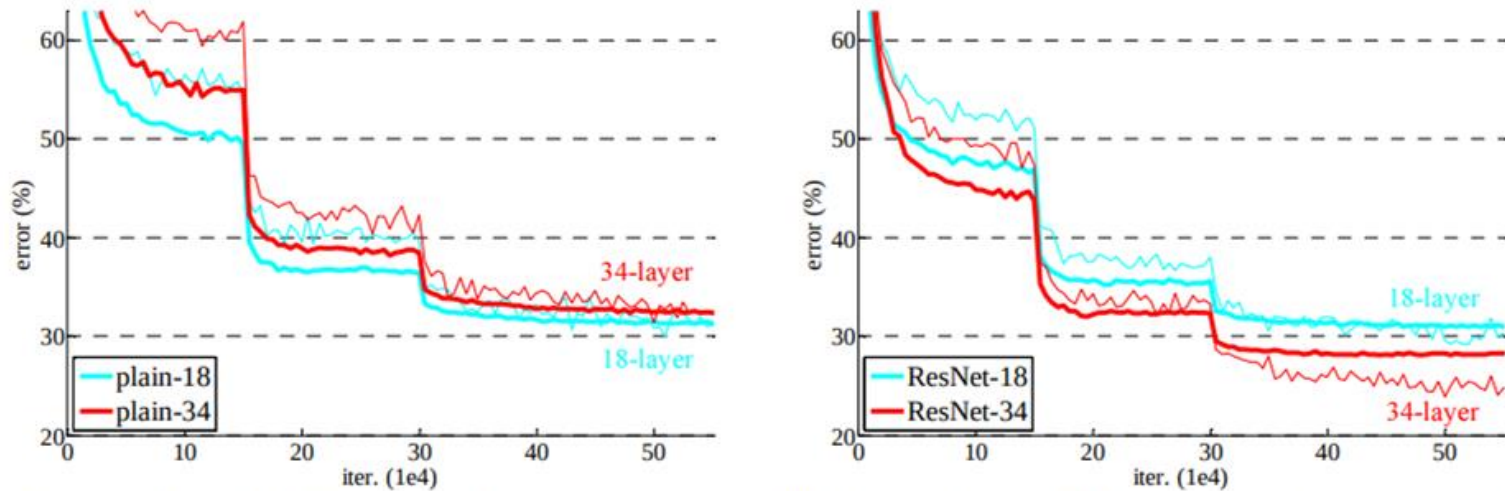
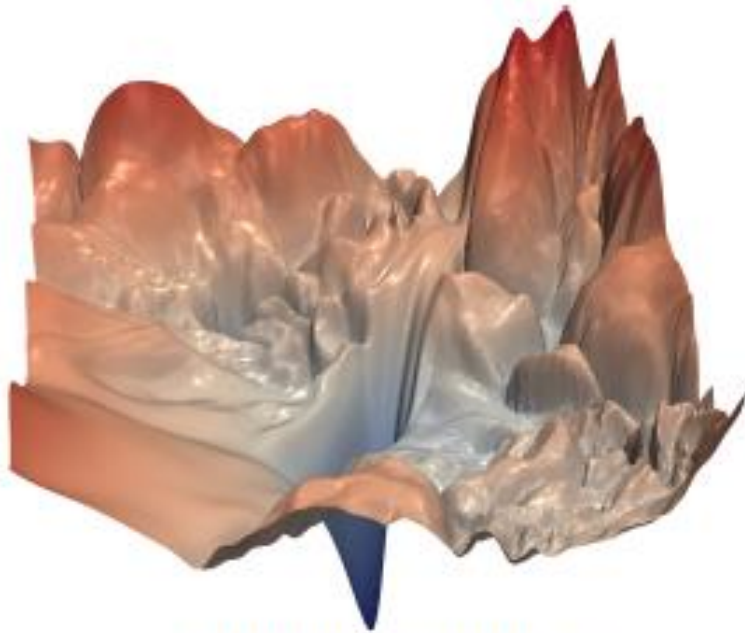
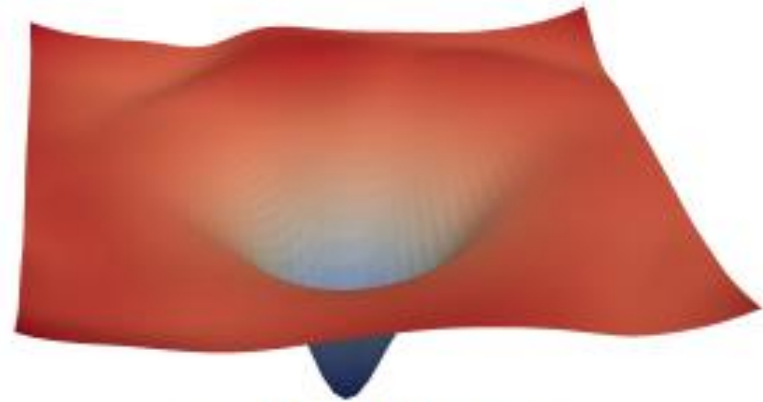


Figure 4. Training on **ImageNet**. Thin curves denote training error, and bold curves denote validation error of the center crops. Left: plain networks of 18 and 34 layers. Right: ResNets of 18 and 34 layers. In this plot, the residual networks have no extra parameter compared to their plain counterparts.

Residual Networks



(a) without skip connections



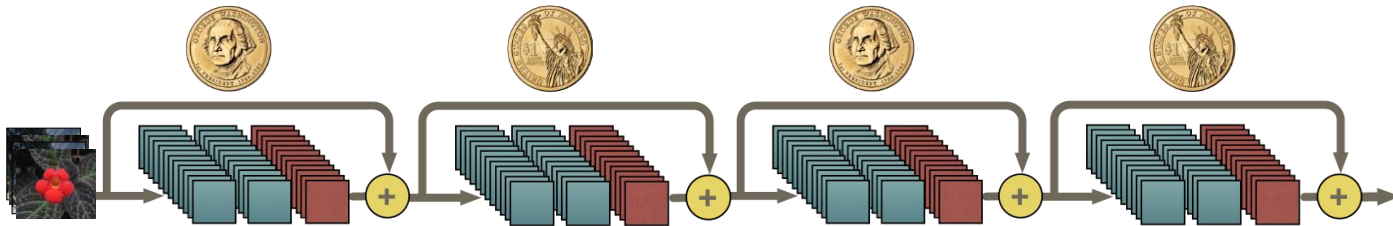
(b) with skip connections

Figure 1: The loss surfaces of ResNet-56 with/without skip connections. The proposed filter normalization scheme is used to enable comparisons of sharpness/flatness between the two figures.

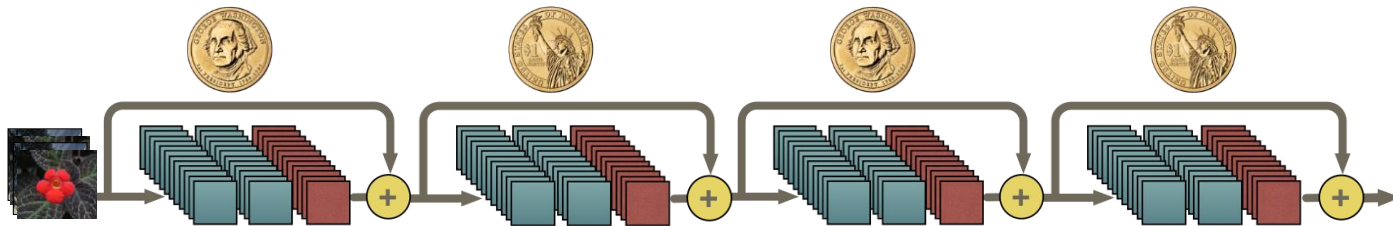
Stochastic depth, Dense Nets

2015 - 2016

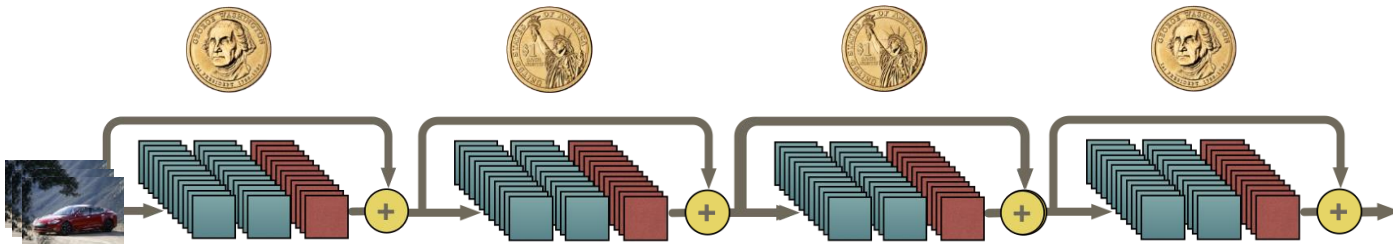
Stochastic Depth



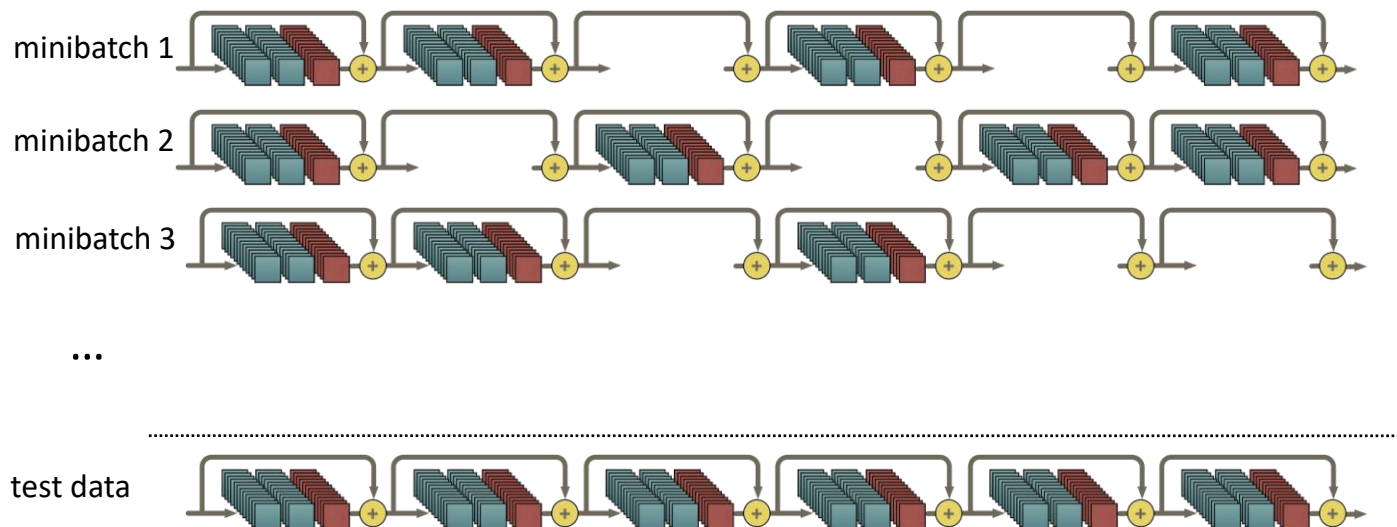
Stochastic Depth



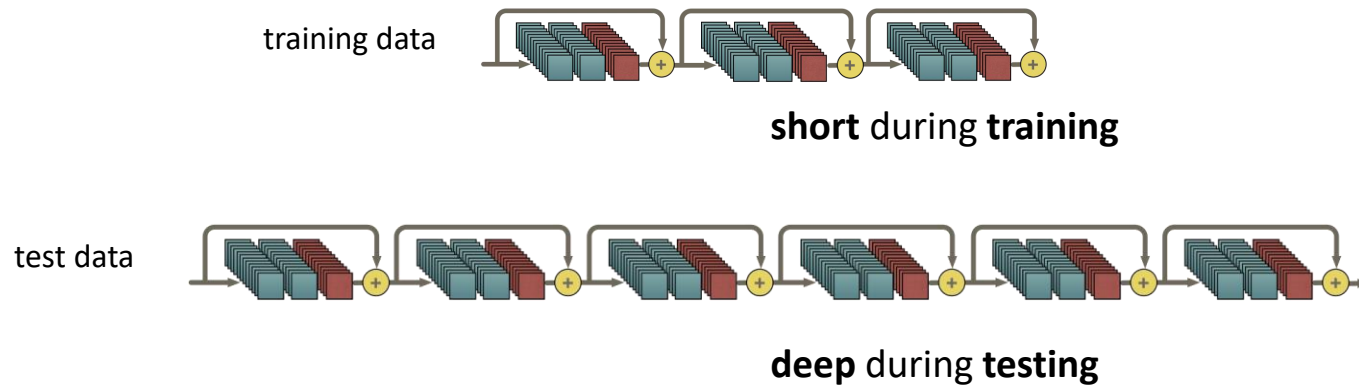
Stochastic Depth



Stochastic Depth



Stochastic Depth

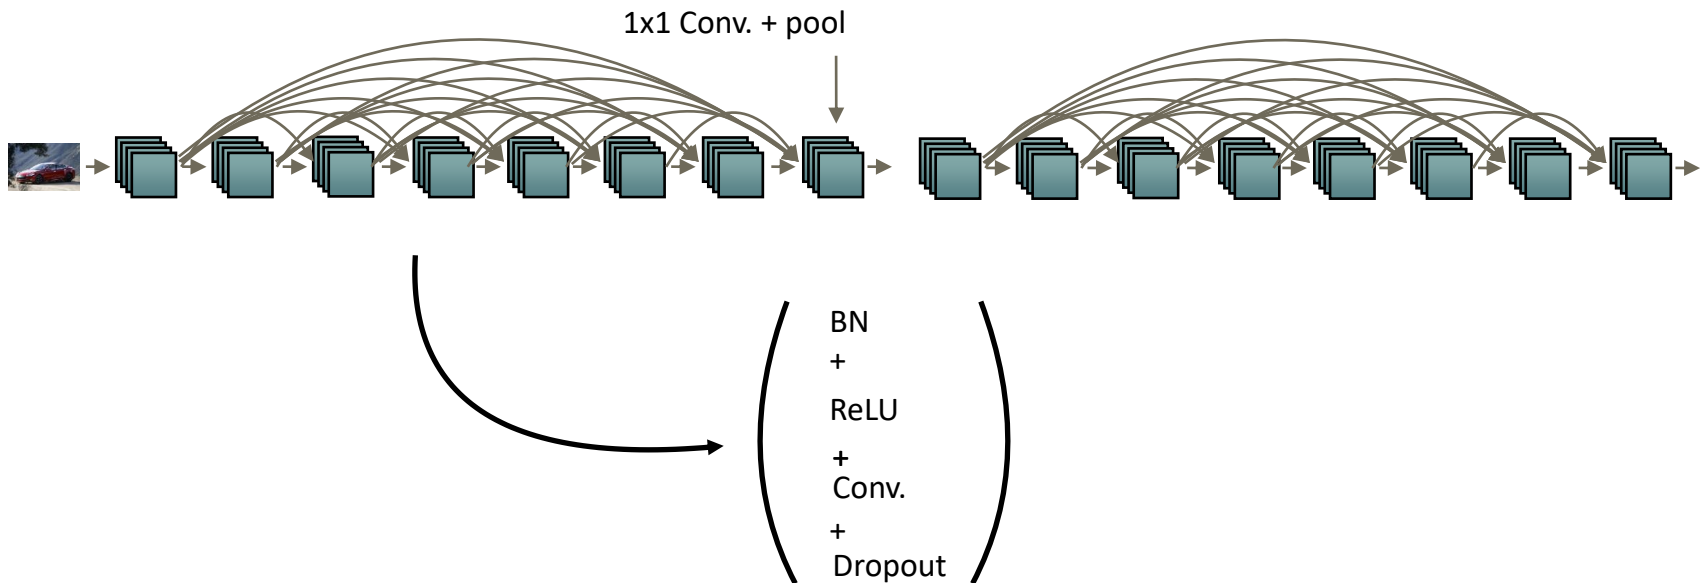


Key aspects:

- Implicit **ensemble** of 2^L models
- Improved **gradient flow**
- 25% speedup during training
- Lower error

DenseNet

Connect every layer to every other layer of the same filter size (dense blocks).

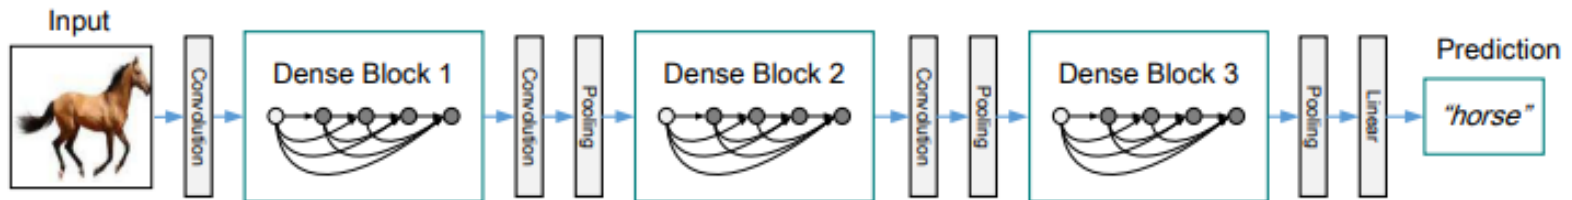


DenseNet

Example Dense Block (4 layers, growth rate 4):

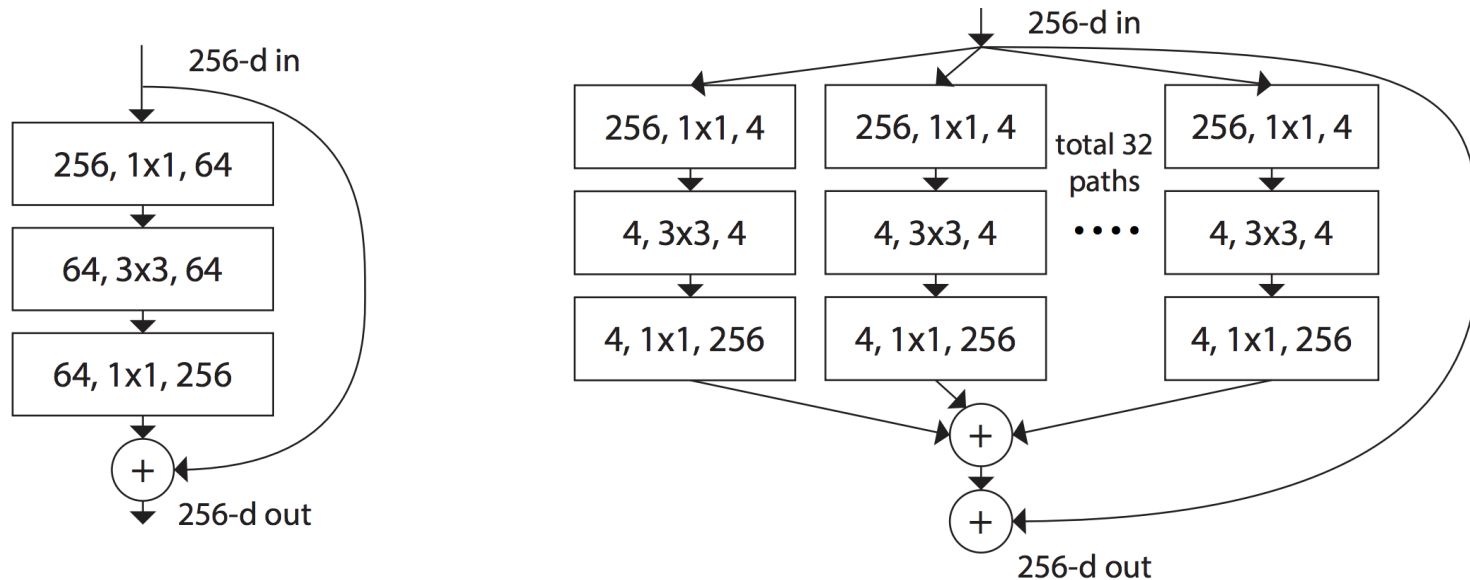


DenseNet can have very narrow layers (**small growth rate**). One explanation for this is that each layer has access to all the preceding feature-maps in its block. Each layer then adds a little bit of information (k feature-maps) to the global state of the network



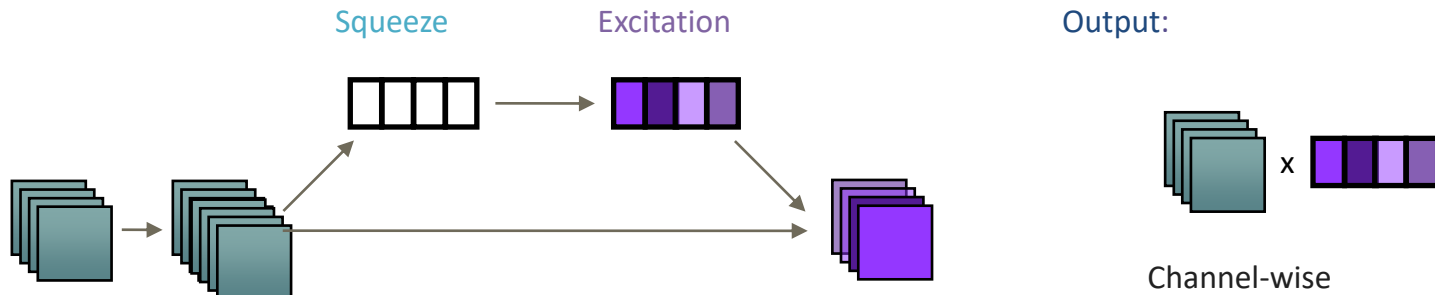
ResNeXt

ResNeXt introduces the concept of “cardinality”. Apart from depth (how many layers) and width (how many filters).



Key concept: create a lot of paths with a very small number of channels using 1×1 convolutions

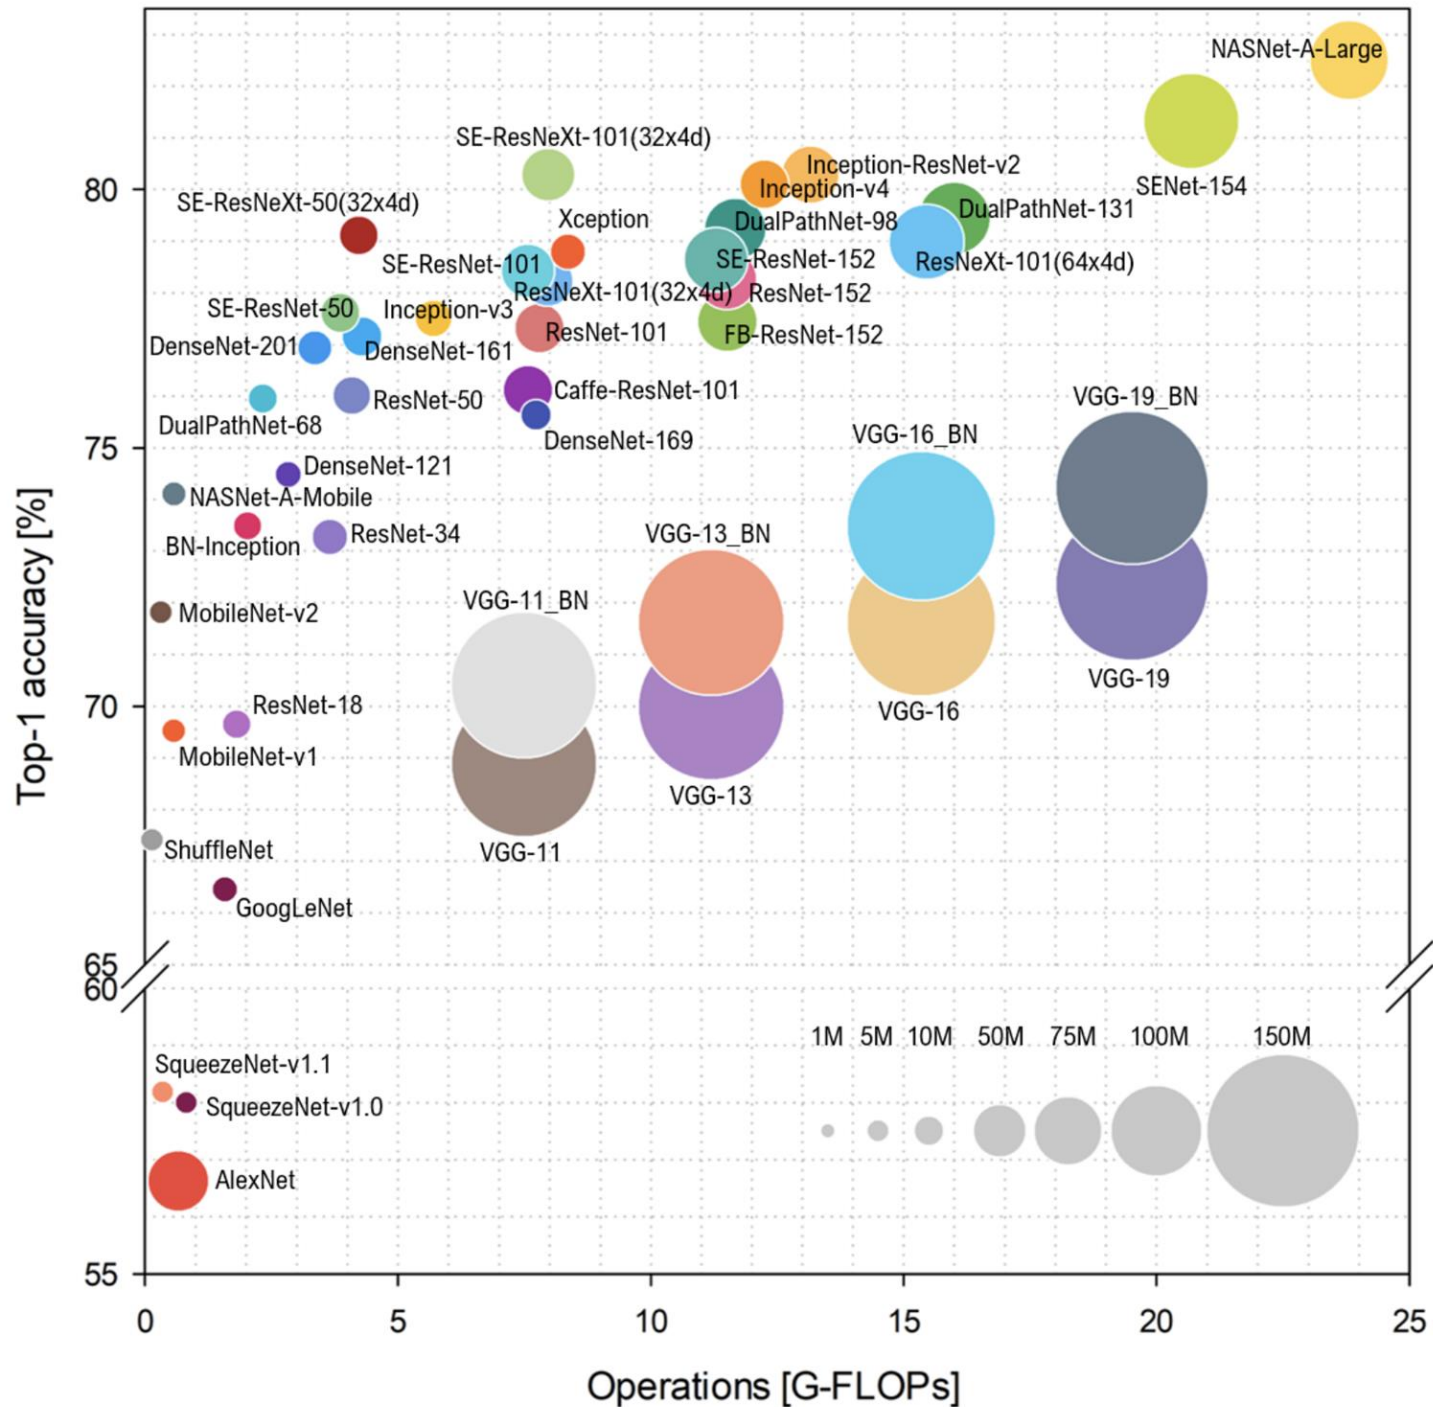
Squeeze-&-Excitation Networks



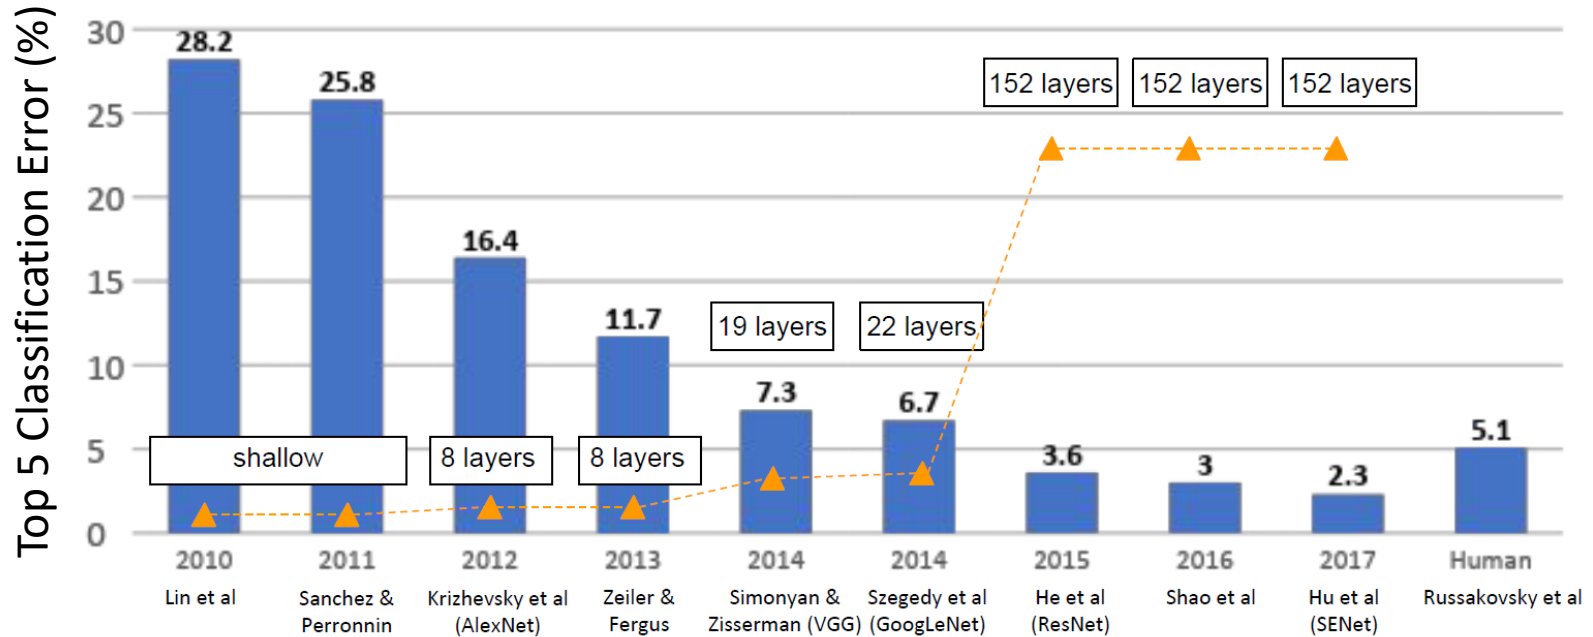
Squeeze: average pooling - captures channel-wise global info

Excitation: non-linear channel interaction + gating (sigmoid)

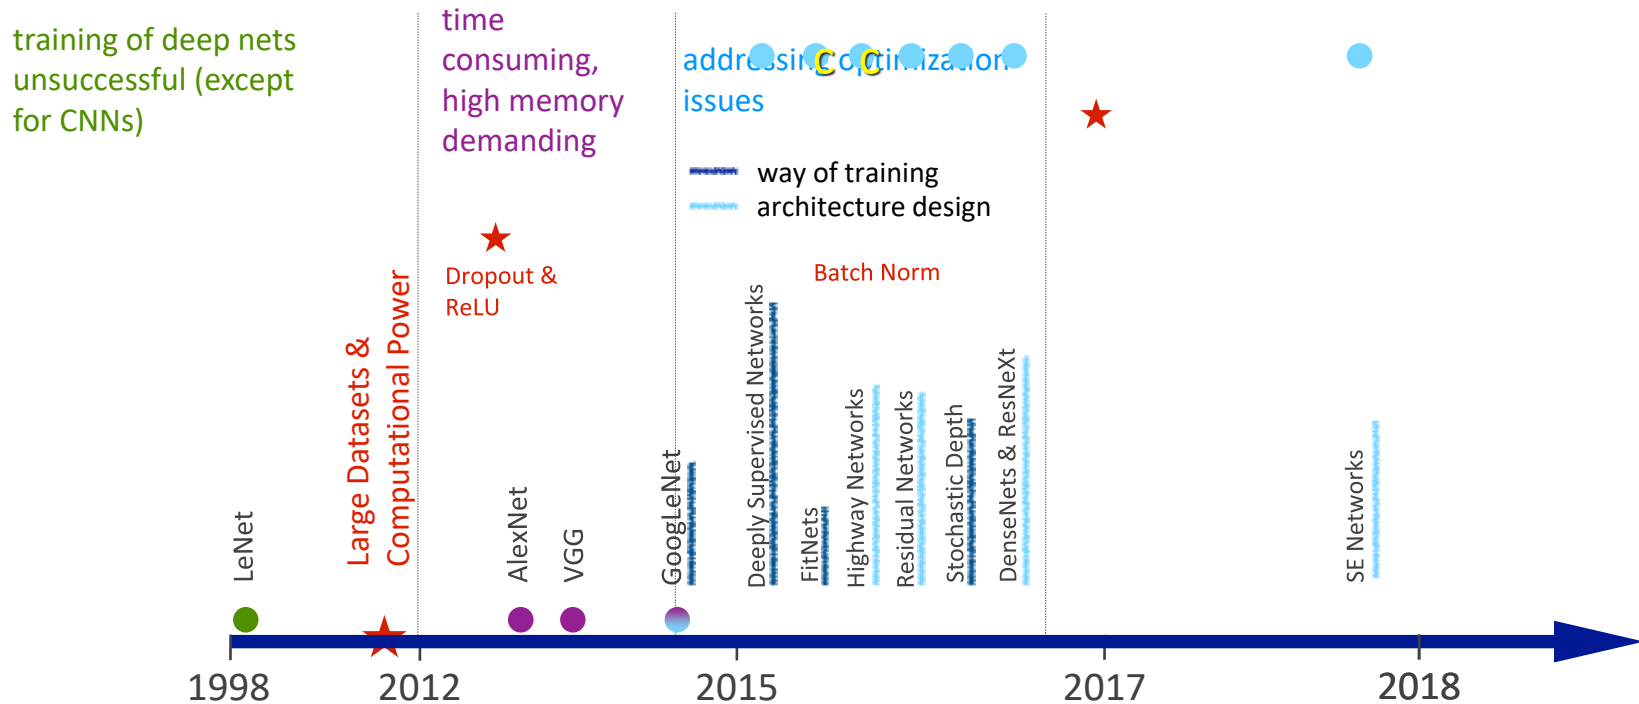
Output: scales activations, leveraging global information at all levels



ImageNet Classification



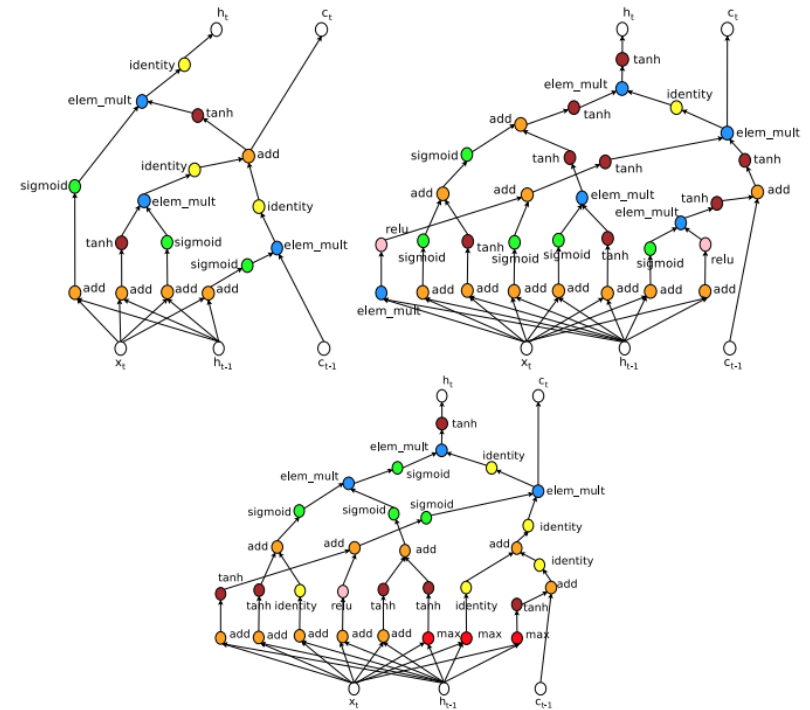
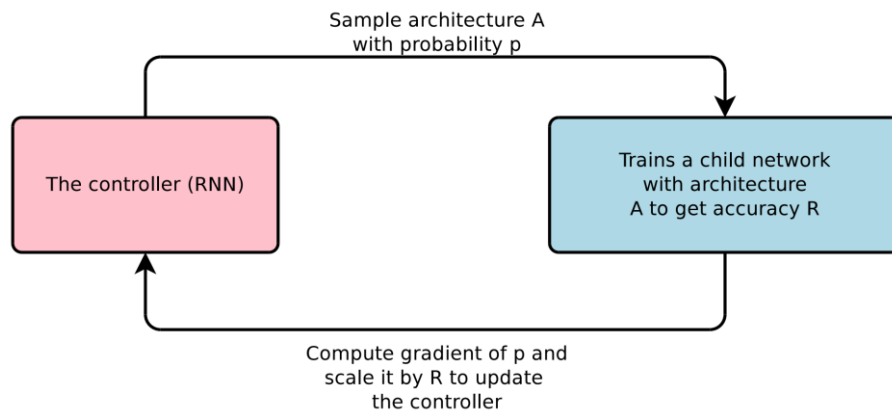
Wrap Up



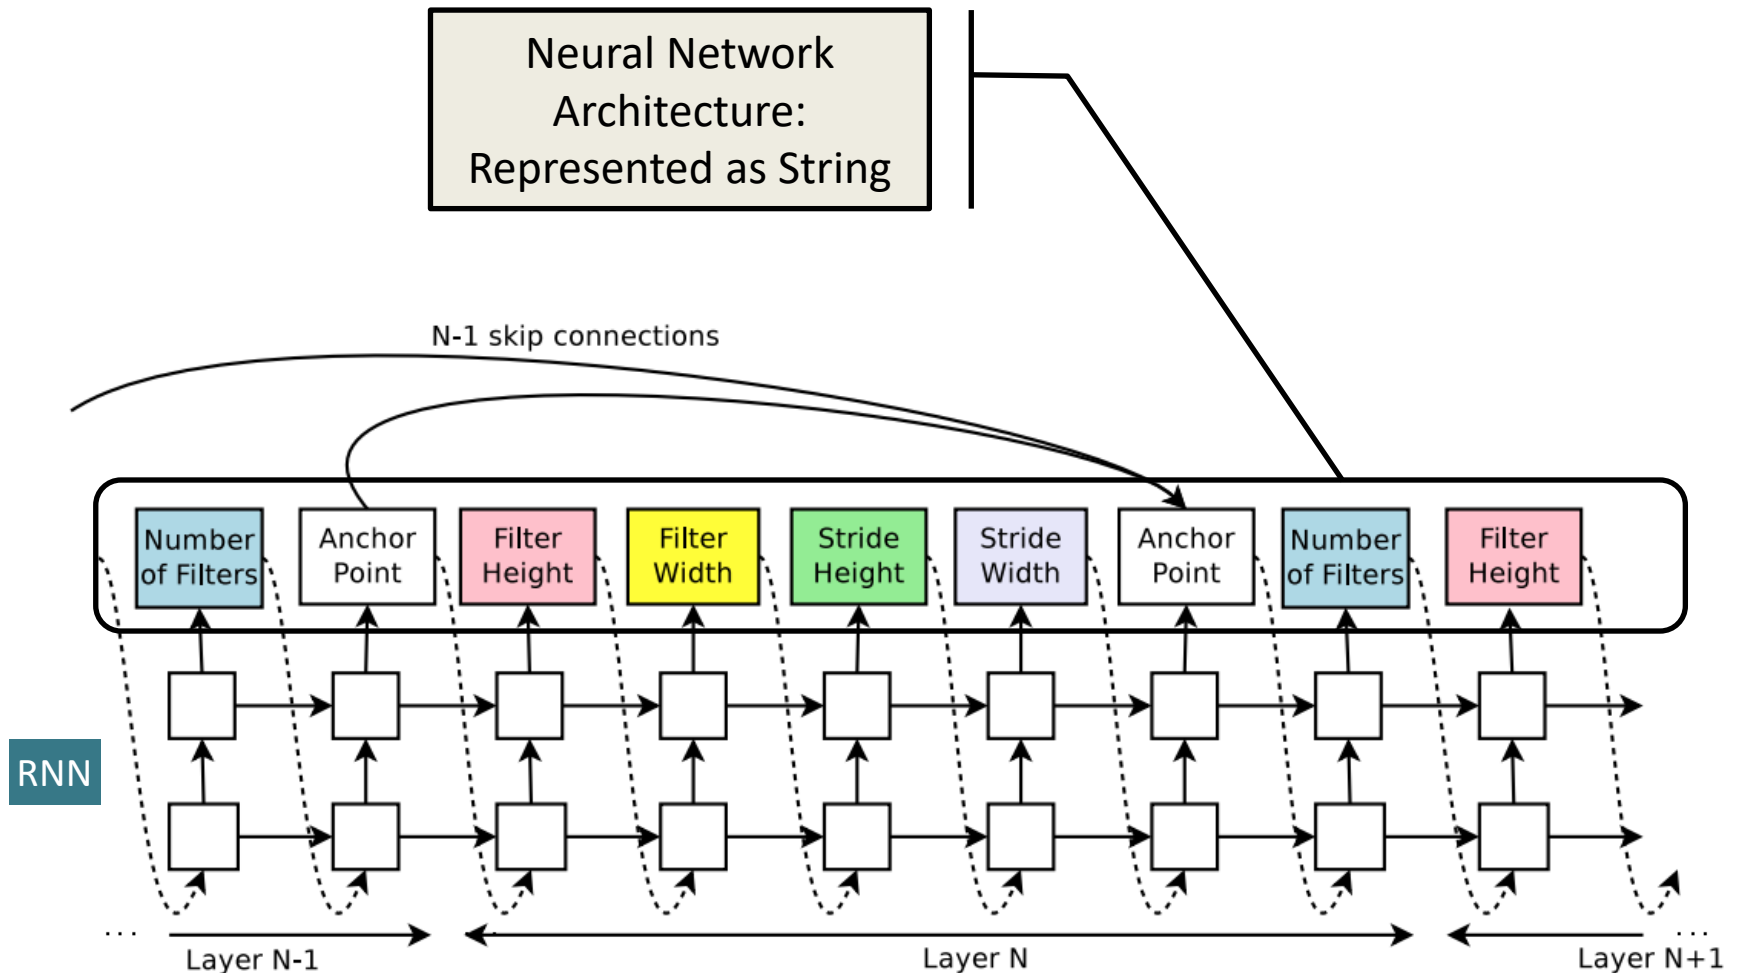
Neural Architecture Search

2018 - 2019

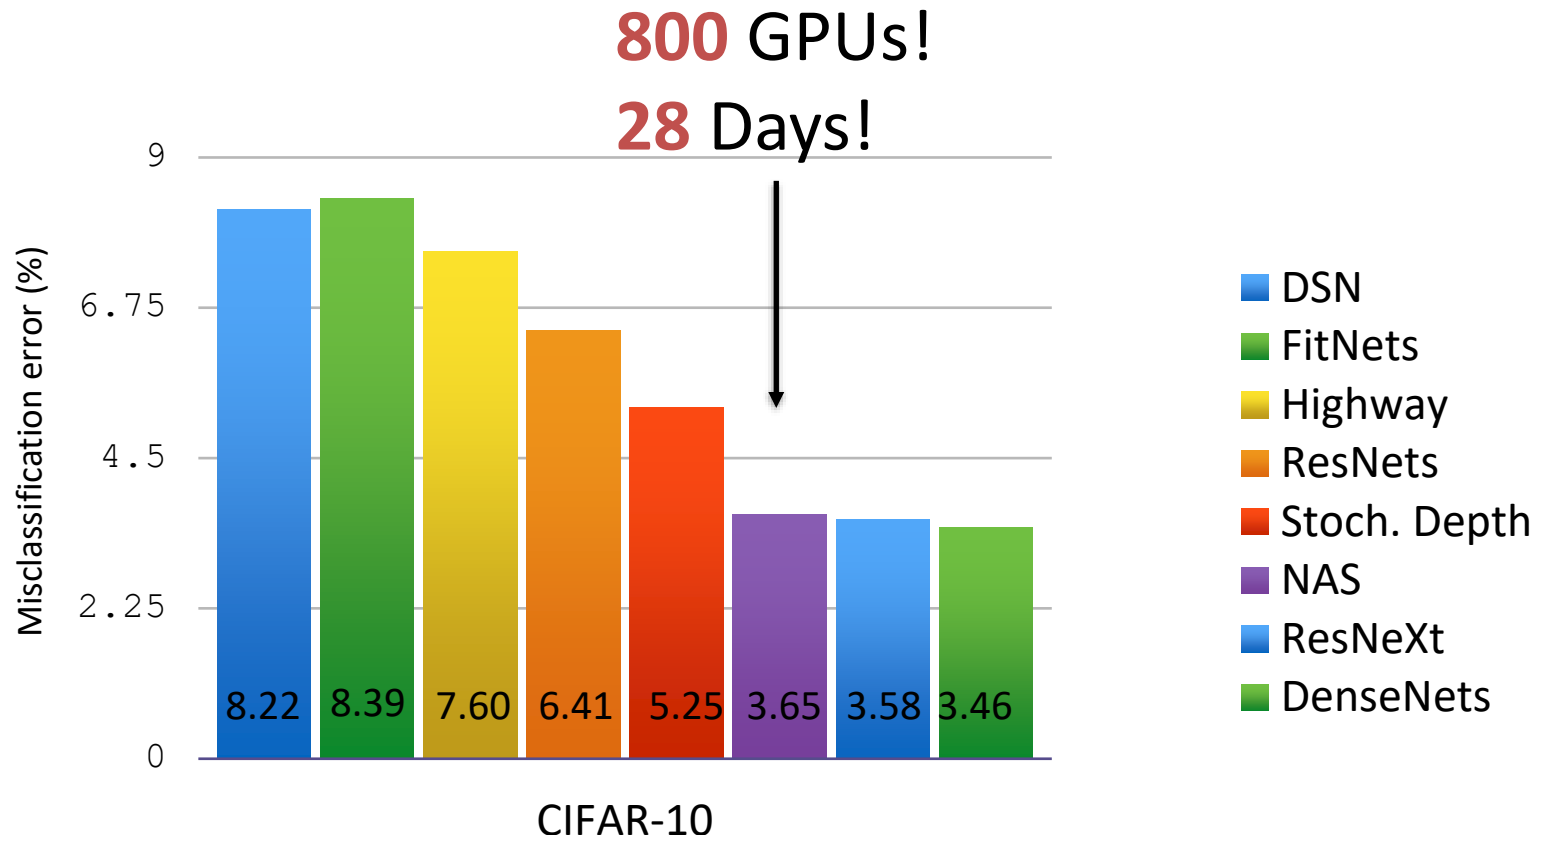
Neural Architecture Search (NAS)



Neural Architecture Search (NAS)

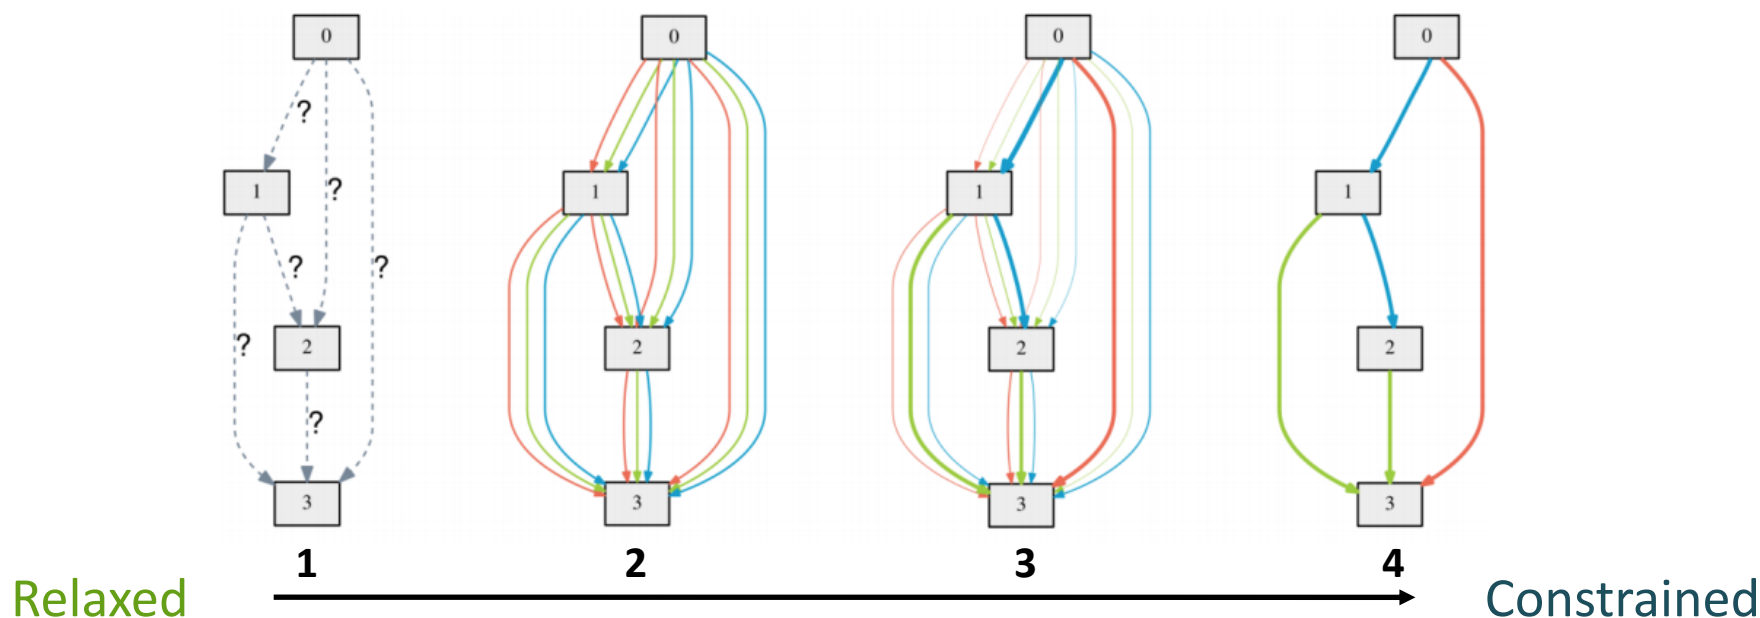


CIFAR-10 Results

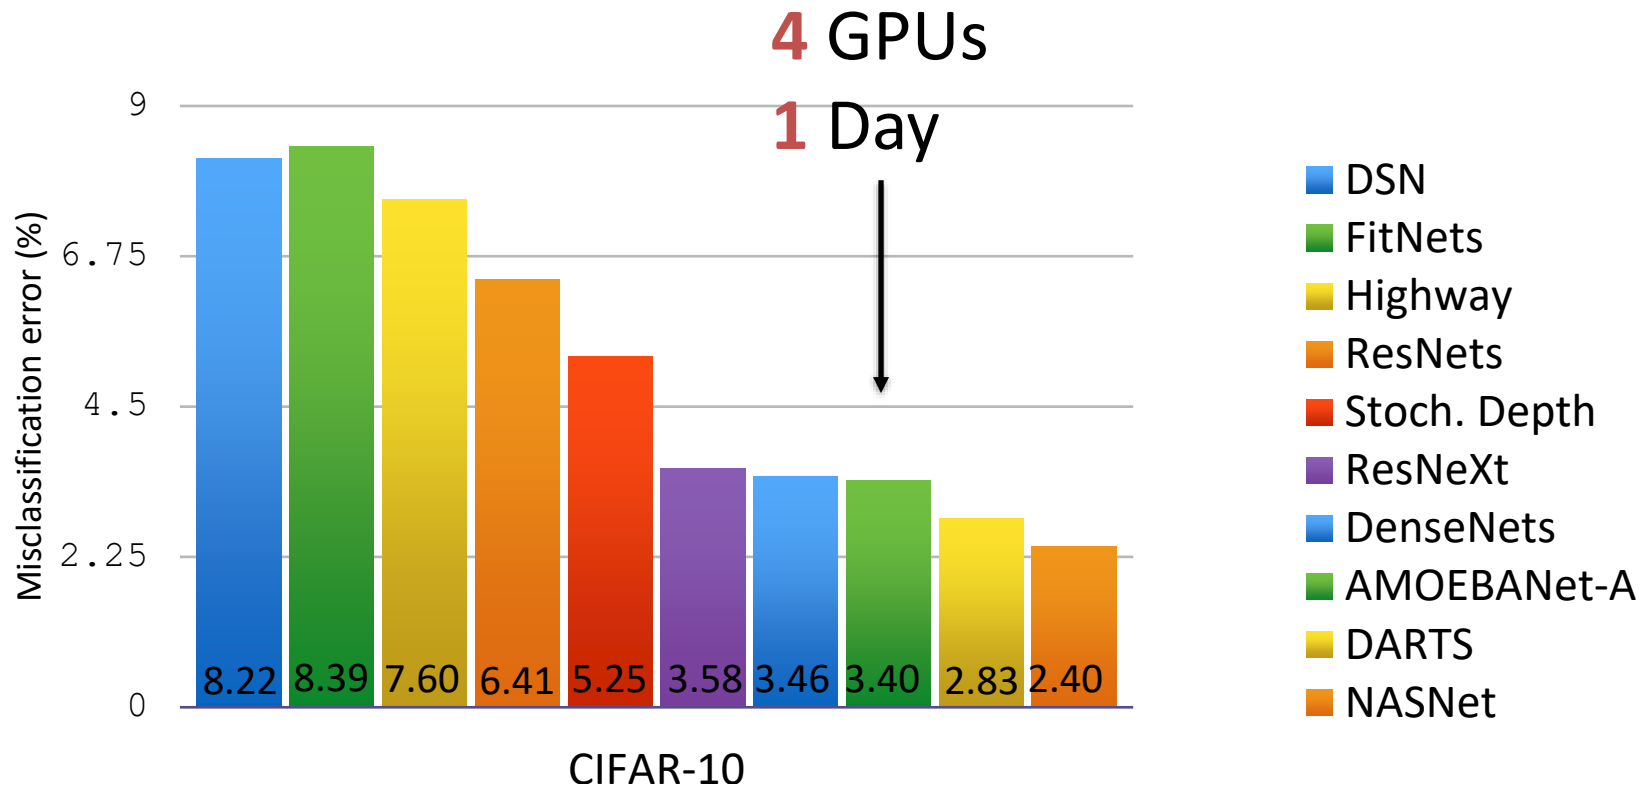


DARTS: Differentiable Architecture Search

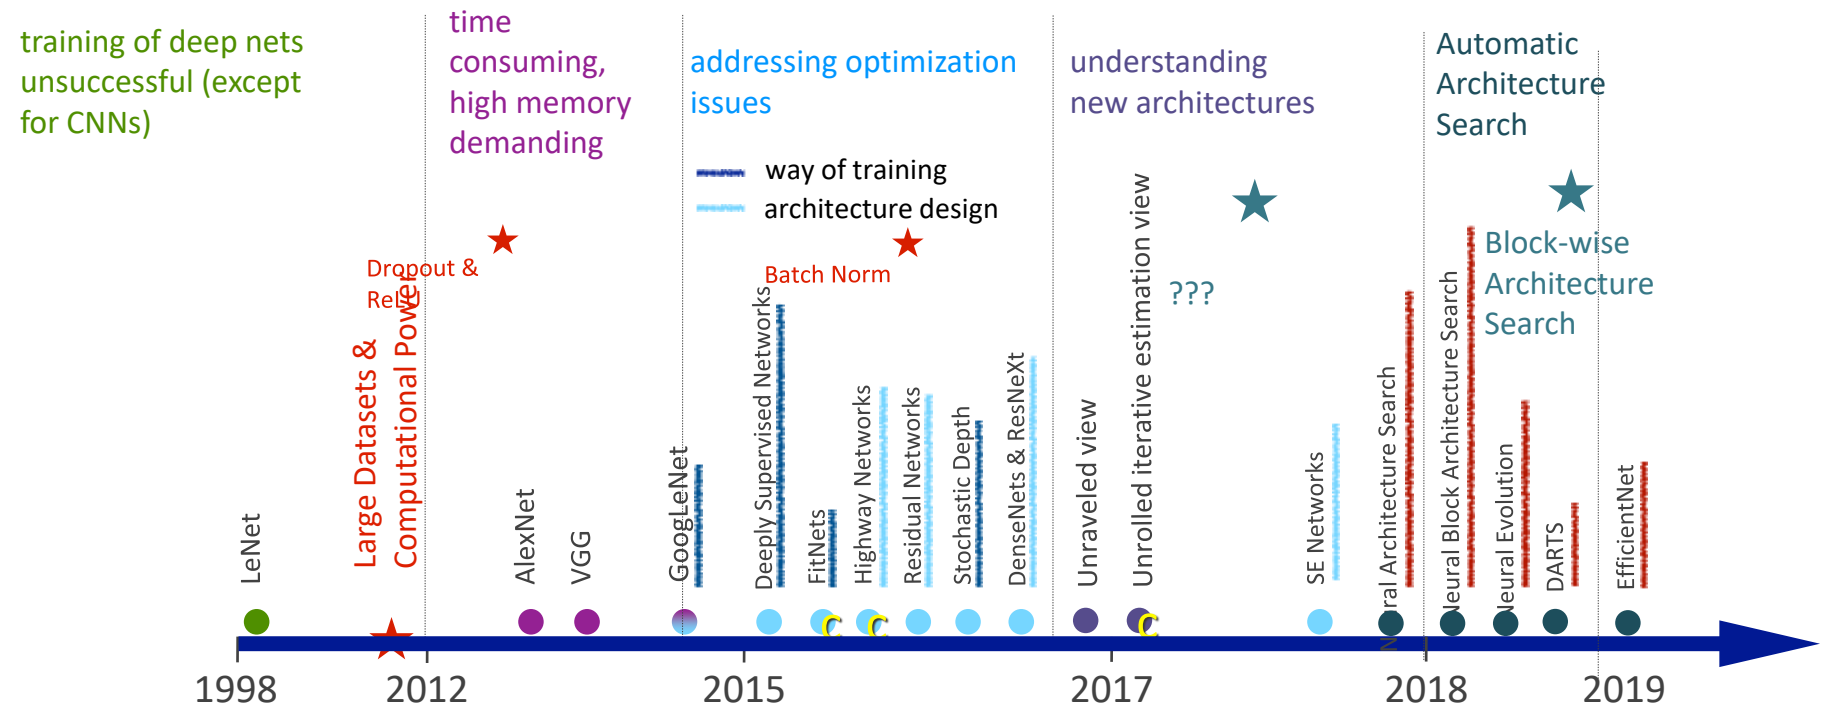
Key concept: Start with all operations connected. Then slowly prune connections that are not being used. Completely remove them at the end for test.



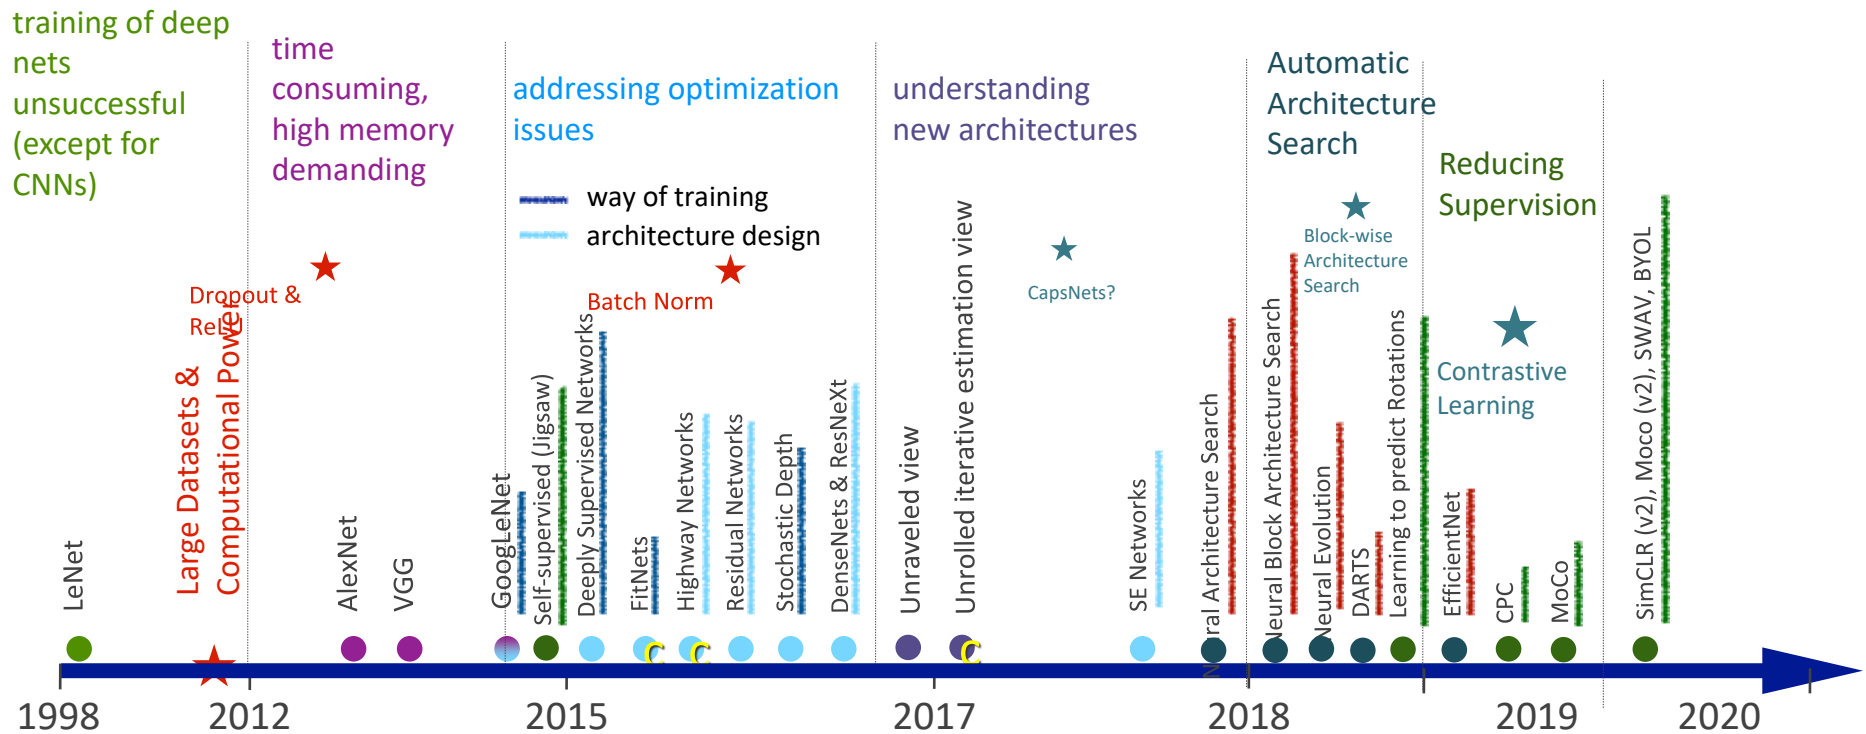
CIFAR-10 Results



Wrap Up

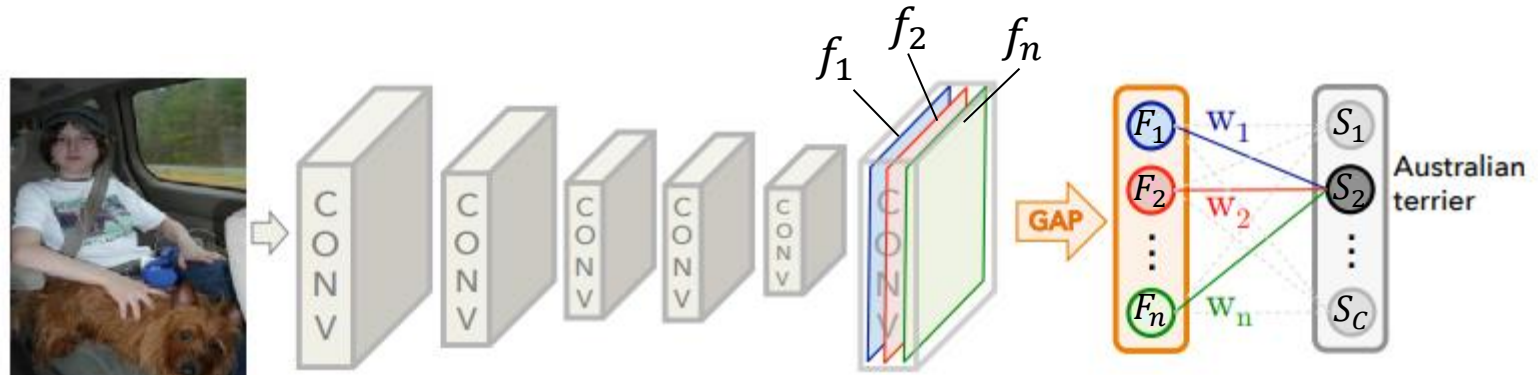


Wrap Up



GETTING INSIGHTS

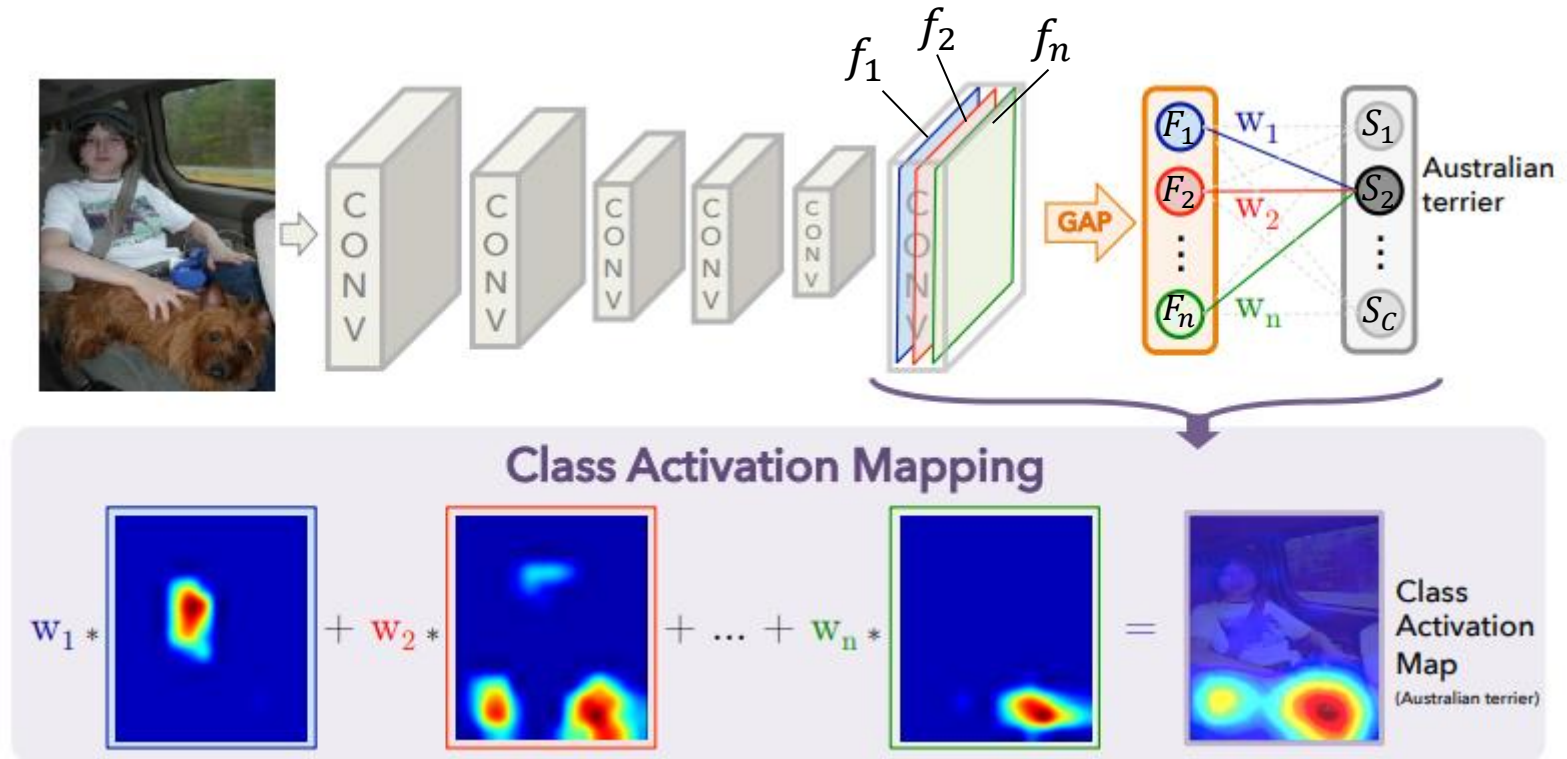
Class Activation Mapping



$$F_k = \sum_{x,y} f_k(x, y)$$

$$S_c = \sum_k w_k^c \sum_{x,y} f_k(x, y)$$

Class Activation Mapping



$$M_c(x, y) = \sum_k w_k^c f_k(x, y)$$

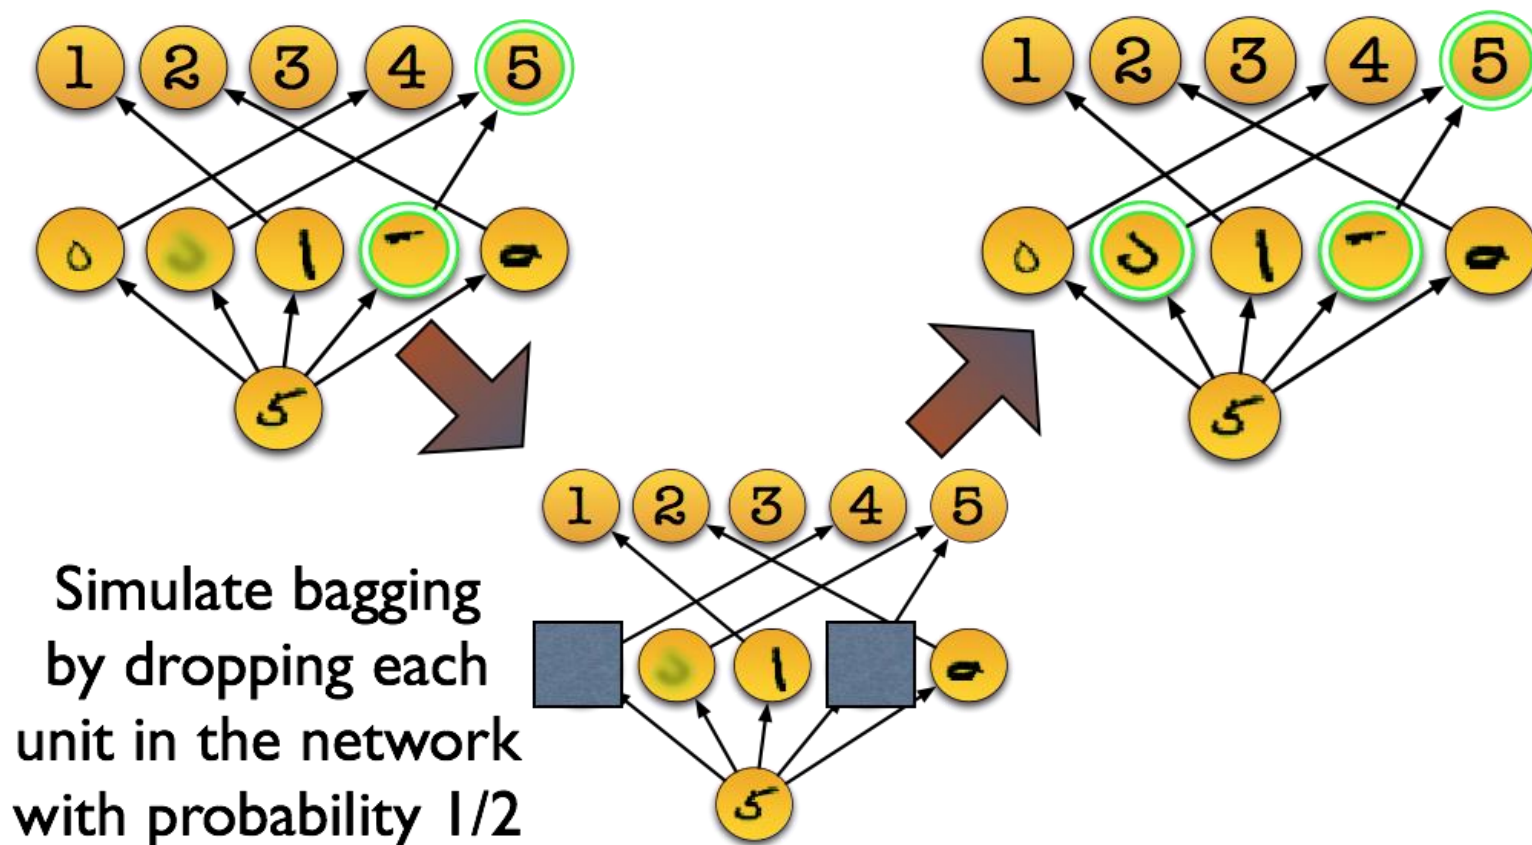
Class Activation Mapping



Seen in previous classes

RECAP OF IMPORTANT IDEAS

Dropout



Dropout in PyTorch

```
torch.nn.Dropout(p=0.5, inplace=False)
```

<https://pytorch.org/docs/stable/generated/torch.nn.Dropout.html>

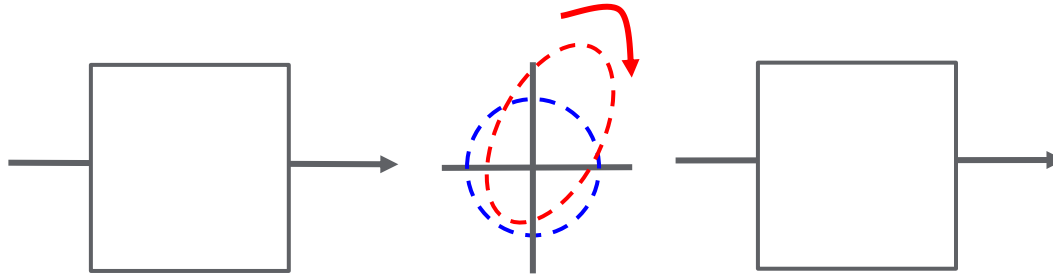


<https://losslandscape.com/>

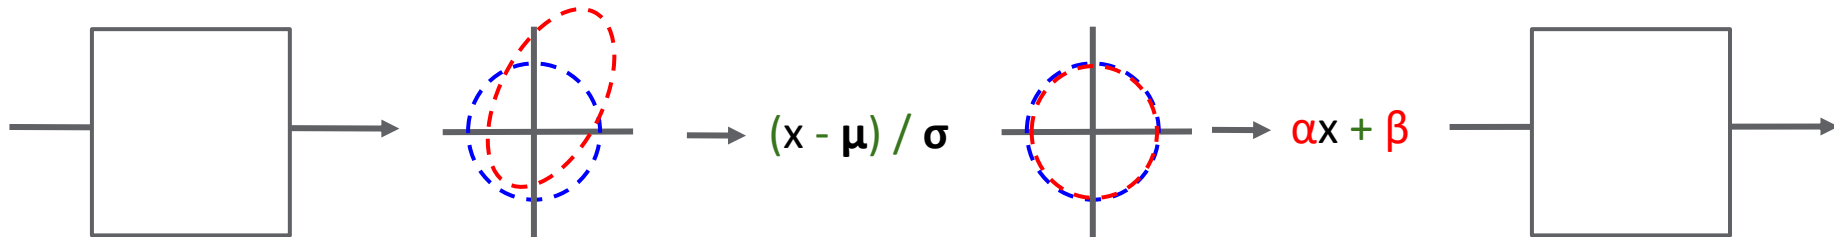
https://youtu.be/2PqTW_p1fls

Batch Normalization

Before:



After:



Batch Normalization in PyTorch

```
torch.nn.BatchNorm1d(num_features, eps=1e-05, momentum=0.1,  
                      affine=True, track_running_stats=True)
```

```
torch.nn.BatchNorm2d(num_features, eps=1e-05, momentum=0.1,  
                      affine=True, track_running_stats=True)
```

<https://pytorch.org/docs/stable/nn.html#normalization-layers>

Data Augmentation

Best way to make a machine learning model generalize better is to train it on more data -> **Create new data!**



Image source: [Building powerful image classification models using very little data](#)

Data Augmentation in PyTorch

```
From torchvision import transforms

rgb_mean = (0.4914, 0.4822, 0.4465)
rgb_std = (0.2023, 0.1994, 0.2010)

transform_train = transforms.Compose([
    transforms.RandomCrop(32, padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize(rgb_mean, rgb_std),
])
```

<https://pytorch.org/vision/stable/transforms.html>

Transfer Learning (supervised pre-training)

- Initialize the weights of your model using the optimal weights learnt on a similar task!
- Fine-tuning can provide a reasonably good model even when we have very few training data.

If you know how to recognize...



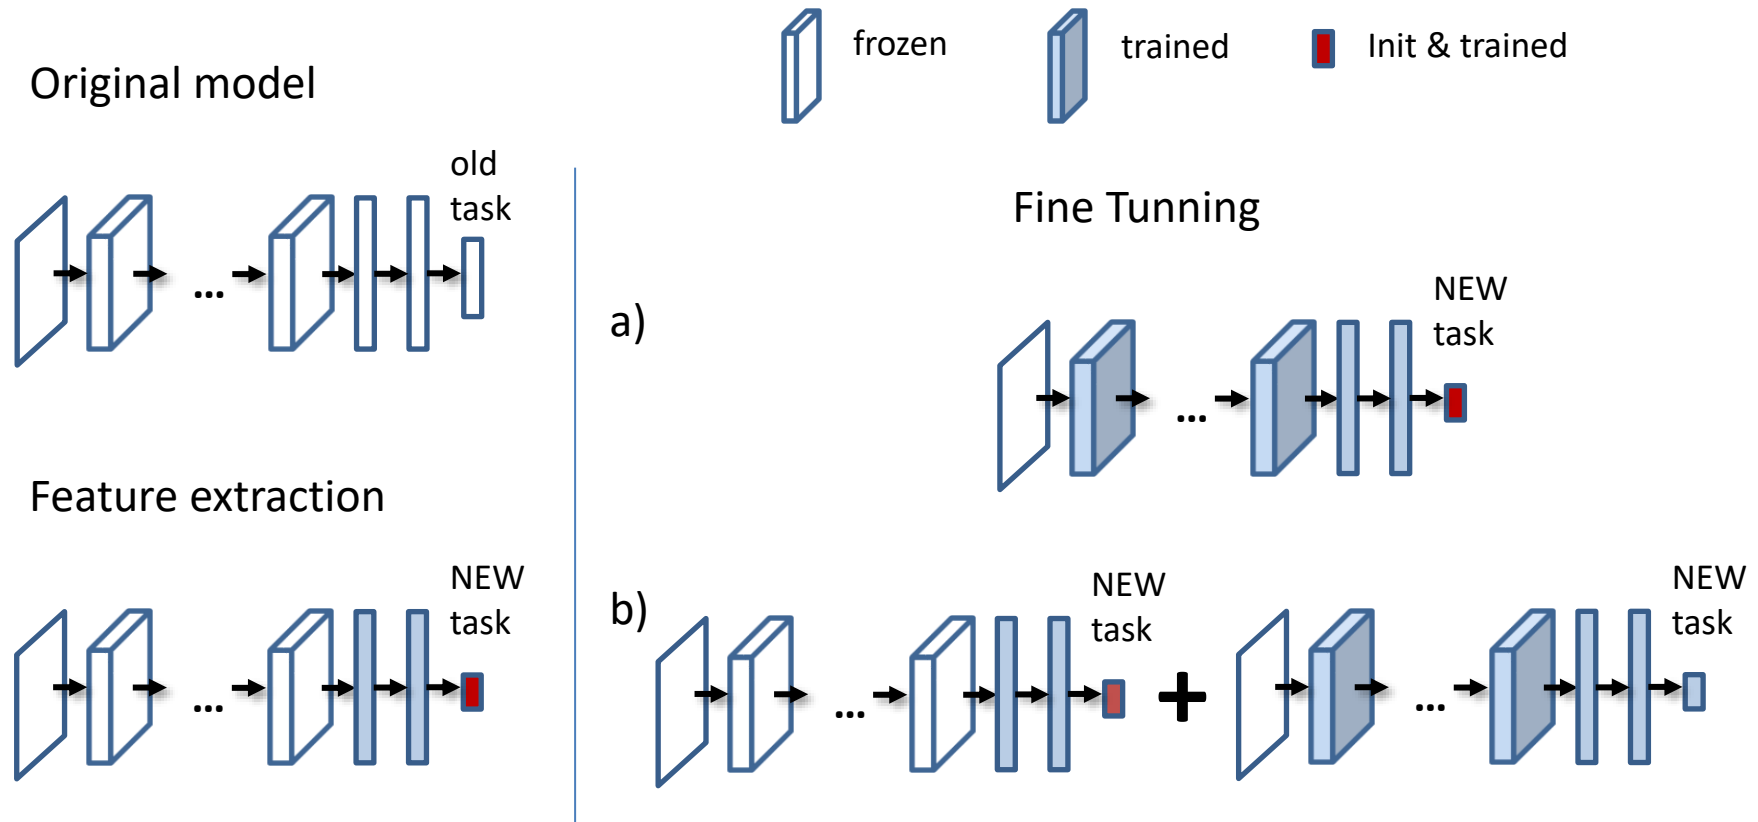
Transfer Learning

You will be able to recognize...



Image source: Yannis Ghazouani, 2016.

Transfer Learning (supervised pre-training)



Sharif Razavian, Ali, et al. "CNN features off-the-shelf: an astounding baseline for recognition." Proceedings of the IEEE CVPRW. 2014.

Transfer Learning

In practice, very few people train a CNN from scratch (with random initialization). Don't have sufficient data and/or resources.

It is common to use pretrained CNNs and use them either as an initialization or a fixed feature extractor for the new task.

When and how to transfer learning?

- New dataset is small and similar to original dataset: feature ext.
- New dataset is large and similar to the original dataset: feat ext+fine tune
- New dataset is small but very different from the original dataset: feat ext
- New dataset is large and very different from the original dataset: pre trained

Practical advice

- Constraints from pretrained models.
- Learning rates.